Qumulo Administrator Guide



Copyright © 2023 Qumulo, Inc.

Table of Contents

Getting Started Supported Configurations and Limits......4 System Upgrades Performing Upgrades......7 Authentication Configuring SAML Single Sign-On (SSO)31 External Services **Network Configuration** Required Networking Ports......57 Network Multitenancy Web UI qq CLI **Snapshots** How Snapshots Work 79

Encryption and Data Security

Managing Encryption at Rest	92
Generating and Storing ECDSA Keys	94
Managing Security Keys	100
Node Addition and Replacement	
Adding Nodes to an Existing Cluster	105
Increasing the Node Fault Tolerance Level during Node-Add Operations	107
Adding Nodes with Adaptive Data Protection	108
Replacing Nodes by Performing a Transparent Platform Refresh	
Improving Performance by Migrating a Cluster to a Different License Class	117
Understanding Offline Nodes and Checking for Free Space	118
Data Replication	
Shift-To Amazon S3	120
Shift-From Amazon S3	131
File System Changes	
How File System Change Notifications Work	143
Watching for Changes with SMB2 CHANGE_NOTIFY	147
Watching for Changes with REST	151
File System Protocols	
Enabling and Using NFSv4.1	157
Managing File Access Permissions with ACLs	164
NFSv4.1 with Kerberos	
How NFSv4.1 Works with Kerberos	173
Prerequisites for Joining to Active Directory	174
Configuring Active Directory	176
Performing Additional Cluster Configuration	180
Using Kerberos Permissions	183
Configuring a Linux Client	189
Configuring Cross-Domain Active Directory Trusts	199
Troubleshooting NFSv4.1 with Kerberos	201

S3 API

Configuring and Using the S3 API	204
Creating and Managing S3 Access Keys	209
Creating and Managing S3 Buckets	215
Managing Access to S3 Buckets	226
Managing Multipart S3 Uploads	232
Supported Functionality and Limits	238
Monitoring and Metrics	
Enabling Cloud-Based Monitoring and Remote Support	244
Connecting to Cloud-Based Monitoring and S3 with a Custom Proxy	250
Restoring the Default Values for Cloud-Based and Nexus Monitoring	252
OpenMetrics API Specification	254

Getting Started

Supported Configurations and Known Limits for Qumulo Core

This section provides an overview of supported configurations and known limits for Qumulo Core.

Supported Configurations

Configuration Type	Supported Value
Protocols	 FTP NFSv3 NFSv4.1 (page 0) S3 API (page 0) SMB 2.002 SMB 2.1 SMB 3.0 SMB 3.1 SMB 3.1.1
Browser	Google Chrome 80 (and higher)
Clients over SMB	macOS 10.14 (and higher)Windows 7 (and higher)
Clients over NFS	macOS 10.14 (and higher)Linux Kernel 2.6.X (and higher)
Linux Configuration	Qumulo Core is up to date with all Ubuntu 18.04 security updates.

Configuration Type	Supported Value
Domain-Functional Level	Microsoft Windows Server 2008 R2 (and higher)
	⑤ Note Qumulo Core doesn't support Samba Domain Controllers.
Kerberos V5 Encryption Types	RC4-HMAC-MD5AES256-CTS-HMAC-SHA1AES128-CTS-HMAD-SHA1
LDAP Servers	OpenLDAP for Group Expansion
Python Version for qq CLI	3.8 (and higher)

Known Limits

Limit Type	Maximum Value
On-Premises Cluster Size	265 nodes
Cloud Cluster Size	100 nodes
Floating IP Addresses for Each Node	10
NFS Exports	64,000
SMB Shares	40,000
Access Control Entries (ACEs) in an Access Control List (ACL)	200
NFS Groups	16 without RFC 2307 with Kerberos
Combined Users and Groups	4 billion
Characters in Cluster Name	2-15, alphanumeric and hyphen (-)
Characters in Full Path (Path Name)	32,760 (limited by protocol)

Limit Type	Maximum Value	
Characters in File Path Component (File or Directory)	255 (limited by protocol)	
Files in a Directory	4.3 billion	
File Size	9 exabytes	
Number of Files	18 quintillion	
Hard Links for Each File	1,024	
LDAP Domains	1	
Active Directory Domains	1	
DNS Servers	3	
Snapshots	40,000	
Quotas	4.3 billion	
	1 Note This approximate value of 2^{32} is equivalent to the maximum number of directories or the entire inode space.	
Number of Replication Rela-	100	
tionships	• Note If a directory is more than 100 levels below the file system root directory, you can't use it as a replication source.	
Usable Provisioned Capacity on Cluster	100%	

System Upgrades

Performing Instant Software Upgrades and Platform Upgrades

This section explains the difference between Instant Software Upgrades and Platform Upgrades. For more information, see Performing Qumulo Core Upgrades by Using the qq CLI on Qumulo Care.

▲ Important

- If you perform multiple upgrades back to back, you might encounter one or more
 platform upgrades in one of the incremental releases; you must install these
 upgrades before you continue. Before performing back to back upgrades, contact the
 Qumulo Care team for guidance.
- If you don't see a rolling reboot option for a platform upgrade, refresh the page in your browser.

For downloads, release notes, and upgrade paths for on-premises and cloud releases of Qumulo Core, see Qumulo Nexus.

Understanding the Differences Between Upgrade Modes

For information about which upgrade modes different Qumulo Core releases use, see Qumulo Core Upgrade Mode Reference (page 10).

Instant Software Upgrade

The more common, faster instant software upgrade requires restarting only the container on your nodes and has a downtime of less than 30 seconds without disruption to the operation of the cluster.

Because in Qumulo Core 3.3.2 (and higher), the Qumulo file and data protection systems are separate from the host in charge of running the operating system and the services specific to each hardware or cloud platform, and because these services run in a lightweight container (by using Ubuntu-native systemd-nspawn containerization) in the user space, it is possible to move quickly from one version of Qumulo Core to another by loading a new container and pointing the runtime environment at updated software.

O Note

- Under certain conditions, an end-to-end instant software upgrade might take a little longer while Qumulo Core performs background tasks. This doesn't impact user experience.
- Instant software upgrades don't impact existing support for the qq CLI or REST API commands.
- A direct upgrade to Qumulo Core 3.3.3 isn't an instant software upgrade (it only establishes the framework for this functionality). Upgrading from Qumulo Core 3.3.2 to 3.3.3 is the first official, minimally disruptive instant software upgrade. Any subsequent upgrade, regardless of release, is an instant software upgrade unless we specify otherwise.

Platform Upgrade

The infrequent, somewhat slower platform upgrade requires either a *complete reboot* (rebooting all nodes in your cluster at the same time with SMB and NFS client interruption) or—in Qumulo Core 5.0.3 (and higher)—a *rolling reboot* (rebooting the nodes in your cluster one at a time with impact only to SMB clients).

The reboot process differentiates platform upgrades from instant software upgrades, where your nodes maintain the Linux operating system (and certain services that Qumulo Core relies on) during the upgrade process and where the underlying host changes less frequently than the container and the file and data protection systems.

▲ Important

- If you perform multiple upgrades back to back, you might encounter one or more
 platform upgrades in one of the incremental releases; you must install these
 upgrades before you continue. Before performing back to back upgrades, contact the
 Qumulo Care team for guidance.
- If you don't see a rolling reboot option for a platform upgrade, refresh the page in your browser.
- Upgrading past a platform upgrade still requires a node reboot, even if you don't
 install the exact build. For example, if your cloud cluster runs Qumulo Core 4.0.0,
 installing Qumulo Core 4.1.0 triggers a platform upgrade, because the installation
 includes all changes contained in Qumulo Core 4.0.6 that is a platform upgrade.

Understanding the Upgrade Phases

Every Qumulo Core upgrade has two phases, preparation and commit.

- 1. Preparation: Qumulo Core stages a new image in an alternate boot drive partition while the current image continues to run. This phase is responsible only for the background work (unpacking and writing the platform image and upgrade firmware, and so on). When the preparation phase is complete, we continue to the commit phase.
- 2. Commit: Qumulo Core does one of the following:
 - · Instant Software Upgrade: Stops the existing container and starts a new one.
 - Platform Upgrade: Initiates a reboot and selectively upgrades the operating system image.

To determine what phase an upgrade is in, use the qq upgrade_status command while your cluster is performing and upgrade. For more information, see Performing Qumulo Core Upgrades by Using the qq CLI on Qumulo Care.

The following is example output from the command.

```
{
  "install_path": "/upgrade/qinstall.qimg",
  "state": "UPGRADE_PREPARING",
  "details": "",
  "error_message": "",
  "error_state": "UPGRADE_ERROR_NO_ERROR",
  "is_blocked": false,
  "blocked_reason": ""
}
```

Qumulo Core Upgrade Mode Reference

This section provides a reference for Qumulo Core upgrade modes from version 3.3.3 onwards.

- An instant software upgrade (page 7) requires restarting only the container on your nodes and has a downtime of less than 30 seconds without disruption to the operation of the cluster.
- A *platform upgrade* (page 8) requires either a complete reboot (rebooting all nodes in your cluster at the same time) or a rolling reboot (rebooting the nodes in your cluster one at a time).
- A *quarterly upgrade* aggregates all improvements and fixes since the last quarterly upgrade. The version number of a quarterly upgrade ends in .0.

A Important

Although the *upgrade types* for on-premises upgrades and cloud upgrades are most often the same, they do occasionally diverge. For example, for Qumulo Core 5.3.1, a cloud deployment allows an instant upgrades, an on-premises deployment requires a platform upgrade.

For downloads, release notes, and upgrade paths for on-premises and cloud releases of Qumulo Core, see Qumulo Nexus.

Qumulo Core Upgrade Modes

For information about the most important features from each release, click the Qumulo Core version.

Version	On-Premises Upgrade Type	Cloud Upgrade Type
6.1.2.2 (page 14)	Platform	
6.1.1 (page 14)	Platform	
6.1.0.1 (Quarterly) (page 15)	Instant	
6.0.2 (page 15)	Instant	
6.0.1.1 (page 15)	Platform	
6.0.0.2 (Quarterly) (page 16)	Instant	
5.3.4 (page 17)	Instant	

Version	On-Premises Upgrade Type	Cloud Upgrade Type
5.3.3.1 (page 17)	Instant	
5.3.2 (page 17)	Instant	
5.3.1 (page 17)	Platform Instant	
5.3.0 (Quarterly) (page 18)	Instant	
5.2.5.1 (page 18)	Inst	ant
5.2.4 (page 18)	Inst	ant
5.2.3 (page 19)	Inst	ant
5.2.2 (page 19)	Inst	ant
5.2.1 (page 19)	Inst	ant
5.2.0.2 (Quarterly) (page 19)	Instant	
5.1.5 (page 19)	Platform	
5.1.4.1 (page 20)	Instant	
5.1.3 (page 20)	Instant	
5.1.2 (page 20)	Instant	
5.1.1 (page 20)	Platform	
5.1.0.1 (Quarterly) (page 21)	Instant	
5.0.6 (page 21)	Instant	
5.0.5 (page 21)	Instant	
5.0.4 (page 21)	Instant	
5.0.3 (page 22)	Instant	
5.0.2 (page 22)	Instant	
5.0.1 (page 22)	Instant	
5.0.0.1 (Quarterly) (page 22)	Instant	

Version	On-Premises Upgrade Type	Cloud Upgrade Type
4.3.4 (page 23)	Instant	
4.3.3 (page 23)	Instant	
4.3.2 (page 23)	Instant	
4.3.1 (page 23)	Instant	
4.3.0 (Quarterly) (page 23)	Inst	tant
4.2.6 (page 24)	Inst	tant
4.2.5 (page 24)	Inst	tant
4.2.4 (page 24)	Platform	
4.2.3 (page 24)	Instant	
4.2.2 (page 24)	Instant	
4.2.1 (page 24)	Platform	
4.2.0 (Quarterly) (page 25)	Instant	
4.1.5 (page 25)	Instant	
4.1.4 (page 25)	Instant	
4.1.3 (page 25)	Instant	
4.1.2 (page 25)	Instant	
4.1.1 (page 25)	Instant	
4.1.0.1 (Quarterly) (page 25)	Instant	
4.0.6 (page 26)	Instant	Platform
4.0.5 (page 26)	Instant	
4.0.4 (page 26)	Instant	
4.0.3 (page 26)	Instant	
4.0.2 (page 26)	Instant	

Version	On-Premises Upgrade Type	Cloud Upgrade Type
4.0.1.1 (page 26)	Instant	
4.0.0.2 (Quarterly) (page 26)	Instant	
3.3.5 (page 27)	Instant	
3.3.4 (page 27)	Instant	
3.3.3 (page 27)	Instant	

Qumulo Core Feature Log

This section lists the most important features from each release.

For information about upgrade types for each release, see Qumulo Core Upgrade Mode Reference (page 10).

Qumulo Core 6.1.2.2

- Added Adaptive Data Protection to the REST API and qq CLI and made performance improvements
- · Added Transparent Platform Refresh to the REST API and qq CLI
- · Added S3 configuration settings to the Web UI
- · Improved hard link creation

A Important

- · This release resolves:
 - An issue that affects parsing Kerberos principal names that contain special characters in Active Directory Kerberos authentication in Qumulo Core 6.1.2.
 This issue doesn't prevent access to SMB shares by using NTLM-based authentication.
 - An issue that affects upgrading Qumulo Core on the Quiver 1U Hybrid Gen2 platform in Qumulo Core 6.1.1.
 - An issue that affects getting a useful core dump in Qumulo Core 6.1.1.
- We've identified a known issue that affects event processing in the Cloud-Based Monitoring (CBM) service Qumulo Core 6.1.0.1. This issue is resolved in Qumulo Core 6.2.0.

Qumulo Core 6.1.1

- Added support for new S3 API actions, snapshot locking for replication target relationships, and lock configuration to the Web UI
- · Enabled the Search Trusted Domains Active Directory configuration option by default
- Upgraded the host and container images
- · Revised the REST API for snapshots
- Added qq CLI commands for lock key functionality

- · Added a qq CLI command for viewing blocked upgrade status
- · Updated the OpenMetrics API Specification
- · Made general improvements to Qumulo Core

Qumulo Core 6.1.0.1 (Quarterly)

O Note

This release resolves an issue where, under certain circumstances, while using Adaptive Data Protection in Qumulo Core 6.1.0 (and lower), it becomes impossible to add nodes to a cluster during normal cluster expansion or Transparent Platform Refresh operations, until you upgrade the cluster to version 6.1.0.1 (or higher).

- · Configured Qumulo Core to use recursive change notifications by default
- · Enabled SMB multichannel by default
- · Added snapshot locking functionality by using the qq CLI
- · Made usability changes to event logging and the Web UI
- · Removed beta multitenancy REST APIs and qq CLI commands

Qumulo Core 6.0.2

- · Added support for streaming file change notifications
- · Improved the S3 API
- · Made improvements to the Qumulo Core allocation algorithm
- Added audit logging and REST API and qq CLI commands for streaming file change notifications
- · Removed v2 Python bindings for NFS export and SMB share APIs
- · Added new fields to the OpenMetrics API
- Fixed issues with creating SMB shares, transparent platform refresh, vSphere clusters, upgrades in the Web UI, making SMB2 copies by using ODX, and creating snapshot policies

Qumulo Core 6.0.1.1

- Added new SMB2 CHANGE NOTIFY functionality
 - Full support for recursive watch-tree notifications
 - Alternate Data Streams

- Numerous performance improvements
- Elimination of missed events
- · Added read combining for NFSv4.1
- · Improved the S3 API
 - S3 API action traversal rights
 - S3 traffic in the Web UI and in qq CLI output
 - Presigned URL functionality
 - HTTP identity transfer encoding for the GetObject API action
- · Added operational information to audit logging in JSON format
- · Made numerous updates to the multitenancy and S3 qq CLI commands and REST APIs
- Fixed an issue with Network Interface Card (NIC) addressing on 240TB and 480TB General Purpose models of the HPE Apollo 4200 Gen10 Plus node

Qumulo Core 6.0.0.2 (Quarterly)

A Important

- This release resolves an issue with transparent platform refresh that we discovered in Qumulo Core 6.0.0 and 6.0.0.1.
- This release restores the status of the /v1/smb/shares REST API endpoints (previously removed in Qumulo Core 5.3.4) to deprecated. We strongly recommend changing over to /v2/smb/shares endpoints as soon as possible.
- · Added Transparent Platform Refresh for node replacement
- · Enabled concurrent node reboots for rolling upgrades
- · Added write combining for NFSv4.1
- · Introduced optional HTTP mode for the S3 API
- · Made SSL certificate changes for the S3 API without a cluster restart
- · Added audit syslog in JSON format
- · Made corresponding changes and additions to the REST API and qq CLI
- Fixed issues with the Web UI, directory access by non-root users on Linux, and changed caching for some OpenMetrics API metrics

Oumulo Core 5.3.4

A Important

The removal of the deprecated /v1/smb/shares REST API endpoints in this release can affect certain third-party backup or migration workflows.

- Added data protection configuration, multitenancy, tenant-aware SMB and NFS, and automatic clean-up of stale multipart S3 uploads
- Removed all REST API endpoints under /v1/smb/shares
- Improved and added REST APIs and qq CLI commands for restriper, multitenancy, and tenant-aware SMB and NFS operations
- · Improved the OpenMetrics APIs
- Fixed an issue with SMB2 attributes for symlinks and a potential security issue with REST API session tokens

Qumulo Core 5.3.3.1

O Note

This release resolves an issue that we discovered in Qumulo Core 5.3.3.

- · Put out the initial release of the S3 Protocol Server
- · Added support for NFSv4.1 with KRB5i
- · Added the ability to disable access tokens
- · Removed /v1/session/kv REST API endpoints
- · Fixed a bug with listing trustees for a role and a bug with LDAP connections

Oumulo Core 5.3.2

- · Added distributed multi-AZ clustering to Qumulo clusters on AWS
- Added SMB2 CHANGE_NOTIFY parity with Windows
- · New Qumulo OpenMetrics API metrics
- · Added expiration time for access tokens
- · Improved the REST API

Oumulo Core 5.3.1

 \cdot Added support for watching for changes in file attributes and directory entries with SMB2

CHANGE NOTIFY filters

- · Updated all firmware for all variants of Mellanox CX5 and CX6 and Broadcom NICs
- · Added metrics to the Qumulo OpenMetrics API
- · Added two-factor authentication for qq-CLI-initiated SAML SSO workflows
- · Improved REST API and qq CLI
- · Added a missing library dependency that allows SAML SSO to work on cloud platforms

Qumulo Core 5.3.0 (Quarterly)

- Added interactive SAML SSO login and SSO for Active Directory users who manage clusters
- · Identified a known issue with SAML SSO on cloud platforms
- Improved SMB CHANGE_NOTIFY requests
- · Released the Qumulo OpenMetrics API
- · Released the Access Tokens API
- · Improved the Web UI
- · Added REST APIs for node identification
- · Improved qq CLI commands

Qumulo Core 5.2.5.1

O Note

This release resolves an issue that we discovered in Qumulo Core 5.2.5.

- · Allowed clusters joined to Active Directory to use SAML 2.0 SSO
- · Added support for displaying NFSv4.1 quotas
- · Created a more balanced usage of free space on clusters after the addition of nodes
- · Improved the Web UI
- · Added a flag to a qq CLI command
- · Fixed bugs related to quotas and Alternative Data Streams

Oumulo Core 5.2.4

The majority of functionality and stability improvements in this Qumulo Core release are internal.

Qumulo Core 5.2.3

The majority of functionality and stability improvements in this Qumulo Core release are internal.

Qumulo Core 5.2.2

- · Made improvements to the Web UI on the NFS Exports and Role Management pages
- Adjusted the REST API endpoint /v1/smb/sessions and the corresponding qq CLI command
- · Fixed an issue that occurs during NFSv4.1 failover
- · Fixed an issue that prevents a cluster from forming due to an incorrect HDD-SSD ratio.

Qumulo Core 5.2.1

- · Added over-the-wire encryption to NFSv4.1 with Kerberos
- · Improved performance for SMB Find requests
- · Added an optional Web UI login banner
- Removed /v1/upgrade/ endpoints
- · Fixed a bug on the NFS Exports page.

Qumulo Core 5.2.0.2 (Quarterly)

- Enabled local trustees for NFSv4.1 with Kerberos
- · Added support for Amazon EC2 M6i instances and Amazon EBS gp3 volumes
- Removed and replaced REST API endpoints related to shutting down and stopping clusters and nodes
- · Fixed an SMB client bug
- Resolved a file system operations latency issue with continuous replication jobs and on clusters with high file creation load
- · Resolved an issue that caused a node to enter a reboot loop under certain conditions.

Qumulo Core 5.1.5

- Enabled creating mounts by using NFSv4.1 with Kerberos
- · Improved NLM and NFSv4.1 locking
- · Updated Mellanox NIC firmware
- Made improvements to the Web UI on the APIs & Tools, Add Nodes, and Active Directory pages

- · Added NFS fields to the REST API and qq CLI
- Deprecated the /v2/upgrade/status REST API
- · Updated security permissions for home directories
- · Prohibited the use of the switch-user command
- · Added support for user principal (UPN) names

Qumulo Core 5.1.4.1

- · Improved workload read performance for streaming
- Deprecated the /v1/session/kv/ REST API endpoints
- Fixed an issue on the FTP Settings page
- Fixed a bug where unexpected behavior from an NFS client can cause issues with memory usage on a Qumulo cluster.

Qumulo Core 5.1.3

- · Added NFSv4.1 byte-range locking support for clusters that have NFSv4.1 enabled
- Added custom naming functionality for snapshots
- · Made improvements to the Web UI
- · Improved node fault tolerance for node-adds
- · Updated the time synchronization for cloud clusters
- Added a REST API to calculate node-add capacity
- Fixed bugs about adding S3 privileges, making OPEN requests over NFSv4.1, and reassigning floating IP addresses.

Oumulo Core 5.1.2

- · Improved user experience for NFSv4.1 ACL editor
- Made improvements to the Web UI
- · Updated time synchronization APIs for QaaS Azure clusters
- · Fixed bugs with malformed RPCBIND requests and replication relationships.

Qumulo Core 5.1.1

- · Improved time management on QaaS Azure clusters
- · Improved resiliency against data center accidents
- · Made improvements to the Web UI

- · Updated the SDK to Python 3.8
- · Disabled time management APIs on QaaS Azure clusters
- · Fixed SMB and Web UI form bugs.

Qumulo Core 5.1.0.1 (Quarterly)

- · Improved write performance for workloads on clusters with high-churn, small-file datasets
- · Added a new Active Directory capability
- · Made a number of pages in the Web UI more clear
- · Improved SMB and NFS share, export management, and status APIs
- · Added inactivity timeout to the Web UI
- · Identified a regression in 5.0.5 and fixed a bug that occurs for rare NFS client applications.

Qumulo Core 5.0.6

- · Improved read performance for previously read large files
- · Improved write performance for certain platforms and workloads
- · Improved cluster resilience on Azure
- · Added support for HPE Broadcom NICs and updated the NIC and SKU reference
- · Made improvements to the Web UI
- Fixed bugs pertaining to snapshots of directories with SMB-illegal characters and localized Active Directory installations.

Qumulo Core 5.0.5

- · Increased the tolerance of uncertainty for NTP
- · Mapped Isilon-style and Samba-style NFS SIDs to Qumulo SIDs
- Made changes to the Python SDK
- · Made an XSS security fix
- · Resolved two bugs

Qumulo Core 5.0.4

- · Enabled advanced data integrity for new clusters
- · Added a check to confirm whether an upgrade can begin
- · Added an initial release of the S3 protocol

- Deprecated reboot qq CLI commands
- · Added access time (atime) API resources and qq CLI commands.

Qumulo Core 5.0.3

- · Added Rolling Reboot for platform upgrades (including new APIs and CLI commands)
- Updated the Web UI for both software and platform upgrades
- Made improvements to write performance over NFSv3
- · Improved the Add Nodes and Snapshots pages in the Web UI

Qumulo Core 5.0.2

- · Changed all HPE firmware updates to be performed using the Qumulo installer
- · Improved cluster resilience against boot drive failures
- · Added the ability to set timestamps through NFSv4.1
- Added the access_time field to the attributes REST API resource.

Qumulo Core 5.0.1

• Enforced the Qumulo-to-Qumulo replication compatibility guarantee to 3 consecutive quarterly releases

☑ Tip

For example, 5.0.1 can't replicate with versions lower than 4.3.0 or higher than 5.2.0.

- · Made improvements to Shift-to-S3
- · Removed unconfigured node APIs

Qumulo Core 5.0.0.1 (Quarterly)

- Changed the Qumulo-to-Qumulo replication compatibility guarantee from 2 to 3 consecutive quarterly releases
- · Enforced Kerberos ticket expiration
- · Enabled NLM persistence by default
- · Added sorting and filtering to the Snapshot Policies page
- · Improved speed of deleted snapshot clean-up
- Removed requirements for QaaS clusters
- · Added support for Python 3.8

- · Improved the API and CLI version and user-querying command
- Resolved an issue that causes an unpredictable node failure if two NFSv4.1 clients have the same hostname.

Qumulo Core 4.3.4

A Important

This release of Qumulo Core resolves a critical security vulnerability. We strongly recommend upgrading your Qumulo cluster to 4.3.4 to resolve this issue.

- · Added work estimation to Shift-from-S3 relationships
- · Added a feature that improves cluster resiliency for unresponsive devices
- · Changed the kiosk timeout window to 15 minutes

Oumulo Core 4.3.3

- · Prevented authentication requests with PAC buffers of unknown type from failing
- · Improved Saved Snapshots page in the Web UI
- Added a mechanism for monitoring status of front-end network on platforms with separate interfaces.

Qumulo Core 4.3.2

Fixed a bug in Qumulo Core's handling of cross-protocol permissions

Qumulo Core 4.3.1

- · Improved write performance
- · Improved speed and duration of terminating outstanding requests to AD Controllers
- · Added file system permission checks for SMB shares and NFS exports
- · Enabled mounting NFS over UDP with floating IP addresses

Qumulo Core 4.3.0 (Quarterly)

- Added Network File System Version 4.1 (NFSv4.1)
- · Improved performance for SMB offloaded data transfer (ODX) and SMB directory query
- · Allowed AWS Sidecar to handle a new class of EBS failures
- · Fixed bugs in Shift-From-S3 and the Web UI

Qumulo Core 4.2.6

- · Made an enhancement to Web UI for Qumulo on Azure
- · Adjusted a REST resource
- Made an important change to export names that end with a slash (/) or are longer than
 64 KiB

Oumulo Core 4.2.5

Added Web UI management support for Qumulo Shift-from-S3 operations

Qumulo Core 4.2.4

- · Enhanced tree-delete operations
- · Made performance improvement for C-72T, C-168T, C-192T, and C-432T platforms
- · Enhanced Shift-to-S3
- · Made performance improvements
- · Made several important bug fixes

Oumulo Core 4.2.3

- · Added Qumulo Shift-from-S3
- · Web UI Update for SMB Shares
- · Enabled DNS configuration for cloud deployments in the Web UI

Qumulo Core 4.2.2

- · Improved the appearance of spinners in the Web UI
- · Made a failover procedure bug fix

Qumulo Core 4.2.1

· Made improvements for security identifiers

A Important

We removed this feature for compatibility reasons. For more information, see the release notes on Qumulo Nexus.

- · Upgraded the Linux kernel to 5.4
- · Improved multistream throughput for 4U nodes
- · Updated qq CLI commands for unconfigured node cluster creation

Fixed a bug

Qumulo Core 4.2.0 (Quarterly)

- · Enhanced snapshot performance
- · Enhanced Active Directory connections
- · Made an important bug fix

Qumulo Core 4.1.5

- · Added details to the Node Add page in the Web UI
- · Updated qq CLI upgrade commands for unconfigured nodes
- · Made improvements to snapshot cleanup
- · Made an important bug fix

Qumulo Core 4.1.4

- · Updated user privileges for the quota qq CLI commands
- · Enhanced autocomplete functionality in the Web UI
- · Made several important bug fixes

Qumulo Core 4.1.3

The majority of functionality and stability improvements in this Qumulo Core release are internal.

Qumulo Core 4.1.2

- Enhanced security of qq CLI commands
- · Made several important bug fixes

Oumulo Core 4.1.1

- · Added the ability to repeat completed jobs for Shift
- · Enhanced the security flag for SMB connections
- · Improved directory-caching performance
- · Updated the HPE Field Verification Tool

Qumulo Core 4.1.0.1 (Quarterly)

- · Enhanced the upgrade experience
- · Improved the Qumulo Sidecar for AWS

· Made several important bug fixes

Qumulo Core 4.0.6

Made a platform upgrade for cloud clusters

Qumulo Core 4.0.5

- · Updated the SMB API and CLI
- · Made several important bug fixes

Qumulo Core 4.0.4

- Updated the firmware for C-72T, C-168T, C-192T, C-432T, K-144T, K-168T, and K-432T platforms
- Enhanced qq CLI commands
- · Fixed an important bug for mixed-node clusters

Qumulo Core 4.0.3

- · Enhanced the web server security
- · Fixed an important bug fix for mixed-node clusters

Qumulo Core 4.0.2

- · Improved the performance of SSD clusters
- · Made minor performance enhancements
- · Fixed an important bug in the Web UI

Oumulo Core 4.0.1.1

Fixed important bugs in:

- · Cloud-Based Monitoring
- · SMB file enumeration
- · Compound API and qq CLI request errors

Qumulo Core 4.0.0.2 (Quarterly)

- · Deprecated the Python 2.7 SDK
- · Made several important bug fixes

Qumulo Core 3.3.5

- · Updated mixed-node compatibility
- · Made a performance improvement for Qumulo Shift
- Improved the Field Verification Tool for HPE Apollo 4200 Gen10 clusters
- · Adjusted audit logging functionality
- · Made a few important bug fixes

Qumulo Core 3.3.4

- · Added support for the HPE ProLiant DL325 Gen10 Plus platform
- · Made general performance enhancements for hybrid clusters

Qumulo Core 3.3.3

- · Added Instant Software Upgrade
- · Made Qumulo Shift performance improvements and Web UI enhancements
- · Updated the HPE Service Pack Pro (SPP) Update for HPE Apollo 4200 Gen10 clusters

Authentication

Installing VPN Keys on a Qumulo Cluster

This section explains how to install VPN keys on your Qumulo cluster over a network. You can install the VPN keys by using the qq CLI from a machine on the same network as your cluster or from one of your nodes.

Prerequisites

Before you begin, make sure that you have done the following.

- · Obtain a .zip file with VPN keys from Qumulo Care
- · Whitelist the following domains in your firewall rules:
 - ∘ ep.qumulo.com
 - missionq.qumulo.com
 - api.missionq.qumulo.com
 - monitor.qumulo.com
 - api.nexus.qumulo.com
- · Permit outbound HTTPS traffic on port 443

O Note

If your firewall performs stateful packet inspection (also known as *SPI* or *deep-packet inspection*), you must allow OpenVPN (SSL VPN) explicitly, rather than only open port 443.

To Install VPN Keys from a Networked Machine

- 1. Copy the .zip file from Qumulo Care to a computer on the same network as your cluster, and decompress the file.
- 2. Install the qq CLI on the same computer. For more information, see QQ CLI: Get Started on Qumulo Care.
- 3. To log in to your cluster, use the qq CLI and specify your cluster's IP address. For example:

```
qq --host 203.0.113.0 login
```

O Note

Your user must have PRIVILEGE SUPPORT WRITE and PRIVILEGE SUPPORT READ.

4. To install the VPN keys on your cluster, specify your cluster's IP address and the path to the directory that contains the VPN keys. For example:

```
qq --host 203.0.113.0 install_vpn_keys /my/path
```

5. To verify that the VPN keys installed correctly, use the get_vpn_keys command. For example:

```
qq --host 203.0.113.0 get_vpn_keys
```

6. Remove any local copies of the VPN key files.

To Install VPN Keys from a Node

O Note

On macOS and Linux, you can use the scp and ssh tools. On Windows Server 2022, Windows Server 2019, and Windows 10 (build 1809 and higher), we recommend installing OpenSSH.

- 1. Copy the .zip file from Qumulo Care to a computer on the same network as your cluster, and decompress the file.
- 2. To copy the VPN key files to one of your nodes, use the scp command. For example:

```
scp /my-path/* admin@203.0.113.0:~/
```

3. To connect to the node to which you copied the VPN key files, use the ssh command. For example:

```
ssh admin@203.0.113.0
```

The qq CLI is available to the admin user. For example:

qq version

4. To install the VPN keys on your cluster, specify the path to the directory that contains the VPN keys. For example:

```
sudo qq install_vpn_keys /my/path/
```

5. To verify that the VPN keys installed correctly, use the get_vpn_keys command. For example:

```
sudo qq get_vpn_keys
```

To Register Cluster with Cloud-Based Monitoring

- 1. To retrieve your cluster ID, use the node_state_get command.
- 2. Send the output of the command to Qumulo Care.
- 3. Use the Web UI to enable Qumulo Care Remote Support.
- 4. Notify Qumulo Care when this process is complete.

Qumulo Care verifies your VPN functionality and then adds your cluster to Cloud-Based Monitoring.

Configuring SAML Single Sign-On (SSO) for Your Qumulo Cluster

This section explains how to integrate your Qumulo cluster with your organization's single sign-on (SSO) service by configuring Security Assertion Markup Language (SAML) 2.0 for Qumulo Core 5.2.5.1 (and higher).

For more information about the SAML standard for exchanging authentication information, see SAML 2.0.

Prerequisites

Before you begin, make sure that you have done the following.

 Join your cluster to an Active Directory (AD) domain. For more information, see Join Your Qumulo Cluster to Active Directory on Qumulo Care.

O Note

Qumulo Core supports SAML authentication only for AD users.

- To allow the cluster to find group memberships for SAML-authenticated users, configure the Base DN in your AD configuration, even if you don't use POSIX attributes.
- Ensure that your SAML Identity Provider (IdP) is linked to the same AD. An *identity* provider (such as Azure AD, Duo, or Okta) is a system that authenticates users (for example, by using passwords and additional factors).

Typically, an IT department manages an IdP centrally and the IdP is linked with AD. Before you can enable SSO, your IT department must register a new Service Provider (SP) in your IdP. A *service provider* is the server which users access, in this case a Qumulo cluster.

O Note

You can use trusts, as long as the Base DN covers all users that might require access to your cluster.

Configure your IdP to return AD User Principal Names (UPNs, for example
 alice@example.com) or an email address as a NameID. A NameID is an identifier for an
 authenticated user. Typically, a NameID uses the format of an email address.

To Configure SAML SSO for Your Qumulo Cluster

This process requires coordination between the cluster administrator and SSO administrator.

- 1. The cluster administrator contacts the SSO administrator and asks the SSO administrator to create a SAML integration for the Qumulo cluster.
- 2. The SSO administrator creates a SAML integration with your organization's SSO identity provider (page 31) (IdP).
 - a. The SSO administrator uses the cluster's fully qualified domain name (FQDN) format for the service provider (page 31) (SP) endpoint (also known as the assertion consumer service URL), in the following format:

https://<my-cluster>.<my-org>.com/saml

O Note

Because the user's browser performs DNS resolution (for example, in a VPN-only scenario), it isn't necessary for an external DNS server to be able to resolve the cluster's FQDN.

- b. If prompted, the SSO administrator enters the HTTP POST binding for the SP endpoint. Typically, this binding is specified by default.
- c. If prompted for SP Entity ID (alternatively named Application Identifier or Audience), the SSO administrator enters .<my-org>.com/saml">https://smy-cluster>.<my-org>.com/saml.
- d. If SAML Signing (depending on the SSO service, this option is named differently) configuration is available, the SSO administrator sets it to Sign SAML response and assertion.

O Note

Qumulo Core requires that IdP sign both the assertion and the entire SAML response.

- 3. After creating the SAML integration, the SSO administrator provides the following information to the cluster administrator.
 - The certificate (public key) of the identity provider, in a __pem_file.
 This certificate lets the cluster verify the authenticity of the messages from the IdP.
 - The IdP SSO URL—to which the Qumulo cluster can send authentication requests—in the following format:

https://<my-org>.<sso-provider>.com/foo

• The IdP issuer or **EntityId**.

Note

Don't confuse EntityId with SP Entity ID.

For example:

```
http://www.<sso-provider>.com/abc12de34fgAB5CDh6i7
```

• The FQDN of the cluster, in the following format:

```
<qumulo-cluster>.<my-org>.com
```

4. To configure and enable SAML login to the Qumulo cluster, the cluster administrator runs the qq saml_modify_settings command. For example:

```
qq saml_modify_settings
    --enable \
    --idp-certificate-file ~/certificate.pem \
    --cluster-dns-name <qumulo-cluster>.<my-org>.com \
    --idp-entity-id http://www.<sso-provider>.com/abc12de34fgAB5CDh6i7 \
    --idp-sso-url https://<my-org>.<sso-provider>.com/abc12de34fgAB5CDh6i7/saml
```

O Note

- To view the current SAML configuration, the cluster administrator can use the qq saml_get_settings command.
- To allow specific changes (for example, correct a typo, update a DNS name or an expired certificate, or temporarily disable SAML SSO without losing any of the other settings), the cluster administrator can use the qq saml_modify_settings command to change individual SAML settings independently.
- For first-time SAML configurations, the cluster administrator must provide all of the required settings.
- Aside from a basic check of the IdP certificate, Qumulo Core doesn't verify the
 configuration parameters. It is the cluster administrator's responsibility to
 ensure that IdP-initiated SAML login works correctly. (This login type initiates
 when the user clicks Continue to SSO login in the Web UI or selects the
 Qumulo cluster on the SSO portal.)

Supported SAML SSO Workflows

Qumulo Core supports three SAML SSO workflows:

- Standard SAML workflows that the IdP (page 31) or SP (page 31) initiates
- A workflow that the qq CLI initiates

Note

- · Members of the built-in Administrators role always have access to the Web UI.
- To allow other users to access the Web UI, you must assign the built-in Observers role to the users.
- Depending on policy, additional verification might be necessary for users. For example, the SSO administrator can enforce mandatory two-factor authentication (2FA) for certain clusters.
- If the user accesses the Web UI by connecting to a node physically, the login page doesn't show doesn't show Continue to SSO login on the Web UI login page, even if SSO is configured.

IdP-Initiated SSO Worfklow

- 1. A user authenticates to her organization's SSO portal and then selects the Qumulo cluster on the SSO portal.
- 2. The SSO portal redirects the user to the cluster's endpoint.

If the user has sufficient privileges, the Web UI logs the user in. Otherwise, the Web UI displays an error message.

SP-Initiated SSO Workflow

- 1. A user navigates to the Qumulo cluster's Web UI endpoint in a browser.
- 2. If the Qumulo cluster has SAML SSO configured, the user can click Continue to SSO login on the Web UI login page.

The Web UI redirects the user to the configured SSO portal. Because the authentication request uses HTTP-Redirect Binding, the login URL appears.

https://<my-org>.<sso-provider>.com/abc12de34fgAB5CDh6i7/saml?SAMLRequest=abcdefgh1234567890...

- 3. The user clicks the login link and the SSO portal authenticates the user.
- 4. The SSO portal redirects the user to the cluster's endpoint.

qq-CLI-Initiated SSO Workflow

In Qumulo Core 5.3.0 (and higher), a user can authenticate a qq CLI session by using SSO.

1. A user uses the qq sso_login CLI command. For example:

```
qq --host 203.0.113.0 sso login
```

The login URL and a prompt appear. The following is an example URL.

https://<my-cluster>.<my-org>.com/saml-login?login-id=12345678-1234-1234-1234-123456789012

O Note

The user must complete the following step within 5 minutes, while the qq CLI pauses for authentication.

- 2. When the user opens the login URL in a browser, the URL redirects the user to a configured SSO portal and one of the following two scenarios takes place:
 - If authentication succeeds, the browser shows a message that contains an eightcharacter verification code and asks the user to return to the CLI session.

The user copies the verification code and enters it into the waiting prompt of the sso_login command.

- If the verification code is correct, the command recognizes that authentication is complete and shows the authenticated username.
- If the verification code is incorrect, the user must retry the workflow.
- · If authentication doesn't succeed, the browser displays an error message.

The user must retry the workflow.

Requiring SSO Authentication for Cluster Management

A Important

- If you use the --require-sso flag, you can no longer use the qq login command with your AD account password. Instead, you must use the qq sso_login command (page 0).
- · This setting doesn't restrict access through file protocols such as SMB.
- Because the FTP protocol sends passwords in plaintext, it is inherently insecure. In addition, many FTP clients don't support Transport Layer Security (TSL) or fall back quietly to the plaintext protocol. For this reason, all Qumulo clusters have FTP disabled by default.

In Qumulo Core 5.3.0 (and higher), you can use the qq saml_modify_settings CLI command to require AD users to use SSO authentication for managing your cluster. For example:

```
qq saml modify settings --require-sso true
```

When the cluster requires SSO authentication, your cluster rejects password-based authentication from AD users in the Web UI, the qq CLI, and the REST API.

Known Issues and Limitations

• Local users (the built-in admin user and any additional users) can always use their passwords to authenticate to the Web UI and the qq CLI.

A Important

We recommend setting a strong password for the built-in admin user and using this account only for emergencies.

- If SSO is required for a Qumulo cluster, it isn't possible to log in to the Interactive API documentation section of the APIs & Tools page in the Web UI.
- · Qumulo Core doesn't support:
 - SAML Single Logout (SLO): We recommend clicking Sign out in the Web UI.
 - Automatic Configuration from Metadata XML: You must specify each parameter by using the qq CLI.
 - Returning to Previous Web UI Page: You can't return to a previous page after reauthenticating (for example, after a timeout).
 - Azure AD SAML Toolkit: Currently, due to a configuration deficiency in the toolkit, IdP-initiated SSO isn't operational for Qumulo as a Service. Use the SP-initiated SSO workflow (page 0).

Troubleshooting SAML SSO Authentication

This section explains troubleshooting common and uncommon SAML SSO authentication issues.

Common Issues

Typically, if SAML authentication fails, Qumulo Core's in-browser error message explains the reasons for failure and you can resolve the issue by setting the right configuration by using the qq saml_modify_settings command. Examples of this issue type include the following scenarios:

- · SAML isn't enabled on the Qumulo cluster.
- There is clock skew between the IdP and the Qumulo cluster (the SSO service sets the clock skew tolerance, typically to 5 minutes).
- The cluster-dns-name or idp-entity-id on the Qumulo cluster aren't configured correctly.
- A user isn't a member of the Observers role that Qumulo Core requires for granting access to the Web UI.

Uncommon Issues

In more complex cases, the in-browser errors are less informative for security reasons. For example, if you configure an incorrect IdP certificate on your cluster, the **Signature validation** failed. SAML Response rejected. error appears.

Several AD configuration issues can cause a User not found error:

· The Qumulo cluster isn't joined to AD.

- The Qumulo cluster is joined to AD that isn't connected to the IdP.
- · IdP sends usernames (NameID) in an unusual format.

To verify that you can use a username, run the qq auth_find_identity command. For example:

```
qq auth_find_identity --name MyUsername
```

· The Configured Base DN doesn't include all users.

To find a security identifier (SID), run the qq auth find identity command. For example:

```
qq auth_find_identity --name MyUsername
```

To verify that a username is discoverable, run the qq ad_sid_to_account command. For example:

```
qq ad_sid_to_account --sid S-1-5-32-544
```

If an error occurs, contact your AD administrator and request the correct Base DN. For more information, see Specifying the Base Distinguished Name (Base DN) (page 174).

Configuring the Search Trusted Domains Option in Active Directory for a Qumulo Cluster

This section explains how to restrict the scope of LDAP queries by using the Search Trusted Domains configuration option for a Qumulo cluster joined to an Active Directory (AD) domain.

During normal AD domain operations, a Qumulo cluster often encounters *LDAP referrals* that indicate to the cluster in what other locations within an AD domain it might locate requested information. Often, these referrals are hints to other trusted AD domains which a cluster accesses through a Domain Trust, such as a Parent Domain Trust or an external Domain Trust.

Reducing Latency by Disabling Search Trusted Domains

In Qumulo Core 6.1.0.1 (and lower), to permit Qumulo clusters to follow LDAP referrals, the Search Trusted Domains configuration option is enabled by default.

In Qumulo Core 6.1.1 (and higher), to reduce the potential latency of AD domain operations that might trigger and follow LDAP referrals unnecessarily (particularly for large, complex AD environments with multiple Domain Trusts), you can disable the **Search Trusted Domains** configuration option.

Disabling this option might benefit your system if you can determine that all relevant user and group accounts—which you might expect to use POSIX attributes, logins with SAML Single Sign-On (SSO), or logins with NFS4.1 and Kerberos—are located entirely in the current domain.

Limitations of Disabling Search Trusted Domains

This section explains the limitations of disabling the Search Trusted Domains configuration option.

Trusted Domains Specified in the Base DN

The Base DN (Distinguished Name) configuration option specifies the path that limits LDAP queries. When you set the Base DN to the top-level domain or base path of a domain, LDAP searches span the entire domain's LDAP structure, including LDAP referrals to other domains that have a Trust with the currently joined domain.

Often, the Base DN configuration ensures that the system searches all Organizational Units (OUs) in the domain, for example when the Administrator team might not have control over the OUs that contain the user accounts to be retrieved. (This is common in a dynamic environment that an external team manages.)

Qumulo Core lets you configure multiple Base DNs by providing their paths in a semicolon-separated list that includes the paths of other trusted domains. This configuration permits the trusted domains to use POSIX attributes and SAML SSO logins.

A Important

Disabling Search Trusted Domains disregards any trusted domains specified in the Base DN.

Ignoring LDAP Referrals and Qumulo Core Authentication Processes

To decide whether your system should ignore LDAP referrals, consider the Qumulo Core authentication processes that this might affect.

Authentication Processes that Trigger LDAP Queries

- Identity mapping from NTFS to POSIX (SMB to NFS) by using the Use Active Directory for POSIX attributes AD configuration option
- SAML single sign-on (SSO) (page 31)
- NFSv4.1 and Kerberos (page 0)
- REST API access tokens (page 41)
- S3 access keys (page 209)

Unaffected Authentication Processes

- Kerberos SMB SSO logins from Domain Local or Trusted Domain users
- NTMLv2 SMB logins (username and password) from Domain Local or Trusted Domain users
- · Domain Local groups that contain users and groups from other Trusted Domains
- Users or groups added to SMB share permissions by using the Qumulo Core Web UI or qq
 CLI
- Security Identifiers (SIDs) resolved to usernames by using client dialog boxes, for example in macOS Finder or Windows File Explorer

External Services

Using Qumulo Core Access Tokens

This section explains how to create and use access tokens—by using the Qumulo REST API, Python SDK, and qq CLI—to authenticate external services to Qumulo Core.

☑ Tip

It is possible to confuse the terms *access token* and *session token*. Unlike access tokens, session tokens are short-lived and require a password to refresh, for example, to authenticate by using the qq_login command. Access tokens are the focus of this section.

In Qumulo Core 5.3.0 (and higher), you can use *access tokens* to let a user authenticate to the Qumulo REST API without having to complete repetitive login procedures.

Access tokens are long-lived. They provide an alternative to session-based authentication that the qq login command and the Web UI use. They also support support authentication for services, long-lived automation processes, and programmatic REST API access that doesn't require user input.

▲ Important

- An attacker can use an access token to authenticate as the token's user to Qumulo Core REST API (through HTTP, the Python SDK, or the qq CLI) and gain all of the user's privileges. Treat access tokens, and the bearer tokens they generate, like passwords. Store your tokens securely, rotate your tokens often, and create a token revocation policy for your organization.
- Because a token allows indefinite authentication to the associated user's account, we strongly recommend against creating tokens for individual Qumulo Core REST API users. For more information, see Best Practices for Using Access Tokens (page 0).

Prerequisites

- PRIVILEGE_ACCESS_TOKEN_WRITE is required for creating, disabling, and deleting access tokens for all users in the system.
- PRIVILEGE_ACCESS_TOKEN_READ is required for listing access tokens.

Creating and Using Access Tokens

PRIVILEGE_ACCESS_TOKEN_WRITE is required for creating, disabling, and deleting access tokens for all users in the system. This section explains how to create access tokens without or with an expiration time by using the qq CLI.

To Create an Access Token without an Expiration Time

Use the auth_create_access_token command and specify the user. For example:

```
$ qq auth_create_access_token jane
```

You can:

- · Specify the user as a name
- · Qualify the user by using a domain prefix, for example:
 - o ad:jane
 - ∘ AD\jane
 - ∘ local:jane
- · Specify ID types, for example:
 - o auth id:1234
 - SID:S-1-1-0

O Note

- · Although you can create groups for users, you can't create access tokens for groups.
- To use an access token in the qq CLI, you must use the --file flag—to specify a path for saving your credentials file in a format that the qq CLI can use—when you create the access token.

The auth_create_access_token command returns a JSON response that contains the bearer token body and the access token ID, which you can use to manage the access token.

```
{
    "bearer_token": "access-v1:abAcde...==",
    "id": "12345678901234567890123"
}
```

A Important

- As soon as you receive your bearer token, record it in a safe place. If you misplace the bearer token, you can't retrieve it at a later time. You must create a new access token.
- Any user can have a maximum of two access tokens. If a user already has two access tokens, creating new tokens fails until you remove at least one token from the user.
 We strongly recommend creating a single access token for each user and using the second access token to perform secret rotation.
- Treat access tokens, and the bearer tokens they generate, like passwords. Store your tokens securely, rotate your tokens often, and create a token revocation policy for your organization.
- To decrease the risk of giving an attacker full administrative access—including access to cluster data—avoid generating tokens for accounts with administrative privileges.

To Create an Access Token with an Expiration Time

In Qumulo Core 5.3.2 (and higher), you can use the auth_create_access_token --expiration-time command and specify the expiration time. You can specify the expiration time in different formats. For example:

```
$ qq auth_create_access_token jane --expiration-time 'Jan 01 2023'
```

```
$ qq auth_create_access_token jane --expiration-time '01/01/2023 00:00'
```

When an access token's expiration time elapses, it isn't possible to use the token for authentication. Any attempt to use the token results in an authentication error. To continue the authentication process, you must either create a new access token (page 0) or update the expiration time for your existing token (page 0).

O Note

The --expiration-time flag interprets arguments as timestamps in the UTC time zone.

Using Bearer Tokens for Authentication

A Qumulo Core access token returns a *bearer token* (page 42), an item in the Authorization HTTP header which acts as the authentication mechanism for the Qumulo Core REST API.

REST API

When you use the Qumulo REST API, add the bearer token to the **Authorization** HTTP header. For example:

```
Authorization: Bearer access-v1:abAcde...==
```

You can also add the bearer token to a curl command. For example:

```
$ curl https://203.0.113.0:8000/v1/session/who-am-i -H 'Authorization: Bearer acces
s-v1:abAcde...=='
```

Python SDK

When you use the Qumulo Python SDK, add the bearer token to a **RestClient** object. For example:

```
from qumulo.rest_client import RestClient
from qumulo.lib.auth import Credentials
client = RestClient('203.0.113.0', 8000, Credentials('access-v1:abAcde...=='))
```

For more information, see the Qumulo Core Python SDK.

qq CLI

To use an access token in the qq CLI, you must use the --file flag—to specify a path for saving your credentials file in a format that the qq CLI can use—when you create the access token. For example:

```
$ qq auth_create_access_token jane --file ./qumulo_credentials
```

To use the credentials file, specify its location by using the **--credentials-store** flag. For example:

```
$ qq --credentials-store ./qumulo_credentials who_am_i
```

Getting Metadata for Access Tokens

PRIVILEGE_ACCESS_TOKEN_READ is required for listing access tokens. This section explains how to get metadata for a specific access token or all access tokens by using the qq CLI.

To Get Metadata for a Specific Access Token

Use the auth_get_access_token command and specify the access token ID. For example:

```
$ qq auth_get_access_token 1234567890123456789012
```

This command returns a JSON object that lists:

- · The access token ID
- · The user that the access token represents
- · The access token's creator
- · The access token's creation time
- · The access token's expiration time
- · Whether the access token is enabled

For example:

```
"creation_time": "2022-12-06T01:14:39.56621474Z",
  "creator": {
    "auth id": "500",
    "domain": "LOCAL",
    "gid": null,
    "name": "admin",
    "sid": "S-1-1-12-12345678-1234567890-1234567890-500",
    "uid": null
 },
  "enabled": true,
  "expiration time": "2023-01-01T00:00:00Z",
  "id": "12345678901234567890123",
  "user": {
    "auth id": "1002",
    "domain": "LOCAL",
    "gid": null,
    "name": "svc",
    "sid": "S-1-1-12-12345678-1234567890-1234567890-1002",
    "uid": null
 }
}
```

To Get Metadata for All Access Tokens

Use the qq auth_list_access_tokens command.

A Important

Listing access tokens *doesn't* return the bearer token required for authentication. If you misplace the bearer token, you can't retrieve it at a later time. You must create a new access token.

The auth_list_access_tokens command returns:

- · The access token ID
- · The user that the access token represents
- · The access token's creator
- · The access token's creation time
- · The access token's expiration time
- · Whether the access token is enabled

For example:

id	user	creator	creation time
	=====	======	
1234567890123456789012	SVC	admin	2022-10-27T15:18:09.725513764Z
0987654321098765432109) svc	admin	2022-10-27T15:18:24.997572918Z
expiration time	enabled		
=======================================	======		
	True		
2023-01-01T00:00:00Z	False		

To filter the command's output by user, use the --user flag and use the same format for the name as for the auth_create_access_token (page 42) command.

Modifying the Expiration Time for an Access Token

PRIVILEGE_ACCESS_TOKEN_WRITE is required for creating, disabling, and deleting access tokens for all users in the system. This section explains how to modify access tokens by using the qq CLI.

Use the auth_modify_access_token command and specify the access token ID and the expiration time. For example:

```
$ qq auth_modify_access_token 1234567890123456789012 --expiration-time 'Jan 01 2023'
```

When an access token's expiration time elapses, it isn't possible to use the token for authentication. Any attempt to use the token results in an authentication error. To continue the authentication process, you must either create a new access token (page 0) or update the expiration time for your existing token (page 0).

Note

The --expiration-time flag interprets arguments as timestamps in the UTC time zone.

Disabling an Access Token

To help you check your system's security posture, Qumulo Core lets you disable an access token without deleting it. This is a good way to check for dependencies on the access token before you delete the token permanently.

PRIVILEGE_ACCESS_TOKEN_WRITE is required for creating, disabling, and deleting access tokens for all users in the system. This section explains how to disable an access token by using the qq CLI.

A Important

After you disable an access token, you can no longer use any bearer tokens associated with the access token to authenticate to Qumulo Core.

To disable an access token, use the auth_modify_access_token command, specify the access token ID, and use the -d flag. For example:

\$ qq auth modify access token 1234567890123456789012 -d

To enable an access token, use the auth_modify_access_token command, specify the access token ID, and use the -e flag. For example:

\$ qq auth_modify_access_token 1234567890123456789012 -e

Deleting Access Tokens

PRIVILEGE_ACCESS_TOKEN_WRITE is required for creating, disabling, and deleting access tokens for all users in the system. This section explains how to delete an access token by using the qq CLI.

A Important

After you delete an access token, you can no longer use any bearer tokens associated with the access token to authenticate to Qumulo Core.

To delete an access token, use the auth_delete_access_token command and specify the access token ID. For example:

\$ qq auth_delete_access_token 1234567890123456789012

Best Practices for Using Qumulo Core Access Tokens

This section lists the best practices for limiting the exposure to lost credentials and working with Qumulo Core access tokens securely.

Avoiding Creation of Tokens for Administrative Accounts

An attacker can use an access token to authenticate as the token's user to Qumulo Core REST API (through HTTP, the Python SDK, or the qq CLI) and gain all of the user's privileges. To decrease the risk of giving an attacker full administrative access—including access to cluster data—avoid generating tokens for accounts with administrative privileges.

Generating Tokens for Service Accounts

When you connect external services to the Qumulo Core REST API, we recommend creating a service account with limited privileges for each individual service and generating an access token for each service account.

To Create a New Service Account

- 1. Log in to Qumulo Core.
- 2. Create a service account.
 - a. Click Cluster > Local Users & Groups.
 - b. In the Users section, click Create.
 - c. In the Create user dialog box, enter a User name and Password, re-enter the password, and then click Create.
- 3. Create a role with privileges.
 - Click Cluster > Role Management.
 - b. In the Role Management section, click Create Role.
 - c. On the Create Role page, enter a Name and Description, click the Privileges for the user, and then click Save.
- 4. Assign the service user to the role.
 - a. On the Role Management page, find the name of the role you created and then click Add Member.
 - b. In the Add Member to <MyRoleName> dialog box, for Trustee, enter the name of the user you created and then click Yes, Add Member.

5. Create access tokens (page 0) for your service account.

Rotating Access Tokens

We strongly recommend rotating access tokens for a service account at a regular interval.

To Rotate an Access Token for a Service Account

- 1. To ensure that there is only one access token for each service account, use the qq auth list access tokens command.
 - If multiple access tokens exist, delete any unused access tokens.
- 2. To create a new access token for the service account, use the qq auth create access token command.
- 3. In the credential store of your service, replace the old access token with the new one.
- 4. Test that your service account can access the Qumulo Core REST API.
- 5. Confirm that there is nothing else relying on the old access token by disabling it first. If this causes any disruptions then you can re-enable it while you resolve the issue.
- 6. To delete the old access token, use the qq auth_delete_access_token command.

Connecting Your Kubernetes Cluster to Your Qumulo Cluster by Using the Qumulo Container Storage Interface (CSI) Driver

This section introduces the Qumulo Container Storage Interface (CSI) driver and explains how you can connect your Kubernetes cluster to your Qumulo cluster by using the Qumulo CSI driver.

To automate container storage, enable dynamic volumes, and help you scale your application container images based on usage and workflows, Qumulo uses its CSI driver to connect the Kubernetes orchestrator to Qumulo persistent storage. (In comparison, for example, the NFS CSI Driver for Kubernetes requires unprivileged NFS access for dynamic volumes and doesn't support volume sizing and expansion.)

For general driver information, see the Container Storage Interface (CSI) Specification.

Supported Features

The Qumulo CSI Driver supports:

- Static and dynamic (expansion) provisioning over NFSv3
- The following Persistent Volume Claim access modes:
 - ReadOnlyMany
 - ReadWriteMany
 - ∘ ReadWriteOnce
 - ReadWriteOncePod
- NFSv4.1

A Important

Even when you enable NFSv4.1 for your Qumulo cluster, you must explicitly configure NFSv4.1 to work with Kerberos (page 0).

Unsupported Features

- Volume cloning
- · Volume snapshot and restore

Requirements

- · A Qumulo cluster
- · Kubernetes 1.19 (and higher)

Connecting Your Qumulo Cluster to Kubernetes

This section explains how you can configure, provision, and mount Qumulo storage for each *Pod* (a logical wrapper for a container) on Kubernetes by using dynamic provisioning. This gives you more control over persistent volume capacity.

Step 1: Install the Qumulo CSI Driver

- 1. Log in to a machine that has kubectl and can access your Kubernetes cluster.
- 2. Download the .zip file or use one of the following commands.
 - · S3:

```
aws s3 cp s3://csi-driver-qumulo/deploy_v1.1.0.zip ./
```

· HTTP:

```
wget https://csi-driver-qumulo.s3.us-west-2.amazonaws.com/deploy_v
1.1.0.zip
```

- 3. Extract the contents of the .zip file.
- 4. Run the shell script and specify the current release version. For example:
 - · Linux:

```
cd deploy_v1.1.0
chmod +x install_driver.sh
./install-driver.sh
```

· Windows:

```
cd deploy_v1.1.0
install-driver.bat
```

The script configures Qumulo's prebuilt Elastic Container Registry (ECR) image (from public.ecr.aws/qumulo/csi-driver-qumulo:v1.1.0) and installs it on your Kubernetes system.

Step 2: Configure Volume and NFS Export Paths

To prepare your Qumulo cluster for connecting to your Kubernetes cluster, you must first configure your volume and NFS export paths on your Qumulo cluster by setting the following parameters for each storage class that you define.

☑ Tip

Write down the paths for the following YAML keys for the storageclass-qumulo.yaml file that you use when you create a storage class (page 0).

1. For storeRealPath, from the root of the Qumulo file system, create a directory for storing volumes on your Qumulo cluster, for example /csi/volumes1.

O Note

Because the CSI driver doesn't create the directory listed in the storeRealPath key automatically, this directory must exist above the NFS export and must not be the NFS export itself.

- 2. For storeExportPath, create the NFS export for hosting the persistent volume.
- 3. When you configure your cluster for multitenancy and the cluster's default tenant doesn't manage the NFS export, the optional tenantId parameter directs the driver to use the provided tenant. For more information, see Configure Multi-Tenancy with Qumulo on Qumulo Care.

O Note

You must provide the value for tenantId as a string. For example: "2".

Step 3: Configure Credentials

To connect your Kubernetes cluster to your Qumulo cluster, you must either use an existing account or create a new account for the CSI driver to communicate with the Qumulo API.

- 1. Configure a username and password for a user on your Qumulo cluster.
- 2. The configured username must have the following permissions:
 - Lookup on storeRealPath
 - · Create directories in storeRealPath
 - · Create and modify quotas:
 - PRIVILEGE_QUOTA_READ
 - PRIVILEGE QUOTA WRITE

- Read NFS exports: PRIVILEGE NFS EXPORT READ
- Perform TreeDelete operations on volume directories:
 PRIVILEGE_FS_DELETE_TREE_WRITE

For more information, see Role-Based Access Control (RBAC) with Qumulo Core on Qumulo Care.

Step 4: Create and Configure Secrets

To allow the CSI driver to operate with your Qumulo cluster, you must create and configure Secrets. You may use either Basic Authentication with a username and password, or an Access Token. Depending on configuration, Basic Authentication may be disallowed and using an Access Token will be required.

- 1. Configure one of the following authentication types.
 - · Basic Authentication:

```
kubectl create secret generic cluster1-login \
    --type="kubernetes.io/basic-auth" \
    --from-literal=username=bill \
    --from-literal=password=SuperSecret \
    --namespace=kube-system
```

· Access Token:

```
TOKEN='access-v1:zNTc5D0zWTdNi/KsZo620fu71TweGh47u+S/5NbV...'
kubectl create secret generic cluster1-login \
    --from-literal=access_token="$TOKEN" \
    --namespace=kube-system
```

2. Give the CSI driver access to the Secrets. For example:

```
kubectl create role access-secrets \
    --verb=get,list,watch \
    --resource=secrets \
    --namespace kube-system
kubectl create rolebinding \
    --role=access-secrets default-to-secrets \
    --serviceaccount=kube-system:csi-qumulo-controller-sa \
    --namespace kube-system
```

Step 5: Create a Storage Class

To link your Kubernetes cluster to your Qumulo cluster, you must create a storage class on your Kubernetes cluster.

1. Begin with the example Qumulo storage class configuration.

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
 name: cluster1
provisioner: qumulo.csi.k8s.io
parameters:
  server: 203.0.113.0
  storeRealPath: "/regions/4234/volumes"
  storeExportPath: "/some/export"
  csi.storage.k8s.io/provisioner-secret-name: cluster1-login
  csi.storage.k8s.io/provisioner-secret-namespace: kube-system
 csi.storage.k8s.io/controller-expand-secret-name: cluster1-login
  csi.storage.k8s.io/controller-expand-secret-namespace: kube-system
reclaimPolicy: Delete
volumeBindingMode: Immediate
mountOptions:
  - nolock
  - intr
  - proto=tcp
allowVolumeExpansion: true
```

- 2. Edit the configuration for your Qumulo cluster.
 - a. Name your storage class.
 - b. Specify server and storeRealPath.
 - c. Specify storeExportPath.
 - d. (Optional) Specify tenantId.
 - e. Configure the following parameters to point to the Secrets that you have created and configured (page 0) in the namespace in which you installed the CSI driver:
 - controller-expand-secret-name
 - controller-expand-secret-namespace
 - provisioner-secret-name
 - provisioner-secret-namespace

f. Specify the NFS mountOptions. For example:

```
mountOptions:
- nolock
- intr
- proto=tcp
- vers=3
```

g. To create the class, apply the configuration. For example:

```
kubectl create -f storageclass-qumulo.yaml
```

Step 6: Create a Persistent Volume Claim (PVC) and Apply it to a Pod

To apply a PVC claim to a Pod dynamically, you must first configure and create it.

1. Begin with the example PVC configuration.

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
    name: claim1
spec:
    accessModes:
    - ReadWriteOnce
    storageClassName: cluster1
    resources:
        requests:
        storage: 1Gi
```

- 2. Edit the configuration for your PVC claim.
 - a. Name your claim.
 - b. Change storeClassName to the name of your claim.
 - c. Specify the capacity in spec.resources.requests.storage. This parameter lets you create a quota on your Qumulo cluster.
 - d. To create the claim, apply the configuration. For example:

```
kubectl apply -f dynamic-pvc.yaml
```

3. Use the claim in a Pod or a Deployment. For example:

▲ Important

When the PVC is released, a tree-delete is initiated on the Qumulo cluster for the directory that the PVC indicates. To prevent this behavior, set reclaimPolicy to Retain.

4. You can launch and use your container image.

Network Configuration

Required Networking Ports for Qumulo Core

This section explains which inbound and outbound networking ports Qumulo Core requires.

O Note

- · Communication between nodes in a cluster must be possible on any port.
- Active Directory authentication services require their own network port range. For an authoritative list, see Active Directory and Active Directory Domain Service Port Requirements in the Windows Server 2008 R2 and Windows Server 2008 documentation.

Networking Ports for Inbound Connections

Port	Protocols	Use	
21	TCP	FTP	
22	TCP	SSH	
80	TCP	HTTP (Web UI)	
111	TCP UDP	rpcbind or portmapper for NFSv3	
443	TCP	HTTPS (Web UI)	
445	ТСР	SMB	
2049	TCP UDP	NFS or MOUNT	
		Qumulo Core supports UDP for the MOUNT protocol for older clients. However, any NFS clients—that specify the TCP mount option or transfer data over NFS after mounting—don't use UDP.	

Port	Protocols	Use	
3712	TCP	Replication	
8000	TCP	REST API	
9000	TCP	S3 API, if you enable the S3 API for your Qumulo cluster (page 204)	
32768-60999	TCP	FTP Passive Mode	

Networking Ports for Outbound Connections

O Note

For cluster formation and inter-node communication Qumulo Core requires the following:

- Hardware Platforms: Unblocked IPv6 traffic in the local subnet—for more information, see Configuring IPv6 in Qumulo Core on Qumulo Care.
- · Cloud Platforms: Unblocked IPv4 traffic in the local subnet

Port	Protocols	Use	
53	UDP	DNS	
88	TCP	Kerberos	
111	III TCP	rpcbind or portmapper for NSM and NLM	
	① Note Depending on the client portmapper configuration, Qumulo Core might require additional ports.		
123	UDP	Synchronization of product and network time, for authentication and time-stamping of artifacts such as audit logs, by using the Network Time Protocol (NTP).	
135	TCP	DCERPC or Netlogon (Domain Controller Binding)	
389, 636	TCP	LDAP to Active Directory or to a standalone LDAP server (by default)	

Port	Protocols	Use
443	ТСР	Qumulo Shift for Amazon S3 (by default)
514	ТСР	Audit with Rsyslog (by default)
3712	TCP	Replication (by default)

Network Multitenancy

Partitioning a Qumulo Cluster into Tenants

This section explains how to enable, disable, and use network multitenancy in Qumulo Core.

In Qumulo Core 5.3.4 (and higher), *network multitenancy* lets you partition a single physical Qumulo cluster into multiple virtual *tenants*. You can define a tenant by using a name and one or more networks. When you assign a network to a tenant, Qumulo Core treats any client that connects from that network as part of that tenant. For more information, see Connect to Multiple Networks in Qumulo Core on Qumulo Care.

For each tenant, you can specify individual management protocol access and configuration (page 65) and file system protocol access and configuration (page 68).

O Note

All tenants share the cluster's underlying file system, identity providers, role-based access control (RBAC) configuration, and other global settings.

Prerequisites

A Important

If your cluster runs a version of Qumulo Core lower than 6.1.0.1, you must use the qq multitenancy enable command to enable multitenancy for your cluster.

To manage network multitenancy and tenants, your user must have membership in a Qumulo role with the following privileges.

- PRIVILEGE NETWORK READ: Viewing networks
- PRIVILEGE_NETWORK_WRITE: Assigning networks to tenants
- PRIVILEGE_TENANT_READ: Viewing tenants
- PRIVILEGE_TENANT_WRITE: Enabling network multitenancy and creating, modifying, and deleting tenants

Creating, Configuring, and Unassigning Tenants by Using the qq CLI.

This section explains the lifecycle of working with tenants, including creating tenants, assigning networks to new and existing tenants, viewing tenant configuration, and unassigning tenants.

Step 1: Create a New Tenant

After you enable multitenancy for your cluster, you can create a new tenant.

Use the qq multitenancy_create_tenant command and specify a name for your tenant. For example:

```
$ qq multitenancy_create_tenant \
--name my_tenant
```

Your cluster creates a new tenant with no networks assigned to it. By default, Qumulo Core disables access to the tenant through all protocols.

Step 2: Assign Networks to a Tenant

To allow a tenant to apply its configuration to clients that connect to the cluster from specific networks, you must associate the tenant with one or more networks. To do this, you can:

- · Create a new tenant with networks assigned to it.
- · Assign networks to, and unassign networks from, an existing tenant.
- · Move networks between tenants.

Creating a New Tenant with Assigned Networks

Use the qq multitenancy create tenant and specify the tenant name and network ID.

▲ Important

The network must not be assigned to another tenant already.

In the following example, we specify a single network.

```
$ qq multitenancy_create_tenant \
    --name my_tenant_name \
    --network-id 2
```

Assigning a Single Network to an Existing Tenant

To assign a single network to a tenant, modify the tenant that belongs to the network. Use the qq network mod network command and specify the network and tenant ID.

```
$ qq network_mod_network \
   --network-id 3 \
   --tenant-id 2
```

Assigning Multiple Networks to an Existing Tenant

To assign multiple networks to a tenant, modify the networks that belong to the tenant. Use the qq multitenancy modify tenant command and specify the tenant and network ID.

A Important

Any existing networks not specified after the --network-id flag become unassigned.

In the following example, we specify three networks.

```
$ qq multitenancy_modify_tenant \
   --id 2 \
   --network-id 2 3 4
```

Unassigning a Single Network from a Tenant

To unassign a single network from a tenant, clear *the tenant that belongs to the network*. Use the qq_network_mod_network command and specify the network and the --clear-tenant-id flag.

O Note

After you unassign a network from a tenant, you can assign it to another tenant.

```
$ qq network_mod_network \
   --network-id 3 \
   --clear-tenant-id
```

Unassigning All Networks from a Tenant

To unassign all networks from a tenant, clear the networks that belong to the tenant. Use the qq multitenancy modify tenant command and specify the tenant and network ID.

Note

- · After you unassign a network from a tenant, you can assign it to another tenant.
- · Don't specify any arguments for the --network-id flag.

```
$ qq multitenancy_modify_tenant \
    --id 2 \
    --network-id
```

Moving Networks between Tenants

Use the qq network_mod_network command and specify the network and target tenant.

```
$ qq network_mod_network \
   --network-id 2 \
   --tenant-id 1
```

Step 3: View Tenant Information

To determine a tenant's network assignments and enabled management and file system protocols, you can view the tenant information.

Viewing Information for a Single Tenant

To view the information for a single tenant, use the qq multitenancy_get_tenants command.

```
$ qq multitenancy_get_tenant \
--id 1
```

Viewing Information for All Tenants

- In the Web UI, log in to Qumulo Core and then click Cluster > Network Multitenancy.
- · In the qq CLI, use the multitenancy list tenants command.

Determining the Tenant Assignment for Networks

• To view the information for a single network, use the qq network get network command.

```
$ qq network_get_network \
   --network-id 2
```

• To view the information for all networks, use the qq network list networks command.

Step 4: Delete a Tenant

A Important

- When you delete a tenant, Qumulo Core removes the tenant's entire configuration from your cluster, including NFS exports and SMB shares associated with the tenant.
- · It isn't possible to delete the last tenant.

To delete a tenant, use the qq multitenancy_delete_tenant command and specify the tenant ID.

```
$ qq multitenancy_delete_tenant \
--id 2
```

Known Network Multitenancy Limitations in Qumulo Core

Currently, Qumulo Core doesn't support:

- · Creating or modifying tenants on cloud-based clusters
- · Using one VLAN on multiple tenants
- Using a separate Active Directory, standalone LDAP, or user-defined mapping configuration for each tenant
- · Using a separate DNS configuration for each tenant
- · Scoping RBAC privileges to each tenant

Configuring Management Protocols on a Tenant

This section explains how to configure management protocols for each tenant on a Qumulo cluster.

After you create tenants (page 60) on your Qumulo cluster, you can manage access for clients that connect to the cluster from the tenant's network by enabling or disabling the management protocols for each tenant.

A Important

- Access to a management protocol lets a client use the protocol to view and modify resources across the entire cluster, not only within the client's tenant.
- If you disable a management protocol, you can still access your cluster by using a physical or remote console.

Prerequisites

To configure management protocols, your user must have membership in a Qumulo role with the following privileges.

- PRIVILEGE TENANT READ: Viewing tenants
- PRIVILEGE TENANT WRITE: Modifying tenants

Enabling and Disabling REST API Access

The Qumulo REST API lets you manage clusters by using the qq CLI, Python bindings, and REST API calls.

- To enable REST API access, use the qq multitenancy_modify_tenant command, specify
 the tenant ID, and use the --enable-rest-api flag.
- To disable access, use the --disable-rest-api flag.

In the following example, the cluster begins to accept REST API traffic on TCP port 8000 on all networks assigned to the tenant.

```
$ qq multitenancy_modify_tenant \
   --id 2
   --enable-rest-api
```

Enabling and Disabling Web UI Access

The Qumulo Core Web UI lets you manage clusters from a browser.

O Note

Enabling Web UI access doesn't require enabling REST API access.

- To enable REST API access, use the qq multitenancy_modify_tenant command, specify the tenant ID, and use the --enable-web-ui flag.
- · To disable access, use the --disable-web-ui flag.

In the following example, the cluster begins to serve the Web UI on TCP ports 80 and 433 on all networks assigned to the tenant.

```
$ qq multitenancy_modify_tenant \
   --id 2
   --enable-web-ui
```

Enabling and Disabling SSH Access

SSH lets you view logs and use the qq CLI by using a client to connect to nodes in a cluster remotely.

O Note

Enabling SSH access doesn't require enabling REST API access.

- To enable SSH access, use the qq multitenancy_modify_tenant command, specify the tenant ID, and use the --enable-ssh flag.
- To disable access, use the --disable-ssh flag.

In the following example, the cluster begins to accept SSH traffic on TCP port 22 on all networks assigned to the tenant.

```
$ qq multitenancy_modify_tenant \
   --id 2
   --enable-ssh
```

Enabling and Disabling Replication Access

Replication lets a cluster transfer data from a directory on one cluster to a directory on another cluster.

- To enable replication access, use the qq multitenancy_modify_tenant command, specify
 the tenant ID, and use the --enable-replication flag.
- · To disable access, use the --disable-replication flag.

In the following example, the cluster begins to accept replication traffic on TCP port 3712 on all networks assigned to the tenant.

```
$ qq multitenancy_modify_tenant \
   --id 2 \
   --enable-replication
```

Configuring File System Protocols on a Tenant

This section describes how to configure file system protocols for each tenant on a Qumulo cluster.

After you create tenants (page 60) on your Qumulo cluster, you can manage access for clients that connect to the cluster from the tenant's network by enabling or disabling file system protocols, or configuring file system protocols specifically for each tenant.

Prerequisites

To configure file system protocols, your user must have membership in a Qumulo role with the following privileges.

- PRIVILEGE NFS EXPORT READ: View NFS exports
- PRIVILEGE_NFS_EXPORT_WRITE: Create, modify, and delete NFS exports
- PRIVILEGE_NFS_SETTINGS_READ: View NFS settings
- PRIVILEGE NFS SETTINGS WRITE: Modify NFS settings
- PRIVILEGE SMB SHARE READ: View SMB shares and settings
- PRIVILEGE SMB EXPORT WRITE: Create, modify, and delete SMB shares; modify SMB settings
- PRIVILEGE_TENANT_READ: View tenants
- PRIVILEGE_TENANT_WRITE: Modify tenants

Enabling NFS and Configuring Settings and Exports

This section explains how to enable the NFS protocol for a tenant, the difference between global settings and settings for each tenant, and how to configure NFS exports for a tenant.

Enabling the NFS Protocol for a Tenant

To let a tenant use NFS to access a specific set of exports, use the qq
multitenancy_modify_tenant command, specify the tenant ID, and use the --enable-nfs flag.



Initially, the tenant has no exports to mount.

In the following example, the cluster begins to accept NFS traffic on TCP and UDP ports 111 and 2019, on all networks assigned to the tenant.

```
$ qq multitenancy_modify_tenant \
    --id 2 \
    --enable-nfs
```

Configuring the NFS Protocol Globally or for a Tenant

You can use *global settings* that apply to all tenants. For example, to enable NFSv4.1 for all tenants, use the **qq nfs_modify_settings --enable-v4** command. For more information about NFS settings see Enabling and Using NFSv4.1 on a Qumulo Cluster (page 157) and Using NFSv4.1 with Kerberos in Qumulo Core (page 0).

To override this global setting, use the qq multitenancy_nfs_modify_settings command to disable NFSv4.1 on a specific tenant, specify the tenant ID, and use the --disable-v4 flag. In the following example, despite the global setting, the specified tenant no longer accepts NFSv4.1 traffic on port 2049, on all networks assigned to the tenant.

```
$ qq multitenancy_nfs_modify_settings \
   --tenant-id 2 \
   --disable-v4
```

To revert to the global setting for a specific tenant, use the qq multitenancy_nfs_delete_settings command and specify the tenant ID. For example:

```
$ qq multitenancy_nfs_delete_settings \
   --tenant-id 2
```

Configuring NFS Exports

To isolate clients within a network in specific directories in the file system, you can configure NFS exports for each tenant.

O Note

Because each new tenant has no exports initially, to give clients within that tenant's network access to the cluster over NFS, you must add exports. For more information, see Create an NFS Export on Qumulo Care.

To create an NFS export for a tenant, use the qq nfs_add_export command and specify the
tenant ID, the export path, and the file system path and use the --no-restrictions flag. In the
following example, a cluster lets clients access directory /my-directory from the export /myexport when the clients are in the network assigned to the tenant.

```
$ qq nfs_add_export \
   --tenant-id 2 \
   --export-path /my-export \
   --fs-path /my-directory \
   --no-restrictions
```

▲ Important

When multiple tenants exist on a cluster, you must specify the tenant ID.

To move an NFS export from one tenant to another, use the qq nfs_mod_export command. In the following example, while the cluster denies access to clients in tenant 2, it lets clients in tenant 1 access the export /my-export.

```
$ qq nfs_mod_export \
   --tenant-id 2 \
   --export-path /my-export \
   --new-tenant-id 1
```

Enabling SMB and Configuring Settings and Shares

This section explains how to enable the SMB protocol for a tenant, the difference between global settings and settings for each tenant, and how to configure SMB shares for a tenant.

Enabling the SMB Protocol for a Tenant

To let a tenant use SMB to access a specific set of shares, use the qq multitenancy_modify_tenant command, specify the tenant ID, and use the --enable-smb flag.

O Note

Initially, the tenant has no shares to mount.

In the following example, the cluster begins to accept SMB traffic on TCP port 445, on all networks assigned to the tenant.

```
$ qq multitenancy_modify_tenant \
   --id 2 \
   --enable-smb
```

Configuring the SMB Protocol Globally or for a Tenant

You can use *global settings* that apply to all tenants. For example, to require encryption for all tenants, use the **qq smb_modify_settings** --encryption-mode required command. For more information about SMB settings, see the following resources in Qumulo Care:

- SMB3 Encryption with Qumulo Core
- SMB Host Restrictions
- · Hide an SMB Share
- Snapshots: Snapshot Directory Mode

To override this global setting, use the qq multitenancy_smb_modify_settings command to allow clients on only specific networks to require encryption. In the following example, despite the global setting, the specified tenant accepts unencrypted SMB requests on all networks assigned to the tenant.

```
$ qq multitenancy_smb_modify_settings \
   --tenant-id 2 \
   --encryption-mode none
```

To revert to the global setting for a specific tenant, use the qq multitenancy_smb_delete_settings command and specify the tenant ID. For example:

```
$ qq multitenancy_smb_delete_settings \
   --tenant-id 2
```

Configuring SMB Shares

To isolate clients within a network in specific directories in the file system, you can configure SMB shares for each tenant.

O Note

Because each new tenant has no shares initially, to give clients within that tenant's network access to the cluster over SMB, you must add shares. For more information, see Create an SMB Share on Qumulo Care.

To create an SMB share for a tenant, use the qq smb_add_share command and specify the tenant
ID, the share path, and the file system path and use the --all-access flag. In the following
example, a cluster lets clients access access directory /my-directory from the share /my-share
when the clients are in the network assigned to the tenant.

```
$ qq smb_add_share \
   --tenant-id 2 \
   --name /my-share \
   --fs-path /my-directory \
   --all-access
```

A Important

When multiple tenants exist on a cluster, you must specify the tenant ID.

To move an SMB share from one tenant to another, use the qq smb_mode_share command. In then following example, while the cluster denies access to clients in tenant 2, it lets clients in tenant 1 access the share /my-share.

```
$ qq smb_mod_share --tenant-id 2 \
   --name foo \
   --new-tenant-id 1
```

Using the Microsoft Management Console (MMC)

To create, modify, and delete SMB shares for each tenant, you can use the MMC Shared Folders snap-in. For more information, see Manage File Shares with Shared Folders MMC Snap-in on Qumulo Care.

▲ Important

Because folder paths always start at the root of the file shstem, using the MMC on a Windows client on a tenant's network causes the MMC to modify that tenant's shares.

Global Configuration of S3 and FTP Protocols

Qumulo Core doesn't permit configuring file system protocols such as S3 and FTP for each tenant. You must enable, disable, and configure these protocols globally.

▲ Important

Enabling these protocols makes them available on all networks.

For more information, see S3 API (page 0) in this guide and FTP in Qumulo Core on Qumulo Care.

Listing Client Connections

To determine tenant activity and client types connected to the tenant, use the qq network_list_connections command. For more information, see Balance of Client Connections on your Qumulo Cluster on Qumulo Care.

The following is example output from the command.

Web UI

Setting the Web UI Login Banner

This section explains how to set a login banner for the Qumulo Core Web UI.

In Qumulo Core 5.2.1 (and higher), clusters have an optional login banner that users must acknowledge before being they can log in to the Web UI.

To Set the Web UI Login Banner

To set the login banner, use the web_ui_modify_settings command. To specify the Markdown file to use for the banner, use the --login-banner flag. For example:

```
qq web_ui_modify_settings --login-banner my-banner.md
```

To Clear the Web UI Login Banner

To clear the login banner, use the web_ui_modify_settings command with the --disable-login-banner flag.

```
qq web_ui_modify_settings --disable-login-banner
```

To View the Current Web UI Login Banner

To view the current login banner, use the web_ui_get_settings command with the --login-banner flag.

```
qq web_ui_get_settings --login-banner
```

Setting the Web UI Inactivity Timeout

This section explains how to set an inactivity timeout for the Qumulo Core Web UI.

In Qumulo Core 5.1.0 (and higher), clusters have an optional *inactivity timeout* that logs users out of the Web UI if they don't interact with it for a specified amount of time.

1 Note

During the final minute of the timeout period, the Your Session is About to Expire dialog box appears. The dialog box shows a countdown and lets the user renew the session or log out immediately. When deciding on the timeout length, take your users' needs into consideration.

To Set the Web UI Inactivity Timeout

To set an inactivity timeout, use the web_ui_modify_settings command. Specify the timeout in minutes by using the --inactivity-timeout flag. For example:

qq web_ui_modify_settings --inactivity-timeout 15

To Clear the Web UI Inactivity Timeout

To clear an inactivity timeout, use the web_ui_modify_settings command with the --disableinactivity-timeout flag.

qq web_ui_modify_settings --disable-inactivity-timeout

To View the Current Web UI Inactivity Timeout

To view the current inactivity timeout, use the web_ui_get_settings command:

qq web_ui_get_settings

qq CLI

Enabling Autocomplete for the qq CLI

This section explains how to enable automatic command completion for the qq CLI and for command aliases.

The qq CLI supports Python argparse completion that helps you use the CLI more effectively. This section explains how to enable automatic command completion for the qq CLI and for command aliases.

A Important

The following procedures apply to running the qq CLI on Linux, macOS, and Windows Subsystem for Linux. Don't run these commands on Qumulo nodes

To Enable Autocomplete for the qq CLI

1. Install the argcomplete Python package.

pip install argcomplete

Note

Qumulo Core supports argcomplete 2.0.0 and higher.

2. Activate the argcomplete package.

sudo activate-global-python-argcomplete

3. Search for any conflicting qq entries.

complete | grep qq

If conflicting entries exist, remove them by specifying the entry name or path. For example:

complete -r /my/path

4. To enable autocompletion for the qq CLI, add the following line to the end of your shell profile (.bashrc, .bash_profile, and so on).

```
eval "$(register-python-argcomplete qq)"
```

5. Reload your shell profile.

```
source ~/.bashrc
```

You can now use the **Tab** key to autocomplete **qq** CLI commands. The **qq** CLI supports autocomplete for all CLI arguments and Qumulo REST API command arguments.

Enabling Autocomplete for qq CLI Command Aliases

To eliminate the need to repeatedly enter qq CLI flags (such as --host or --credentials-store), for example when dealing with multiple Qumulo clusters, you can add aliases for qq CLI commands to your shell profile. In the following example, we alias a complex qq CLI command to the simple alias qqcreds.

```
alias qqcreds='qq --host my.qumulo.com --credentials-store ~/.my_creds'
```

When you reload your profile, you can append a parameter to the complex command by appending it to the alias. For example:

```
qqcreds my_credentials
```

To ensure that your argcomplete configuration works with qq CLI command aliases, you must perform additional configuration and add a third-party helper script to your system.

▲ Important

Before you begin, review the source code of the complete-alias helper script. Qumulo doesn't contribute to, maintain, or take responsibility for this script.

To Enable Autocomplete for qq CLI Command Aliases

1. Add a qq CLI command alias and the COMPAL_AUTO_UNMASK configuration parameter to your shell profile (.bashrc, .bash profile, and so on). For example:

```
#qq CLI Autocomplete
eval "$(register-python-argcomplete qq)"
COMPAL_AUTO_UNMASK=1
source ~/.bash_completion.d/complete_alias
```

☑ Tip

Don't reload your shell profile yet.

2. Create a directory for the complete-alias daemon and download the script to it.

3. Add your alias to the complete_alias file.

```
echo "complete -F _complete_alias qqcreds" >> ~/.bash_completion.d/complete_al
ias
```

4. Search for any conflicting complete entries.

```
complete | grep complete
```

If conflicting entries exist, remove them by specifying the entry name or path. For example:

```
complete -r /my/path
```

5. Reload your shell profile.

```
source ~/.bashrc
```

You can now use the Tab key to autocomplete qq CLI command aliases.

Snapshots

How Snapshots Work in Qumulo Core

This section explains snapshots, their storage usage, and their locking functionality in Qumulo Core.

How Snapshots Work

Qumulo Core 2.5.0 (and higher) can take instant snapshots of the file system. A *snapshot* is an entry for every version of file system elements such as files, directories, creation and modification timestamps, permissions, and so on. Each new entry points only to changed data and, to allow original and new entries to share data, Qumulo Core writes the entries alongside each other.

Taking a snapshot doesn't consume storage or incur a performance penalty. There is only a negligible performance penalty for reading and writing snapshotted file system data.

How Snapshots Grow Over Time

The following example shows how Qumulo Core allocates storage to data and links it to file metadata as file system data changes.

In this scenario:

- 1. A user creates a file with 4 MB of data.
- 2. Qumulo Core takes a snapshot of the file.
- 3. A user modifies 1 MB of data within the file.
- 4. Qumulo Core allocates a new 1 MB region to the modified data.



Now, the following is true:

- 5 MB: The total storage that the file occupies
 - o 3 MB: Data shared between the original and new versions of the file
 - o 1 MB: Original data that exists only in the saved (snapshotted) version of the file
 - 1 MB: New data that exists only in the live (latest) version of the file

Next, the following conditions take effect:

- If the user rewrites that particular 1 MB of data, the system overwrites the existing live data without allocating new space.
- · If the user rewrites a different region of the file, Qumulo Core allocates additional storage.

Determining Snapshots' Storage Usage

When Qumulo Core tracks the difference between the *saved* (snapshotted) and *live* (latest) versions of a file, it creates a *lineage* of snapshots independent from each other. To determine the amount of data that a single snapshot references, use the qq

snapshot_get_capacity_used_per_snapshot command and specify the snapshot ID. For example:

```
qq snapshot_get_capacity_used_per_snapshot \
   --id 1682119059
```

More than one snapshot can reference *covered data*. It isn't possible to release covered data until you delete all *covering snapshots* that reference it.

- To determine the total covered data, including data no longer present in the snapshot, use the qq snapshot_get_capacity_used_per_snapshot command and specify multiple, comma-separated snapshot IDs.
- To determine the total amount of data, including covered data that multiple snapshots reference, use the qq snapshot_get_total_used_capacity command.

When you delete a snapshot, Qumulo Core removes the data which that snapshot references but retains the data which any other snapshot references. This ensures a full file representation within the remaining snapshots. Qumulo Core uses a background process to recover the storage that the snapshot had consumed.

O Note

When you delete a snapshot, the background process might take some time. To track the reclaimed storage, run the qq snapshot_get_total_used_capacity command.

Example: Tracking Covering Snapshots and Data Changes

For example, if you use the qq snapshot_get_total_used_capacity command, Qumulo Core shows that storage usage is 1,319,413,953,331 Bytes (1.2 TiB). This amount includes the total snapshot data and the covering snapshots.

If you add up the usage for all snapshots currently in the file system (by using the qq snapshot_get_capacity_used_per_snapshot command), Qumulo Core shows that total snapshot storage usage is 2,147,483,648 Bytes (2 GiB). This amount includes the data changes that each snapshot stores but doesn't include the unchanged file portions within each snapshot.

Example: Tracking File Snapshot Changes Over Time

For example, you have a 1 TiB file that you modify over time.

- Snapshot 1: This snapshot is 1,099,511,627,776 Bytes in size and contains the full 1 TiB file.
- · Snapshot 2: This snapshot is 1,073,741,824 Bytes in size and contains 1 GiB of data changes.
- Snapshot 3: This snapshot is 1,073,741,824 Bytes in size and contains an additional 1 GiB of data changes.

If you delete snapshot 1, only 1,023 GiB of data (covered by snapshots 2 and 3) remain. Qumulo Core doesn't release this 1,023 GiB of data until you delete all snapshots that reference the original file.



Without the data that snapshots 2 and 3 cover, no full file representation is possible.

Managing Snapshots in Qumulo Core

This section explains how to create on-demand snapshots and snapshot policies, view and search for existing snapshots, and delete snapshots by using the Web UI. It also explains how to create snapshots on a schedule, create a snapshot with an expiration time, and modify a snapshot's expiration time.

Managing Snapshots by Using the Web UI

This section explains how to create on-demand snapshots and snapshot policies, view and search for existing snapshots, and delete snapshots by using the Web UI.

To Create an On-Demand Snapshot

- 1. Log in to the Qumulo Core Web UI.
- 2. Click Cluster > Saved Snapshots.
- 3. On the Saved Snapshots page, in the upper right, click Take Snapshot.
- 4. In the On Demand Snapshot dialog box, do the following.
 - a. Enter the Snapshot Name.
 - b. For Apply to Directory, enter the directory to snapshot.
 - c. For Delete Snapshot, specify whether Qumulo Core should never delete the snapshot or delete it after a specified time period.
 - d. Click Save.

To Create a Snapshot Policy

- 1. Log in to the Qumulo Core Web UI.
- 2. Click Cluster > Policy.
- 3. On the Snapshot Policies page, in the upper right, click Create Policy.
- 4. On the Create Snapshot Policy page, do the following:
 - a. Enter the Policy Name.
 - b. For Apply to Directory, enter the directory to snapshot.
 - c. In the Run Policy on the Following Schedule section, specify the snapshot frequency and when to delete snapshots.
 - d. Click Enable policy upon creation.
 - e. Click Create Policy.

To View Existing Snapshots

The Snapshots page lets you navigate a large number of snapshots.

- 1. Log in to the Qumulo Core Web UI.
- 2. Click Cluster > Saved Snapshots.
- In Qumulo Core version 4.3.3 (and higher), if you have more than 50 snapshots, useto navigate the snapshot pages.

You can also use the controls at the bottom of the table to navigate to a specific page or change the number of rows for each page.

To Find a Specific Snapshot

In Qumulo Core version 4.3.3 (and higher), you can search for a specific snapshot by name, creation time, and so on.

- 1. Log in to the Qumulo Core Web UI.
- 2. Click Cluster > Saved Snapshots.
- 3. At the top of the table, click enable filters.

The Search... field appears.

4. Enter a search query.

The table rows match your query as you type.

5. (Optional) To turn off filtering, click disable filters.

To Delete a Single Snapshot

- 1. Log in to the Qumulo Core Web UI.
- 2. Click Cluster > Saved Snapshots.
- 3. On the right side of a snapshot's row, click

To Delete Multiple Snapshots

In Qumulo Core version 4.3.3 (and higher), you can delete multiple snapshots at once.

- 1. Log in to the Qumulo Core Web UI.
- 2. Click Cluster > Saved Snapshots.
- 3. On the left side of the table, select every snapshot to delete.

When you select more than one row, the Bulk Delete button appears.

4. When you finish selecting snapshots, click Bulk Delete.

O Note

Because all selection and deletion controls operate only on the current page, it isn't possible to delete a snapshot accidentally if it isn't listed on the current page.

Managing Snapshots by Using the qq CLI

This section explains how to create snapshots on a schedule, create a snapshot with an expiration time, and modify a snapshot's expiration time by using the qq CLI.

A Important

Creating and modifying snapshot policies with an associated lock requires the SNAPSHOT_LOCK permission in addition to policy permissions.

Creating Snapshots on a Schedule by Using a Snapshot Policy

Use the qq snapshot_create_policy command to create a snapshot policy and specify the interval at which Qumulo Core takes and deletes snapshots.

In the following example, we create a policy named every_day that takes a snapshot every midnight in the Pacific time zone and retains the snapshot for two days. Every new snapshot that this policy creates is locked with a key named my-key-name. For more information, see Locking and Unlocking Snapshots in Qumulo Core (page 87).

Note

The timezone flag uses values from the tz database. If you don't specify a time zone, the snapshot policy uses UTC time.

```
qq snapshot_create_policy daily \
    --name every_day \
    --days-of-week all \
    --at 00:00 \
    --timezone America/Los_Angeles \
    --time-to-live 7days
    --lock-key my-key-name
```

In the following example, we change a previously created policy with ID 1 to a policy named hourly that takes a snapshot every hour, but only during business hours (Monday to Friday, 8am to 6pm in the Pacific time zone), and retains snapshots for two days. Every new snapshot that this policy creates is unlocked (previously created snapshots remain locked). For more information, see Locking and Unlocking Snapshots in Qumulo Core (page 87).

```
qq snapshot_modify_policy change_to_hourly_or_less \
   -i 1 \
   --name hourly \
   --period lhours \
   --days-of-week MON,TUE,WED,THU,FRI \
   --start-time 08:00 \
   --end-time 18:00 \
   --timezone America/Los_Angeles \
   --time-to-live 2days
   --clear-lock-key
```

Creating an On-Demand Snapshot with an Expiration Time

Use the qq_snapshot_create_snapshot command to specify an expiration date or expiration time before Qumulo Core deletes the snapshot.

O Note

If you don't specify an expiration date or expiration time before deletion, Qumulo Core never deletes the snapshot.

In the following example, the snapshot expires on December 31, 2030, at midnight, in UTC time.

```
qq snapshot_create_snapshot \
--expiration 2030-12-31T00:00Z
```

In the following example, Qumulo Core deletes the snapshot in one year from the snapshot's creation time.

```
qq snapshot_create_snapshot \
   --time-to-live 12months
```

Modifying a Snapshot's Expiration Time

Use the qq snapshot_modify_snapshot command and specify the snapshot ID from the Saved Snapshots page in the Web UI (for example, for the 1234567_replication_from_prod snapshot, the ID is 1234567).

In the following example, the snapshot with ID 1234567 never expires.

```
qq snapshot_modify_snapshot \
  -i 1234567
  -e ''
```

In the following example, the snapshot with ID 1234567 expires after one month.

```
qq snapshot_modify_snapshot \
  -i 1234567 \
  -t 1month
```

Locking and Unlocking Snapshots in Qumulo Core

This section explains how to lock or unlock a snapshot by using a key located in the Qumulo file system key store and the qq CLI. In addition, it explains how to lock policy-created snapshots for local policies and for policies that are part of a replication target relationship.

For more information, see Managing Security Keys in the Qumulo File System Key Store (page 100).

Locking and Unlocking Snapshots

A Important

- Unlocking a snapshot requires a cryptographic signature. Before you lock a snapshot, make sure that you have access to your private keys and that you understand the unlocking procedure.
- It isn't possible to delete or shorten the expiration time of a locked snapshot. However, you can extend the expiration time of a locked snapshot.
- Qumulo Core removes both locked and unlocked snapshots at their expiration time automatically.

In Qumulo Core 6.1.0.1 (and higher), you can lock a snapshot by using a key located in the Qumulo file system key store (page 100). You can also ensure that a snapshot policy locks all new snapshots with a particular key (page 0) by associating the key with the snapshot policy.

In Qumulo Core 6.1.1 (and higher), you can ensure that a replication target relationship locks all new policy snapshots with a specific key (page 0) by associating the key with the replication target.

To Lock a Snapshot by Using the qq CLI

Use the qq snapshot_lock_snapshot command and specify the snapshot ID and either the key ID or key name. For example:

```
qq snapshot_lock_snapshot \
    --id 1682119059 \
    --lock-key my-key-name
```

To Unlock a Snapshot by Using the qq CLI

Unlocking a snapshot requires proving that you can sign a challenge by using the same key that locked the snapshot. You can do this by using either of the following methods.

If You Have Direct Access to the Private Key

O Note

To use this method, you must install the Python cryptography library.

Use the qq snapshot_unlock_snapshot command and specify the snapshot ID and the path to the private key file. For example:

```
qq snapshot_unlock_snapshot \
    --id 1682119059 \
    --private-key-file /path/to-my-file.pem
```

If You Don't Have Direct Access to the Private Key

If you can use the private key only to sign data, take the following steps.

1. To receive the unlock challenge, run the qq snapshot_get_unlock_challenge command and specify the snapshot ID. For example:

```
qq snapshot_get_unlock_challenge \
--id 1682119059
```

▲ Important

Any change to the snapshot after you receive the unlock challenge, but before you perform the unlock, causes the challenge to change.

2. To generate a verification signature, use the response from the challenge with your private key.

For more information about creating a verification signature by using a private key or key management service, see Signing a Security Challenge by Using an ECDSA Private Key (page 97).

3. To unlock the snapshot, use the qq snapshot_unlock_snapshot command and specify the snapshot ID and the Base64-encoded unlock challenge that your private key signed. For example:

```
qq snapshot_unlock_snapshot \
    --id 1682119059 \
    --signature "VGhpcyBpcyBteSB1bmxvY2sgY2hhbGxlbmdlLg=="
```

Associating a Lock Key with a Replication Target Relationship

To lock all policy-created snapshots by using a lock key, you can associate the key with a replication target relationship. Consider the following system behavior:

- · Qumulo Core locks only policy-created snapshots that have an expiration time.
- If you reverse the relationship by switching the source and target, the new target can't use the existing key and you must set a key for the new target. However, if you revert the relationship by returning the source and target to their original assignments, Qumulo Core lets you use the key from the original source-target relationship.
- If a target replication relationship uses a key, you can't disable or delete the key, unless you reverse the relationship.
- If you disable or delete a key while a target replication relationship is reversed and then return the source and target to their original assignments, you must set a new key to be able to lock future snapshots.

To Associate a Lock Key with a Replication Target Relationship

Run the qq replication_set_target_relationship_lock command and specify the relationship ID and key name or ID. For example:

```
qq replication_set_target_relationship_lock \
    --relationship-id 12345a6b-7c89-0d12-3456-78fe9012f345
    --lock-key my-key-name
```

To Disassociate a Lock Key from a Replication Target Relationship

Run the qq replication_set_target_relationship_lock command and specify the relationship ID and and --clear-lock-key flag. For example:

```
qq replication_set_target_relationship_lock \
    --relationship-id 12345a6b-7c89-0d12-3456-78fe9012f345
    --clear-lock-key
```

Recovering Files by Using Snapshots

This section explains how to use snapshots to recover files.

In Qumulo Core 2.5.0 (and higher), you can recover files by accessing the hidden .snapshot directory over SMB or NFS.

Inside the ...snapshot directory, directories with snapshot IDs represent various snapshots. The *modified* timestamp of a directory is the time at which Qumulo Core took the snapshot.

O Note

- When you use NFS on Linux and macOS (even if you configure your system to show hidden files), the .snapshot directory doesn't appear when you list a directory's contents. You must navigate to the .snapshot directory explicitly.
- When you use SMB, the .snapshot directory appears only at the root of the share in Finder or File Explorer. In other directories, you must navigate to the .snapshot directory explicitly.

To Recover Files on Linux or macOS by Using the Command Line

1. Navigate to the .snapshot directory. For example:

cd /Volumes/MyShareName/.snapshot

2. Locate the file or directory to recover and copy it to a new location.

☑ Tip

To see the .snapshot directory at the root of the share, show hidden files by pressing # + Shift + .

To Recover Files on macOS by Using Finder

- 1. On the Finder menu, click Go > Go to Folder....
- 2. In the dialog box, enter the path to the .snapshot directory. For example:

/Volumes/MyShareName/.snapshot

O Note

You must specify the .snapshot directory from the root of the share.

3. Locate the file or directory to recover and copy it to a new location.

To Recover Files on Windows by Using File Explorer

- 1. On Windows 7 (and higher), configure Windows Explorer (or File Explorer) as follows:
 - a. Disable Hide protected operating system files.
 - b. Enable Show hidden files, folders, and drives.
- 2. Navigate to the .snapshot directory.
- 3. Locate the file or directory to recover and copy it to a new location.

Encryption and Data Security

Managing Encryption at Rest in Qumulo Core

This section explains how encryption at rest works in Qumulo Core, how to rotate master keys, and how to ensure that the master keys across your cluster are secured correctly.

How Master Keys and Data Keys Work in Qumulo Core

In addition to encrypting data *in transit* (for example, to clients that use SMBv.3.1), in Qumulo Core 3.1.5 (and higher) software-based encryption secures data *at rest* for on-premises clusters. Qumulo Core encrypts all data and metadata in the file system. Removing or reinserting drives and replication doesn't affect encryption at rest. For more information, see Encryption Limitations (page 93).

To encrypt data and data keys, Qumulo Core uses *master keys* that it stores on the boot drive of every node in the cluster, in a file that only the **root** user can access. The *master key* protects the *data key* that encrypts the data on the cluster. This lets Qumulo Core protect your data from potential threats such as a malicious actors' access to stolen or decommissioned disks.

▲ Important

- While the *master key* on your boot drive encrypts your *data keys*, the master key *itself* isn't encrypted.
- Qumulo Care team members can help you rotate your master keys (page 0).
 However, they don't have access to your encryption keys and can't retrieve them for you.
- The boot drive contains the disk image, the installed build of Qumulo Core, and configuration files. In the unlikely event that your boot drive fails and requires replacement, remove the encrypted data keys associated with the master key from the boot drive by rotating the master key (page 0). When you complete the key rotation process, you can dispose of the failed boot drive securely.
- To avoid potential decryption, ensure that your data keys eventually *age out* by rotating the master key any time you replace a drive in your cluster.

While encryption at rest is available by default on all clusters that you create by using Qumulo Core 3.1.5 (and higher), upgrading from a version of Qumulo Core lower than 3.1.5 doesn't provide encryption capability. You must rebuild your cluster to take advantage of this feature. When you create a new cluster, Qumulo Core enables encryption at rest automatically and distributes the master key to all nodes in the cluster.

For replication, Qumulo Core maintains the encryption type after a transfer. Although source and target clusters don't require encryption for replication, we strongly recommend encrypting both source and target clusters.

Rotating the Master Key

This section explains how to rotate the master key and check the encryption status for your cluster by using the qq CLI and Web UI.

To Rotate the Master Key by Using the qq CLI

1. Use the qq rotate_encryption_keys command.

When the command completes, it shows the following message.

Key rotation complete

2. To check your cluster's encryption status and the last key rotation time, use the qq encryption_status command.

To Check Your Cluster's Encryption Status by Using the Web UI

- 1. Log in to Qumulo Core.
- 2. On the Dashboard page, in the Cluster Overview section, click More details.
- 3. If encryption is enabled for your cluster, the Cluster page shows the message Data Encrypted.

Encryption Limitations

- Qumulo Core doesn't encrypt host file system data on the node (such as system logs, core files, and so on).
- · Qumulo Core doesn't support removing encryption from encrypted clusters.
- Single-stream write throughput and latency might experience up to 10-10% degradation and reads might experience up to 5% degradation.
- Qumulo Cloud clusters don't support encryption at rest and should use cloud-native solutions for this functionality.

Generating and Storing ECDSA Keys on a Qumulo Cluster

This section explains how to generate Elliptic Curve Digital Signature Algorithm (ECDSA) keys and ECDSA verification signatures that are compatible with the Qumulo file system key store.

In Qumulo Core 6.1.0 (and higher), you can store multiple ECDSA public keys in the Qumulo file system key store and use these keys to protect filesystem resources.

▲ Important

- Currently, Qumulo Core supports only 256-bit ECDSA keys in .pem and .der formats. Qumulo Core doesn't support storing ECDSA keys of other lengths and formats.
- No KMS system shows the private key. To sign messages later, write down the key ID in the responses from key generation commands.

Generating an ECDSA Private Key

This section explains how to generate a 256-bit ECDSA private key by using Linux CLI tools and AWS, GCP, and Azure CLI or API.

To Generate a Private Key by Using Linux CLI Tools

To generate a key in the .pem format, use the openssl or ssh-keygen tools.

· Use the opensal command and specify the path to the private key. For example:

```
openssl ecparam \
-genkey \
-name prime256v1 \
-out /private-key-path
```

· Use the ssh-keygen command and specify the path to the private key. For example:

```
ssh-keygen \
  -f /private-key-path \
  -t ecdsa \
  -m PEM
```

The following is an example private key.

Generating a Key Pair by Using the AWS Key Management Service (KMS)

Use the AWS Management Console, AWS CLI, or AWS KMS API. For more information, see the following resources:

- · Console: Creating asymmetric KMS keys
- · AWS CLI: create-key in the AWS CLI Command Reference.
- · AWS KMS API: CreateKey in the AWS Key Management Service API Reference

When you create a key pair, specify the following details:

- · Key Type: Asymmetric key
- · Usage: Sign and verify
- Key Specification: ECC_NIST_P256

Generating a Key Pair by Using the GCP Cloud Key Management Service (Cloud KMS)

Use the GCP Cloud Console, Cloud CLI, or Cloud API. For more information, see Create a key in the Cloud Key Management Service documentation.

When you create a key pair, specify the following details:

- Protection Level: software or HSM
- · Purpose: Asymmetric sign
- · Algorithm: Elliptic Curve P-256 SHA256 Digest

Generating a Key Pair by Using the Azure Key Vault

Use the Azure Key Vault and the Azure CLI. For more information, see az keyvault key create in the Azure documentation.

When you create a key pair, specify the following details:

· Key Type: EC

Curve: P-256

Key Size: 256

Extracting the Public Key from an ECDSA Private Key

After you create a 256-bit ECDSA private key, you can extract a public key from it by using Linux CLI tools and AWS, GCP, and Azure CLI and API. You can store the public key in the Qumulo file system key store (page 0).

To Extract the Public Key by Using Linux CLI Tools

1. Use the openssl tool and specify the path to the private key and the path for saving the public key. For example:

```
openssl pkey \
  -in /private-key-path \
  -pubout > /public-key-path
```

2. If your private key is in OpenSSH format, export the public key into the .pem format. Use the ssh-keygen tool and specify the path to the private key and the path for saving the public key. For example:

```
ssh-keygen \
-e \
-f /private-key-path \
-m PEM > /public-key-path
```

3. To convert your private key to .pem format, you can use the ssh-keygen tool and specify the path to the private key. For example:

```
ssh-keygen \
-p \
-f /private-key-path \
-m pem
```

The following is an example public key in .pem format.

```
----BEGIN PUBLIC KEY----
EXAMPLEabCDef0GHIJKL4MN0PqRStUV5wXyz491abc1d2efGijklmN0P0qrsTUv
WXYza1BCdEfGHIjk0lMn0pqr1STUvW3XYzAB6c8DefghIJkEXAMPLE==
----END PUBLIC KEY----
```

Extracting the Public Key by Using the AWS Key Management Service (KMS)

Use the AWS Management Console, AWS CLI, or AWS KMS API. For more information, see the following resources:

Console: Displaying KMS key details

AWS CLI: get-public-key

AWS API: GetPublicKey

Extracting the Public Key by Using the GCP Cloud Key Management Service (Cloud KMS)

Use the GCP Cloud Console, Cloud CLI, or Cloud API. For more information, see Retrieve the public key in the Cloud Key Management Service documentation.

Extracting the Public Key by Using the Azure Key Vault

Use the Azure Key Vault and the Azure CLI. For more information, see az keyvault key download in the Azure documentation.

Signing a Security Challenge by Using an ECDSA Private Key

When you perform actions such as adding a new key to the Qumulo file system key store, replacing an existing key in the key store, or unlocking a snapshot, you must verify that you have access to the private key by signing a security challenge.

You can use your private key to generate a verification signature and then provide this signature to Qumulo Core in Base64 encoding.

Note

- · When you add a new key, the key name is the security challenge.
- When you replace an existing key or unlock a snapshot, the qq CLI command provides the challenge.

To Sign a Security Challenge by Using Linux CLI Tools

1. Save the security challenge to be signed to a file. For example:

```
echo -n "This is my challenge." > /tmp/challenge.out
```

O Note

The -n flag ensures that there are no newline characters following the challenge.

2. To sign the challenge, use the openssl tool and specify the path to the private key. For example:

```
openssl dgst \
  -sha256 \
  -r \
  -sign /path-to-private-key \
  -out /tmp/signature.sha256 /tmp/challenge.out
```

3. To encode the signature in Base64 format, use the openssl tool. For example:

```
openssl base64 \
  -in /tmp/signature.sha256 \
  -out /tmp/key_signature.b64
```

To Sign a Security Challenge by Using the AWS Key Management Service (KMS)

- 1. Sign a security challenge by using the AWS CLI or AWS KMS API. For more information, see the following resources:
 - · AWS CLI: sign in the AWS CLI Command Reference
 - · AWS KMS API: Sign in the AWS Key Management Service API Reference
- 2. Specify the ECDSA_SHA_256 algorithm.

The response returns a Base64-encoded verification signature.

To Sign a Security Challenge by Using the GCP Cloud Key Management Service (Cloud CMS)

- 1. Sign a security challenge by using the GCP Cloud Console, Cloud CLI, or Cloud API. For more information, see Creating a signature in the Cloud Key Management Service documentation.
- 2. Specify the sha256 digest algorithm.
- 3. If the signature in the response is comprised of raw bytes (not a Base64-encoded string) encode the signature file by using the base64 CLI tool on the signature file that the Cloud CLI generates. For example:

```
base64 /gcp-output-path/signature.bytes
```

To Sign a Security Challenge by Using the Azure Key Vault

A Important

The Azure API requires a security challenge as a UrlBase64-encoded SHA-256 digest.

- 1. Sign a security challenge by using the Azure Key Vault and the Azure API. For more information, see sign in the Azure documentation.
- 2. Convert your plaintext challenge into the correct format. For example:

- 3. Use the re-encoded challenge to call the Azure API.
- 4. Specify ES256 as the algorithm.

The signature in the response is encoded in UrlBase64 format.

5. Encode the signature in Base64 format. For example:

```
echo $(echo -n VGhpcyBpcyBteSBzaWduYXR1cmUu | tr '_-' '/+')==
```

Managing Security Keys in the Qumulo File System Key Store

This section explains how to manage security keys in the Qumulo file system key store by using the qq CLI.

In Qumulo Core 6.1.0 (and higher), you can store multiple ECDSA public keys in the Qumulo file system key store and use these keys to protect filesystem resources.

▲ Important

- Currently, Qumulo Core supports only 256-bit ECDSA keys in .pem and .der formats. Qumulo Core doesn't support storing ECDSA keys of other lengths and formats.
- Qumulo Core retains only the public key. We strongly recommend storing the corresponding private key safely, outside of your Qumulo cluster and according to your organization's security policy.

For information about protecting your snapshots by using a key from the Qumulo file system key store, see Locking and Unlocking Snapshots (page 87).

☑ Tip

The response of the qq fs_security_add_key command includes the key identifier. When you use qq fs_security commands, you can specify either the key identifier (by using the --id flag) or the key name (by using the --name flag).

Adding a Public Key

This section explains how to add a public key to the Qumulo file system key store. To store a public key in the key store, you must have a pair of asymmetric keys. For more information, see Generating an ECDSA Private Key (page 0).

If You Have Access to the Private Key

Note

To use this method, you must install the Python cryptography library.

Use the qq fs_security_add_key command and specify the key name, the path to the private key file, and an optional comment. For example:

```
qq fs_security_add_key \
    --name my-key-name \
    --private-key-file /path/to-my-file.pem \
    --comment "This is an optional comment."
```

If You Don't Have Direct Access to the Private Key

Use the qq fs_security_add_key command and specify the key name, the public key contents, the Base64-encoded verification signature (the key name signed with the private key), and an optional comment. For example:

```
qq fs_security_add_key \
    --name my-key-name \
    --public-key "VGhpcyBpcyBteSBwdWJsaWMga2V5IGNvbnRlbnRzLg==" \
    --verification-signature "VGhpcyBpcyBteSB1bmxvY2sgY2hhbGxlbmdlLg==" \
    --comment "This is an optional comment."
```

For more information, see Extracting the Public Key from an ECDSA Private Key (page 96) and Signing a Security Challenge by Using an ECDSA Private Key (page 97).

Retrieving Public Key Information

- To retrieve information for a single public key, use the qq fs_security_key command and specify the key identifier or name.
- To retrieve information for all public keys, use the qq fs_security_list_keys command.

The output displays information in a table format. To view the output in JSON format, use the --json flag.

Retrieving Public Key Usage Information

Use the qq fs_security_get_usage command and specify the key identifier or name.

The output displays information in a table format. To view the output in JSON format, use the -json flag.

Modifying a Public Key

To modify the name or comment for a public key, use the qq fs_security_modify_key command and specify the key identifier or name and the flags for the fields to modify.

Rotating a Public Key

This section explains how to rotate a public key in the Qumulo file system key store.

O Note

- Key rotation doesn't affect the resources that the key protects or change the identifier or name of the public key.
- When the key rotation is complete, only the replacement public key can unlock the protected resources.

If You Have Access to the Existing and Replacement Private Keys

O Note

To use this method, you must install the Python cryptography library.

Use the qq fs_security_replace_key command and specify the key identifier or name, the path to the existing private key, and the path to the replacement private key. For example:

```
qq fs_security_replace_key \
    --name my-key-name \
    --old-private-key-file /path/to-existing-key.pem \
    --replacement-private-key-file /path/to-replacement-key.pem
```

If You Don't Have Direct Access to the Existing and Replacement Private Keys

- To receive the key replacement challenge, use the qq fs_security_get_key_replace_challenge command and specify the identifier or name of the key to replace.
- 2. To generate a verification signature, use the response from the challenge with the existing private key and another verification signature by using the challenge and the replacement private key.
 - For more information, see Signing a Security Challenge by Using an ECDSA Private Key (page 97).
- 3. To rotate the key, use the qq fs_security_replace_key command and specify the key identifier or name, the replacement public key contents, the replacement key verification signature (Base64-encoded key replacement challenge signed with the replacement private key), and the existing key verification signature (Base64-encoded key replacement challenge signed with the existing private key). For example:

```
qq fs_security_replace_key \
    --name my-key-name \
    --replacement-public-key "VGhpcyBpcyBteSByZXBsYWNlbWVudCBwdWJsaWMga2V5Lg=="
    --replacement-key-verification-signature "UmVwbGFjZW1lbnQga2V5IHZlcmlmaWNhdG
lvbiBzaWduYXR1cmU=" \
    --old-key-verification-signature "RXhpc3Rpbmcga2V5IHZlcmlmaWNhdGlvbiBzaWduYX
R1cmU="
```

For more information, see Extracting the Public Key from an ECDSA Private Key (page 96).

▲ Important

Because the key version is part of the challenge message, and this version changes when a user writes or modifies the key, any change to the key name or comment after you receive the challenge message makes the message stale.

Disabling a Public Key

When you add a key to the Qumulo file system key store, Qumulo Core enables it automatically.

- To disable a key, use the qq fs_security_modify_key command and specify the key identifier or name and the --disable flag.
- · To re-enable a key, use the --enable flag.

Note

- It isn't possible to lock a Qumulo file system resource with a disabled key. However, you can still unlock resources that this key locked previously.
- It isn't possible to disable a key that Qumulo Core uses to create new resources. For example, you can't disable a key associated with a snapshot policy when the snapshot policy creates new snapshots by using the key. In this scenario, you must disassociate the key from the snapshot policy before you can disable it. For more information, see Retrieving Public Key Usage Information (page 0).

Deleting a Public Key

Use the qq fs_security_delete_key and specify the key identifier or name.

Note

It isn't possible to delete a key that a Qumulo file system resource uses. For more information, see Retrieving Public Key Usage Information (page 0).

Node Addition and Replacement

Adding Nodes to an Existing Qumulo Cluster

This section explains how to add new HPE, Supermicro, or Quiver nodes to an existing cluster.

After you connect and power on your new nodes, Qumulo Core discovers any unconfigured nodes automatically and prompts you to add nodes in the Web UI.

If Qumulo Core doesn't discover any unconfigured nodes, it displays the message **No unconfigured nodes found**. If you expect to see nodes, contact the Qumulo Care team.

O Note

- · Qumulo Core requires a short time to update the total available storage.
- · Existing nodes retain their numbering.

Prerequisites

- Sufficient Static IP Addresses: The number of static IP addresses must be equal to or
 greater than the number of nodes in your cluster. For more information, see IP Failover
 with Qumulo Core on Qumulo Care.
- Same Qumulo Core Version on All Nodes: For information about upgrading Qumulo Core, see Performing a Clean Installation of Qumulo Core and Performing Qumulo Core Upgrades by Using the Web UI on Qumulo Care.

Step 1: Resolve Drive Compatibility Issues

▲ Important

If the version of Qumulo Core on your existing nodes predates the Qumulo-certified drives that you received with your new nodes, you can't install a lower version of Qumulo Core on your new node and Qumulo Core displays the message Installation failed. Use the cluster logs to identify any incompatible drives.

To receive support for new, Qumulo-certified drives, do one of the following:

- Upgrade your existing cluster to the latest version of Qumulo Core and then install the same version of Qumulo Core on your new node. For more information, see Performing Qumulo Core Upgrades by Using the Web UI on Qumulo Care.
- Update the Supported Drive List on your new node. For more information, contact the Qumulo Care team.

Step 2: Add Your New Nodes to an Existing Qumulo Cluster

- 1. Log in to the Web UI.
- 2. Click Cluster > Add Nodes.
- 3. On the Add Nodes page, select unconfigured nodes to add to your cluster.
- 4. Click Add Selected Nodes to Cluster.
- 5. In the Add <N> nodes to cluster <name>? dialog box, click Yes.

If you add one or more node model types, a message reminds you about Qumulo Core adding a new model type to your cluster.

Qumulo Core configures your new nodes and adds them to your cluster.

On the Cluster page, the Web UI shows the banner Successfully added <N> nodes to the cluster and the total available storage.

Increasing the Node Fault Tolerance Level for Your Qumulo Cluster during Node-Add Operations

This section explains how to increase the node fault tolerance level for your Qumulo cluster during node-add operations.

- In Qumulo Core 5.1.3 (and higher) you can increase the node fault tolerance level for an existing cluster during the cluster expansion process.
- In Qumulo Core 6.1.0.1 (and higher), you can manage the drive and node fault tolerance levels during the cluster expansion process by changing the data protection configuration (page 108).

A Important

We strongly recommend contacting the Qumulo Care team before proceeding with cluster expansion.

To Reconfigure Your Cluster's Node Fault Tolerance Level

1. Follow the instructions in Adding Nodes to an Existing Qumulo Cluster (page 105).

O Note

- When a tradeoff between increasing node fault tolerance and maximizing usable capacity is available for your cluster, the Cluster > Add Nodes page shows the Accept a trade-off in the increase of usable capacity checkbox.
- When a tradeoff isn't available, the page shows an message which indicates that it is possible to increase the node fault tolerance level.
- 2. Before you click Yes in the Add <N> nodes to cluster <MyCluster>? dialog box, check that the projected capacity matches the expected capacity.
- 3. After the cluster expansion process finishes, Qumulo Core begins data protection reconfiguration automatically.
- 4. To monitor this process, click Cluster > Overview. On the Cluster page, in the protection status section, you can view the rebalance phase status and the estimated time to completion.
- 5. When the restriper completes the data protection reconfiguration, the Data Protected section of the Cluster page shows the increased node fault tolerance level.

Adding Nodes with Adaptive Data Protection to a Qumulo Cluster

This section explains how to configure Adaptive Data Protection for your Qumulo cluster during node-add operations.

- Qumulo Core 6.1.0.1 and 6.1.1 support Adaptive Data Protection by letting you reconfigure
 your cluster's fault tolerance level and storage efficiency only when you add nodes to your
 cluster.
 - Depending on your cluster, Qumulo Core shows configuration options that offer better fault tolerance levels, better storage efficiency, or both benefits. To enable Adaptive Data Protection for your cluster, you must contact the Qumulo Care team.
- In Qumulo Core 6.1.2 (and higher), you can change your cluster's data protection configuration when you add or replace nodes by using the qq CLI.

1 Note

Your cluster's usable capacity doesn't increase until the data protection reconfiguration is complete. Because data protection reconfiguration is a long process (that can take days or weeks), we don't recommend beginning this process when your cluster is nearly running out of usable capacity. In this scenario, contact the Qumulo Care team.

To Add a Node with Adaptive Data Protection to a Qumulo Cluster

A cluster's data protection configuration includes the stripe configuration (for example, 10.8) and the node fault tolerance level.

A Important

Ensure that your configuration and node order are correct. Unless you add more nodes, it isn't possible to revert this operation or reorder nodes after adding them to a cluster.

- 1. To select the new data protection configuration for your nodes, contact the Qumulo Care team.
- 2. Rack and wire your new nodes and then power them on.
- 3. To determine the UUIDs of the nodes to add to your cluster, use the qq unconfigured_nodes_list command.
- 4. Write down the UUIDs of the nodes that you want to add to the cluster, in the order that you want to add them.

5. To add your nodes to the cluster, use the qq add_nodes command and specify the node UUIDs (their ordering in the command determines their ordering in the cluster), the configuration type, and the node fault tolerance level.

In the following example, we change the stripe configuration to 10.8 and the node fault tolerance level to 2 nodes.

```
qq add_nodes \
    --node-uuids 12345a6b-7c89-0d12-3456-78fe9012f345 \
    abcde1f2-g3hi-j4kl-mnop-qr56stuv7wxy \
    --target-stripe-config 10 8 \
    --target-max-node-failures 2
```

The following is example output from the command.

```
Current cluster:
    Usable capacity: 600 TB
    Node fault tolerance level: 1 node
With the selected node-add operation and data protection reconfiguration:
    Usable capacity: 800 TB
    Node fault tolerance level: 2 nodes
```

6. To confirm the configuration with the selected node-add and data protection configuration operations, enter yes.

Monitoring the Data Protection Reconfiguration Process

To view the progress of the three stages of the data protection reconfiguration process, log in to the Qumulo Core Web UI and click **Cluster**.

- 1. Qumulo Core begins to move data to new nodes in the cluster and the Web UI displays the message Rebalancing for data protection reconfiguration.
- 2. Qumulo Core reencodes all data on your cluster and the Web UI displays the message Reconfiguring data protection.

O Note

In certain scenarios, this stage might appear to pause while the system performs preparatory work on the cluster.

When this stage is complete, your data is protected according to the cluster's new configuration and the system begins to use the new drive and node fault tolerance levels.

3. Qumulo Core adds new capacity to your cluster and the Web UI displays the message Rabalancing.

If you initiated the reconfiguration process as part of a node replacement step, the system migrates data from the existing nodes in the cluster.

Cluster Availability During the Reconfiguration Process

Your cluster remains available throughout the data protection reconfiguration process.

- · You can upgrade Qumulo Core.
- · Your cluster maintains the ability to recover from node and drive failure automatically.

During the reconfiguration process, drive and node fault tolerance levels remain at the minimums that the existing and new configurations specify. For example, if your existing cluster has 2-node and 2-drive fault tolerance, and you initiate reconfiguration where the new configuration has 1-node and 3-drive fault tolerance, your cluster has 1-node and 2-drive fault tolerance during the reconfiguration process.

Note

- To avoid impact to frontend workloads, Qumulo Core slows down the reconfiguration process automatically.
- When Qumulo Core finds missing nodes or drives, it pauses the reconfiguration process. When you replace or bring the nodes or drives online, the reconfiguration process continues.
- · It isn't possible to add or replace nodes during the reconfiguration process.

Replacing Nodes in a Qumulo Cluster by Performing a Transparent Platform Refresh

This section explains how to replace nodes that have reached retirement or end of life by performing a two-stage transparent platform refresh on clusters that run Qumulo Core 6.1.0.1 (and higher).

O Note

- · Qumulo Core doesn't support replacing nodes in clusters with more than 100 nodes.
- In Qumulo Core 6.1.2.2 (and higher), you can use the qq CLI to replace nodes. To replace nodes on a lower version of Qumulo Core, contact the Qumulo Care team.

How Transparent Platform Refresh Works

Transparent platform refresh comprises two stages. For help with your node replacement plan, contact the Qumulo Care team.

Stage 1: Register a Node Replacement Plan

In this stage, you register a node replacement plan (page 0) with your cluster. The plan includes information about the nodes to replace and the data protection configuration.

In following example, we use a four-node cluster and:

- · Replace nodes 1-4 with five new nodes in a single step
- Change the data protection configuration to the 8.6 stripe configuration with 1-node fault tolerance

Stage 2: Execute the Node Replacement Plan Steps

In this stage, you execute the node replacement plan's steps (page 0) and Qumulo Core performs data protection reconfiguration.

Note

It isn't possible to add nodes or begin another node replacement step while a node replacement step is already in progress.

There are two node replacement plan types:

 Single-Step Node Replacement: Qumulo Core adds all new nodes and removes all nodes marked for replacement in a single step. Use this approach when the node replacement speed is a priority. • Multi-Step Node Replacement: Each step of the plan adds some the new nodes and removes some nodes marked for replacement. Use this approach when rack space or switch port capacity in your data center is limited.

Cluster Properties During Node Replacement

- When a replacement step begins, Qumulo Core distributes floating IP addresses among the nodes in the combined cluster (page 112). After Qumulo Core removes nodes marked for replacement, it redistributes any client connections that use floating IP addresses among the nodes that remain in the cluster.
- While a node replacement step is in progress, both new nodes and nodes marked for replacement appear on the Cluster page of the Web UI and clients can connect to any of the nodes in the combined cluster while the step is in progress.
- When a node replacement step is complete, Qumulo Core reassigns static IP addresses to different nodes. To view the new assignment, click Cluster > Network Configuration in the Web UI.
- When Qumulo Core adds nodes to a cluster, it assigns node IDs sequentially, without reusing or changing IDs.

For example, if you have a four-node cluster with node IDs 1-4, and you replace node IDs 2 and 3 with two new nodes, after node replacement the cluster contains node IDs 1, 4, 5, and 6. If you add another node, it has the ID 7.

- · A cluster's usable capacity doesn't increase until:
 - Any data protection reconfiguration is complete
 - The last step of the node replacement plan is in progress

For example, if you replace nodes in a single step without data protection reconfiguration, usable capacity increases as soon as Qumulo Core begins the step.

Prerequisites

Ensure that the number of static and floating IP addresses is equal to or greater than the number of nodes in the combined cluster (page 112).

Step 1: Register a Node Replacement Plan by Using the qq CLI

1. Use the qq replace_nodes register_plan command and the --nodes-to-be-replaced flag to specify the nodes to replace and the --target-stripe-config flag to specify the stripe configuration. For example:

```
qq replace_nodes register_plan \
    --nodes-to-be-replaced 1 2 3 4 \
    --target-stripe-config 8 6
```

Qumulo Core stores the node replacement plan on your cluster.

O Note

- If your plan includes data protection reconfiguration, Qumulo Core records only the stripe configuration. You specify the node fault tolerance when you execute the plan steps.
- If your plan doesn't include data protection reconfiguration, you can omit the --target-stripe-config flag.
- To replace all nodes in the cluster, use the --replace-all flag instead of the --nodes-to-be-replaced flag.
- 2. Rack and wire your new nodes and then power them on.
- 3. To determine the UUIDs of the nodes to add to your cluster, use the qq unconfigured_nodes_list command.
- 4. Write down the UUIDs of the nodes that you want to add to the cluster, in the order that you want to add them.

Step 2: Execute the Node Replacement Plan Steps by Using the qq CLI

1. Use the qq replace_nodes add_nodes_and_replace command to initiate each step, the --nodes-being-replaced flag to specify the nodes to replace, and the --node-uuids flag to specify the nodes to add during the current step.

▲ Important

Qumulo Core adds nodes to the cluster in the order in which you list their UUIDs after the --node-uuids flag. When you begin the node replacement step, it isn't possible to revert this operation or reorder nodes after adding them to a cluster.

If your plan includes data protection reconfiguration, use the --reconfigure-data-protection and --target-max-node-failures flags to initiate the reconfiguration during the current step. For example:

```
qq replace_nodes add_nodes_and_replace \
    --nodes-being-replaced 1 2 3 4 \
    --node-uuids 12345a6b-7c89-0d12-3456-78fe9012f345 abcde1f2-g3hi-j4kl-mnop-qr
56stuv7wxy \
    --reconfigure-data-protection \
    --target-max-node-failures 1
```

The following is example output from the command:

```
Current cluster:
    Usable capacity: 200 TB
    Node fault tolerance level: 1 node
With the selected node replacement step:
    Usable capacity: 220 TB
    Node fault tolerance level: 1 node
```

O Note

To replace all nodes in the cluster, use the --replace-all flag instead of the --nodes-being-replaced flag.

2. To confirm the reconfiguration with the selected node-replace and data protection configuration operations, enter yes.

For more information, see Monitoring the Data Protection Reconfiguration Process (page 0).

3. Wait for the node replacement step to complete.

After each node replacement step, Qumulo Core begins to migrate data from existing nodes in the backround.

O Note

This is a long process (that can take days or weeks). When the data migration is complete, Qumulo Core removes the nodes marked for replacement from the cluster. These nodes no longer appear on the Cluster page of the Web UI.

- 4. Unrack the removed nodes from your data center.
- 5. Initiate the next node replacement step.

Viewing, Editing, and Cancelling the Node Replacement Plan

• To view the current node replacement plan, use the qq" replace_nodes get_plan command.

If a node replacement step is in progress, the command shows the list of nodes in process of being replaced during the current step.

- To edit the node replacement plan after you register it with your cluster, use the qq replace_nodes_register_plan command with a new node replacement plan (page 0).
- To cancel the current node replacement plan, use the qq replace_nodes cancel_plan command.

A Important

Caneling a node replacement plan after executing one or more steps might make it impossible to reregister and complete the plan.

Monitoring the Data Protection Reconfiguration Process

To view the progress of the three stages of the data protection reconfiguration process, log in to the Qumulo Core Web UI and click Cluster.

- 1. Qumulo Core begins to move data to new nodes in the cluster and the Web UI displays the message Rebalancing for data protection reconfiguration.
- 2. Qumulo Core reencodes all data on your cluster and the Web UI displays the message Reconfiguring data protection.

O Note

In certain scenarios, this stage might appear to pause while the system performs preparatory work on the cluster.

When this stage is complete, your data is protected according to the cluster's new configuration and the system begins to use the new drive and node fault tolerance levels.

3. Qumulo Core adds new capacity to your cluster and the Web UI displays the message Rabalancing.

If you initiated the reconfiguration process as part of a node replacement step, the system migrates data from the existing nodes in the cluster.

Cluster Availability During the Reconfiguration Process

Your cluster remains available throughout the data protection reconfiguration process.

· You can upgrade Qumulo Core.

Your cluster maintains the ability to recover from node and drive failure automatically.
 During the reconfiguration process, drive and node fault tolerance levels remain at the minimums that the existing and new configurations specify. For example, if your existing cluster has 2-node and 2-drive fault tolerance, and you initiate reconfiguration where the new configuration has 1-node and 3-drive fault tolerance, your cluster has 1-node and 2-drive fault tolerance during the reconfiguration process.

O Note

- To avoid impact to frontend workloads, Qumulo Core slows down the reconfiguration process automatically.
- When Qumulo Core finds missing nodes or drives, it pauses the reconfiguration process. When you replace or bring the nodes or drives online, the reconfiguration process continues.
- · It isn't possible to add or replace nodes during the reconfiguration process.

Improving Performance by Migrating a Qumulo Cluster to a Different License Class

This section explains how you can improve the performance of a cluster that runs Qumulo Core 6.1.0.1 (or higher) by migrating it a different license class (for example, from hybrid to all-NVMe nodes).

A Important

Because it isn't possible to have nodes with different license classes in the same cluster, you must migrate all nodes in your cluster to the same license class.

For help with migrating your cluster to a different license class, contact the Qumulo Care team.

Understanding Offline Qumulo Core Nodes and Checking for Free Space

This section explains what happens when a Qumulo Core node goes offline and how you can check the remaining free space.

What Happens When a Qumulo Core Node Goes Offline

Qumulo Core uses *erasure coding* to let multiple drives or nodes to go offline but continue to serve data. For more information, see Qumulo Drive Failure Protection on Qumulo Care.

On hybrid Qumulo nodes (that have HDDs and SSDs), Qumulo Core attempts to maintain cluster functionality for as long as possible. When a node goes offline, the cluster evicts existing data promoted to SSDs on the remaining nodes and makes all writes to the free space on the SSDs. When the node comes online, Qumulo Core begins to push writes to the HDDs that back the SSDs.

A Important

Writes, deletes, and changes count towards SSD space. For more information, see Checking Remaining Free Space on Your Cluster (page 118).

A cluster can operate with an offline node until its SSD space fills up. When a cluster has no more free space, the cluster goes into read-only (ENOSPC) state until the node comes back online and all nodes rejoin the quorum.

Note

The amount of time that users have before the cluster enters the ENOSPC state depends on the rate of change in new rates to the cluster.

Checking Remaining Free Space on Your Cluster

You can estimate the amount of time before your cluster enters the **ENOSPC** state by running multiple iterations of the **debug_metrics_get** command (**metrics_get** in Qumulo Core 4.2.0 and lower) on every node in your cluster.

The output of the command shows the number of valid tokens remaining for writes to the cluster. At 50,000 (or fewer) tokens, an **ENOSPC** event is imminent.

Note

Because SSD space on the cluster is limited, new writes, deletes, and changes consume any reclaimable tokens.

Qumulo Core Version	Command
4.2.1 (and higher)	
	<pre>sudo qsh -c /opt/qumulo/qq debug_metrics_get \ measurement space_agent grep reclaimable; sleep 60; \ sudo qsh -c /opt/qumulo/qq debug_metrics_get \ measurement space_agent grep reclaimable</pre>
3.3.2 - 4.2.0	
	<pre>sudo qsh -c /opt/qumulo/qq metrics_get \ measurement space_agent grep reclaimable; sleep 60; \ sudo qsh -c /opt/qumulo/qq metrics_get \ measurement space_agent grep reclaimable</pre>
3.1 - 3.3.1	
	<pre>sudo /opt/qumulo/qq metrics_get \ measurement space_agent grep reclaimable; sleep 60; \ sudo /opt/qumulo/qq metrics_get \ measurement space_agent grep reclaimable</pre>

You can also contact the Qumulo Care team for a time estimate.

Data Replication

Using Qumulo Shift-To to Copy Objects to Amazon S3

This section explains how to use Shift-To to copy objects from a directory in a Qumulo cluster to a folder in an Amazon Simple Storage Service (Amazon S3) bucket and how to manage Shift relationships.

For more information about copying objects from S3 to Qumulo, see Using Qumulo Shift-From for Amazon S3 to Copy Objects (page 131).

Prerequisites

- · A Qumulo cluster with:
 - Qumulo Core 3.2.1 (and higher) for the CLI and 3.2.5 (and higher) for the Web UI
 - HTTPS connectivity to s3.<region>.amazonaws.com though one of the following means:
 - Public Internet
 - VPC endpoint
 - AWS Direct Connect

For more information, see AWS IP address ranges in the AWS General Reference.

- · Membership in a Qumulo role with the following privileges:
 - PRIVILEGE_REPLICATION_OBJECT_WRITE: This privilege is required to create a Shift relationship.
 - PRIVILEGE_REPLICATION_OBJECT_READ: This privilege is required to view the status of a Shift relationship.

O Note

- For any changes to take effect, user accounts with newly assigned roles must log out and log back in (or their sessions must time out).
- Use special care when granting privileges to roles and users because certain privileges (such as replication-write privileges) can use system privileges to overwrite or move data to a location where a user has greater permissions.
 This can give a user access to all directories and files in a cluster regardless of any specific file and directory settings.
- An existing bucket with contents in Amazon S3

- · AWS credentials (access key ID and secret access key) with the following permissions:
 - s3:AbortMultipartUpload
 - s3:GetObject
 - o s3:PutObject
 - o s3:ListBucket

For more information, see Understanding and getting your AWS credentials in the AWS General Reference

Example IAM Policy

In the following example, the IAM policy gives permission to read from and write to the my-folder folder in the my-bucket. This policy can give users the permissions required to run Shift-To jobs.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Action": "s3:ListBucket",
      "Effect": "Allow",
      "Resource": "arn:aws:s3:::my-bucket"
    },
    {
      "Action": [
        "s3:AbortMultipartUpload",
        "s3:GetObject",
        "s3:PutObject"
      "Effect": "Allow",
      "Resource": "arn:aws:s3:::my-bucket/my-folder/*"
    }
  ]
}
```

How Shift-To Relationships Work

Qumulo Core performs the following steps when it creates a Shift-To relationship.

- 1. Verifies that the directory exists on the Qumulo cluster and that the specified S3 bucket exists, is accessible by using the specified credentials, and contains downloadable objects.
- 2. Creates the Shift-To relationship.
- 3. Starts a job by using one of the nodes in the Qumulo cluster.

O Note

If you perform multiple Shift operations, Qumulo Core uses multiple nodes.

- 4. To ensure that the copy is point-in-time consistent, takes a temporary snapshot of the directory (for example, named replication_to_bucket_my_bucket).
- 5. Recursively traverses the directories and files in the snapshots and copies each object to a corresponding object in S3.
- 6. Preserves the file paths in the local directory in the keys of replicated objects.

For example, the file <code>/my-dir/my-project/file.text</code>, where <code>my-dir</code> is the directory on your Qumulo cluster, is uploaded to S3 as the following object, where <code>my-folder</code> is the specified S3 folder.

https://my-bucket.s3.us-west-2.amazonaws.com/my-folder/my-project/file.txt

O Note

This process doesn't encode or transform your data in any way. Shift-To replicates only the data in a regular file's primary stream, excluding alternate data streams and file system metadata such as access control lists (ACLs). To avoid transferring data across the public Internet, a server-side S3 copy operation also copies any hard links to files in the replication local directory to S3 as full copies of objects, with identical contents and metadata.

The following table explains how entities in the Qumulo file system map to entities in an S3 bucket.

Entity in the Qumulo File System	Entity in an Amazon S3 Bucket
Access control list (ACL)	Not copied
Alternate data streams	Not copied
Directory	Not copied (directory structure is preserved in the object key for objects created for files)
Hard link to a non-regular file	Not copied

Entity in the Qumulo File System	Entity in an Amazon S3 Bucket
Hard link to a regular file	Copy of the S3 object
Holes in sparse files	Zeroes (holes are expanded)
Regular file	S3 object (the object key is the file system path and the metadata is the field data)
SMB extended file attributes	Not copied
Symbolic link	Not copied
Timestamps (mtime , ctime , atime , btime)	Not copied
UNIX device file	Not copied

7. Checks whether a file is already replicated. If the object exists in the remote S3 bucket, and neither the file nor the object are modified since the last successful replication, its data isn't retransferred to S3.

O Note

Shift never deletes files in the remote S3 folder, even if the files are removed from the local directory since the last replication.

8. Deletes the temporary snapshot.

Storing and Reusing Relationships

The Shift-To relationship remains on the Qumulo cluster. You can monitor the completion status of a job, start new jobs for a relationship after the initial job finishes, and delete the relationship (when you no longer need the S3-folder-Qumulo-directory pair). To avoid reuploading objects that a previous copy job uploaded, relationships take up approximately 100 bytes for each object. To free this storage, you can delete relationships that you no longer need.

If you repeatedly copy from the same Qumulo directory, you can speed up the upload process (and skip already uploaded files) by using the same relationship.

A new relationship for subsequent uploads doesn't share any tracking information with previous relationships associated with a directory and might recopy data that is already uploaded.

Using the Qumulo Core Web UI to Copy Files and Manage Relationships

This section describes how to use the Qumulo Core Web UI 3.2.5 (and higher) to copy files from a Qumulo cluster to Amazon S3, review Shift relationship details, stop a running copy job, repeat a completed copy job, and delete a relationship.

To Copy Files to Amazon S3

- 1. Log in to Qumulo Core.
- 2. Click Cluster > Copy to/from S3.
- 3. On the Copy to/from S3 page, click Create Copy.
- 4. On the Create Copy to/from S3 page, click Local ⇒ Remote and then enter the following:
 - a. The Directory Path on your cluster (/ by default)
 - b. The S3 Bucket Name
 - c. The Folder in your S3 bucket
 - d. The Region for your S3 bucket
 - e. Your AWS Region (/ by default)
 - f. Your AWS Access Key ID and Secret Access Key.
- 5. (Optional) For additional configuration, click Advanced S3 Server Settings.
- 6. Click Create Copy.
- 7. In the Create Copy to S3? dialog box, review the Shift relationship and then click Yes, Create.

The copy job begins.

To View Configuration Details and Status of Shift Relationships

- 1. Log in to Qumulo Core.
- 2. Click Cluster > Copy to/from S3.

The Copy to/from S3 page lists all existing Shift relationships.

3. To get more information about a specific Shift relationship, click : > View Details.

The Copy to/from S3 Details page displays the following information:

- · Throughput: average
- · Run Time
- · Data: total, transferred, and unchanged
- · Files: total, transferred, and unchanged

To Stop a Copy Job in Progress

- 1. Log in to Qumulo Core.
- 2. Click Cluster > Copy to/from S3.
- 3. To stop a copy job for a specific relationship, click : > Abort.
- 4. In the Abort copy from? dialog box, review the Shift relationship and then click Yes, Abort.

 The copy job stops.

To Repeat a Completed Copy Job

- 1. Log in to Qumulo Core.
- 2. Click Cluster > Copy to/from S3.
- 3. To stop a copy job for a specific relationship, click : > Copy Again.
- 4. In the Copy again? dialog box, review the Shift relationship and then click Yes, Copy Again.

 The copy job repeats.

To Delete a Shift Relationship

- 1. Log in to Qumulo Core.
- 2. Click Cluster > Copy to/from S3.
- 3. To stop a copy job for a specific relationship, click : > Delete.
- 4. In the Delete copy from? dialog box, review the Shift relationship and then click Yes, Delete.

The copy job is deleted.

Using the Qumulo CLI to Copy Files and Manage Relationships

This section describes how to use the Qumulo CLI 3.2.5 (and higher) to copy files from a Qumulo cluster to Amazon S3, review Shift relationship details, stop a running copy job, repeat a completed copy job, and delete a relationship.

Copying Files from Amazon S3

To copy files, use the replication_create_object_relationship command and specify the following:

- · Local directory path on Qumulo cluster
- · Copy direction (copy-to)
- · S3 object folder
- · S3 bucket
- · AWS region

- · AWS access key ID
- AWS secret access key

The following example shows how to create a relationship between the directory /my-dir/ on a Qumulo cluster and the S3 bucket my-bucket and folder /my-folder/ in the us-west-2 AWS region. The secret access key is associated with the access key ID.

```
qq replication_create_object_relationship \
    --source-directory-path /my-dir/ \
    --direction COPY_TO_OBJECT \
    --object-folder /my-folder/ \
    --bucket my-bucket \
    --region us-west-2 \
    --access-key-id AKIAIOSFODNN7EXAMPLE \
    --secret-access-key wJalrXUtnFEMI/K7MDENG/bPxRfiCYEXAMPLEKEY
```

The CLI returns the details of the relationship in JSON format, for example:

```
"access_key_id": "ABC",
"bucket": "my-bucket",
"object_store_address": "s3.us-west-2.amazonaws.com",
"id": "lc23b4ed-5c67-8f90-le23-a4f5f6ceff78",
"object_folder": "my-folder/",
"port": 443,
"ca_certificate": null,
"region": "us-west-2",
"source_directory_id": "3",
"direction": "COPY_TO_OBJECT",
}
```

Viewing Configuration Details and Status of Shift Relationships

- To view configuration details for all Shift relationships, use the replication list object relationships command.
- To view configuration details for a specific relationship, use the
 replication_get_object_relationship command followed by the --id and the Shift
 relationship ID (GUID), for example:

```
qq replication_get_object_relationship --id 1c23b4ed-5c67-8f90-1e23-a4f5f6cef
f78
```

- To view the status of a specific relationship, use the
 replication_get_object_relationship_status command followed by the --id and the
 Shift relationship ID.
- To view the status of all relationships, use the replication_list_object_relationship_statuses command.

The CLI returns the details of all relationships in JSON format, for example:

```
Γ
 {
    "direction": "COPY TO OBJECT",
    "access key id": "AKIAIOSFODNN7EXAMPLE",
    "bucket": "my-bucket",
    "object store address": "s3.us-west-2.amazonaws.com",
    "id": "1c23b4ed-5c67-8f90-1e23-a4f5f6ceff78",
    "object_folder": "my-folder/",
    "port": 443,
    "ca certificate": null,
    "region": "us-west-2",
    "source_directory_id": "3",
    "source directory path": "/my-dir/",
    "state": "REPLICATION RUNNING",
    "current job": {
      "start time": "2020-04-06T17:56:29.659309904Z",
      "estimated_end_time": "2020-04-06T21:54:33.244095593Z",
      "job progress": {
        "bytes transferred": "178388608",
        "bytes unchanged": "0",
        "bytes remaining": "21660032",
        "bytes total": "200048640",
        "files transferred": "17",
        "files unchanged": "0",
        "files remaining": "4",
        "files total": "21",
        "percent_complete": 89.0368314738253,
        "throughput_current": "12330689",
        "throughput overall": "12330689"
      }
    "last job": null
  }
]
```

The state field shows the REPLICATION_RUNNING status and the current_job field shows the job's progress. When Qumulo Core copies files from S3, details for the most recently completed job become available in the last_job field, the state field changes to REPLICATION_NOT_RUNNING, and the current_job field reverts to null.

O Note

If you already ran a job for a relationship, it is possible for both the current_job and last job fields to be non-null while you run a new job.

The bytes_total and files_total fields represent the total amount of data and number of files to be transferred by a Shift job. The bytes_remaining and files_remaining fields show the amount of data and number of files not yet transferred. The values of these four fields don't stabilize until the work estimation for the job is complete.

The percent_complete field displays the overall job progress and the estimated_end_time field displays the time at which the job is estimated to be complete. The values of these two fields are populated when the work estimation for the job is complete.

Stopping a Copy Job in Progress

To stop a copy job already in progress, use the replication_abort_object_relationship command followed by the --id and the Shift relationship ID.

Repeating a Completed Copy Job

To repeat a completed copy job, use the replication_start_object_relationship command followed by the --id and the Shift relationship ID.

This command begins a new job for the existing relationship and downloads any content that changed in the S3 bucket or on the Qumulo cluster since the time the previous job ran.

Deleting a Shift Relationship

After your copy job is complete, you can delete your Shift relationship. To do this, run the replication_delete_object_relationship command followed by the --id and the Shift relationship ID.

Note

You can run this command only against a relationship that doesn't have any active jobs running.

This command removes the copy job's record, leaving locally stored objects unchanged. Any storage that the relationship used to track downloaded objects becomes available when you delete the relationship.

Troubleshooting Copy Job Issues

Any fatal errors that occur during a copy job cause the job to fail, leaving a partially copied set of files in the directory in your S3 bucket. However, to let you review the Shift relationship status any failure messages, the Shift relationship continues to exist. You can start a new job to complete the copying of objects to the S3 bucket—any successfully transferred files from the previous job aren't retransferred from your Qumulo cluster.

Whenever Qumulo Core doesn't complete an operation successfully and returns an error from the API or CLI, the error field within the last_job field (that the replication_list_object_relationship_statuses command returns) contains a detailed failure message. For more troubleshooting details, see qumulo-replication.log on your Qumulo cluster.

Best Practices

We recommend the following best practices for working with Qumulo Shift-To for Amazon S3.

- Bucket Lifecycle Policy: To abort any incomplete uploads older than several days and
 ensure the automatic clean-up of any storage that incomplete parts of large objects (left
 by failed or interrupted replication operations) use, configure a bucket lifecycle policy. For
 more information, see Uploading and copying objects using multipart upload in the
 Amazon Simple Storage Service User Guide.
- VPC Endpoints: For best performance when using a Qumulo cluster in AWS, configure a VPC endpoint to S3. For on-premises Qumulo clusters, we recommend AWS Direct Connect or another high-bandwidth, low-latency connection to S3.
- Unique Artifacts: To avoid collisions between different data sets, specify a unique object folder or unique bucket for each replication relationship from a Qumulo cluster to S3.
- Object Versioning: To protect against unintended overwrites, enable object versioning. For more information, see Using versioning in S3 buckets in the Amazon Simple Storage Service User Guide.
- Completed Jobs: If you don't plan to use a Shift relationship to download updates from S3, delete the relationship to free up any storage associated with it.
- Concurrent Replication Relationships: To increase parallelism, especially across distinct datasets, use concurrent replication relationships to S3. To avoid having a large number of concurrent operations impact client I/O to the Qumulo cluster, limit the number of concurrent replication relationships. While there is no hard limit, we don't recommend creating more than 100 concurrent replication relationships on a cluster (including both Shift and Qumulo local replication relationships).

Restrictions

 Object-Locked Buckets: You can't use buckets configured with S3 Object Lock and a default retention period for Shift-To. If possible, either remove the default retention period and set retention periods explicitly on objects uploaded outside of Shift or use a different S3 bucket without S3 Object Lock enabled. For more information, see How S3 Object Lock works in the *Amazon Simple Storage Service User Guide*.

- File Size Limit: The size of an individual file can't exceed 5 TiB (this is the maximum object size that S3 supports). There is no limit on the total size of all your files.
- File Path Limit: The length of a file path must be shorter than 1,024 characters, including the configured object folder prefix, excluding the local directory path.
- Hard Links: Qumulo Core 3.2.3 (and higher) supports hard links, up to the maximum object size that S3 supports.
- Objects Under the Same Key: Unless an object contains Qumulo-specific hash metadata
 that matches a file, any object that exists under the same key that a new relationship
 replicates is overwritten. To retain older versions of overwritten objects, enable versioning
 for your S3 bucket. For more information, see Using versioning in S3 buckets in the
 Amazon Simple Storage Service User Guide.
- Object Checksums: All files replicated by using S3 server-side integrity verification (during upload) use a SHA256 checksum stored in the replicated object's metadata.
- S3-Compatible Object Stores: S3-compatible object stores aren't supported. Currently, Qumulo Shift-To supports replication only to Amazon S3.
- HTTP: HTTP isn't supported. All Qumulo connections are encrypted by using HTTPS and verify the S3 server's SSL certificate.
- Anonymous Access: Anonymous access isn't supported. You must use valid AWS credentials.
- Replication without Throttling: Replication provides no throttling and might use all available bandwidth. If necessary, use Quality of Service rules on your network.
- Amazon S3 Standard Storage Class: Qumulo Shift-To supports uploading only objects stored in the Amazon S3 Standard storage class. You can't download objects stored in the Amazon S3 Glacier or Deep Archive storage classes and any buckets that contain such objects cause a copy job to fail.
- Content-Type Metadata: Because all objects are stored in S3 using the default binary/
 octet-stream content type, they might be interpreted as binary data if you download
 them by using a browser. To attach content-type metadata to your objects, use the AWS
 Console.

Using Qumulo Shift-From to Copy Objects from Amazon S3

This section explains how to use Shift-From to copy objects from a folder in an Amazon Simple Storage Service (Amazon S3) bucket (cloud object store) to a directory in a Qumulo cluster and how to manage Shift relationships.

For more information about copying objects from Qumulo to S3, see Using Qumulo Shift-To for Amazon S3 to Copy Objects (page 120) on Qumulo Care.

O Note

From Qumulo Core 4.3.4, Shift-From estimates the work that a copy job performs.

Prerequisites

- · A Qumulo cluster with:
 - Qumulo Core 4.2.3 (or higher)
 - HTTPS connectivity to s3.<region>.amazonaws.com though one of the following means:
 - Public Internet
 - VPC endpoint
 - AWS Direct Connect

For more information, see AWS IP address ranges in the AWS General Reference.

- · Membership in a Qumulo role with the following privileges:
 - PRIVILEGE_REPLICATION_OBJECT_WRITE: This privilege is required to create a Shift relationship.
 - PRIVILEGE_REPLICATION_OBJECT_READ: This privilege is required to view the status of a Shift relationship.

1 Note

- For any changes to take effect, user accounts with newly assigned roles must log out and log back in (or their sessions must time out).
- Use special care when granting privileges to roles and users because certain privileges (such as replication-write privileges) can use system privileges to overwrite or move data to a location where a user has greater permissions.
 This can give a user access to all directories and files in a cluster regardless of any specific file and directory settings.
- · An existing bucket with contents in Amazon S3
- · AWS credentials (access key ID and secret access key) with the following permissions:
 - s3:GetObject
 - o s3:ListBucket

For more information, see Understanding and getting your AWS credentials in the AWS General Reference

Example IAM Policy

In the following example, the IAM policy gives permission to read from and write to the my-folder folder in the my-bucket. This policy can give users the minimal set of permissions required to run Shift-From jobs. (Shift-To jobs require a less-restrictive policy. For more information and an example, see Using Qumulo Shift-To for Amazon S3 to Copy Objects (page 120).)

How Shift-From Relationships Work

Qumulo Core performs the following steps when it creates a Shift-From relationship.

- 1. Verifies that the directory exists on the Qumulo cluster and that the specified S3 bucket exists, is accessible by using the specified credentials, and contains downloadable objects.
- 2. Creates the Shift-From relationship.
- 3. Starts a job by using one of the nodes in the Qumulo cluster.

O Note

If you perform multiple Shift operations, Qumulo Core uses multiple nodes.

- 4. Lists the contents of the S3 folder and downloads the objects to the specified directory on your Qumulo cluster.
- 5. Forms the full path of the file on the Qumulo custer by appending the path of the object (relative to the S3 folder) to the directory path on the Qumulo cluster.

For example, the following object is downloaded to /my-dir/my-project/file.text, where my-folder is the specified S3 folder and my-dir is the directory on your Qumulo cluster.

https://my-bucket.s3.us-west-2.amazonaws.com/my-folder/my-project/file.txt

O Note

This process doesn't encode or transform your data in any way. Shift-From attempts only to map every S3 object in the specified folder to a file on your Qumulo cluster.

6. Avoids redownloading an unchanged object in a subsequent job by tracking the information about an object and its replicated object.

1 Note

If you rename or move an object or local file between jobs, or if there are any metadata changes in S3 or Qumulo, the object is replicated again.

Storing and Reusing Relationships

The Shift-From relationship remains on the Qumulo cluster. You can monitor the completion status of a job, start new jobs for a relationship after the initial job finishes, and delete the relationship (when you no longer need the S3-folder-Qumulo-directory pair). To avoid

redownloading objects that a previous copy job downloaded, relationships take up approximately 100 bytes for each object. To free this storage, you can delete relationships that you no longer need.

If you repeatedly download from the same S3 folder, you can speed up the download process (and skip already downloaded files) by using the same relationship.

A new relationship for subsequent downloads doesn't share any tracking information with previous relationships associated with a directory and might recopy data that is already downloaded.

Using the Qumulo Core Web UI to Copy Files and Manage Relationships

This section describes how to use the Qumulo Core Web UI 4.2.5 (and higher) to copy files from Amazon S3 to a Qumulo cluster, review Shift relationship details, stop a running copy job, repeat a completed copy job, and delete a relationship.

To Copy Files from Amazon S3

- 1. Log in to Qumulo Core.
- 2. Click Cluster > Copy to/from S3.
- 3. On the Copy to/from S3 page, click Create Copy.
- 4. On the Create Copy to/from S3 page, click Local ← Remote and then enter the following:
 - a. The Directory Path on your cluster (/ by default)
 - b. The S3 Bucket Name
 - c. The Folder in your S3 bucket
 - d. The Region for your S3 bucket
 - e. Your AWS Region (/ by default)
 - f. Your AWS Access Key ID and Secret Access Key.
- 5. (Optional) For additional configuration, click Advanced S3 Server Settings.
- 6. Click Create Copy.
- 7. In the Create Copy from S3? dialog box, review the Shift relationship and then click Yes, Create.

The copy job begins and Qumulo Core estimates the work to be performed. When the estimation is complete, the Web UI displays a progress bar with a percentage for a relationship on the Replication Relationships page. The page also displays the estimated total work, the remaining bytes and files, and the estimated time to completion for a running copy job.

O Note

For work estimates, Shift-From jobs calculate the total number of files and bytes in a job's bucket prefix. This requires the job to use the ListObjectV2 S3 action once for every 5,000 objects (or 200 times for every 1 million objects).

To View Configuration Details and Status of Shift Relationships

- 1. Log in to Qumulo Core.
- 2. Click Cluster > Copy to/from S3.

The Copy to/from S3 page lists all existing Shift relationships.

3. To get more information about a specific Shift relationship, click : > View Details.

The Copy to/from S3 Details page displays the following information:

- · Throughput: average
- · Run Time
- · Data: total, transferred, and unchanged
- · Files: total, transferred, and unchanged

To Stop a Copy Job in Progress

- 1. Log in to Qumulo Core.
- 2. Click Cluster > Copy to/from S3.
- 3. To stop a copy job for a specific relationship, click : > Abort.
- 4. In the Abort copy from? dialog box, review the Shift relationship and then click Yes, Abort.

 The copy job stops.

To Repeat a Completed Copy Job

- 1. Log in to Qumulo Core.
- 2. Click Cluster > Copy to/from S3.
- 3. To stop a copy job for a specific relationship, click : > Copy Again.
- In the Copy again? dialog box, review the Shift relationship and then click Yes, Copy Again.
 The copy job repeats.

To Delete a Shift Relationship

- 1. Log in to Qumulo Core.
- 2. Click Cluster > Copy to/from S3.
- 3. To stop a copy job for a specific relationship, click : > Delete.

4. In the Delete copy from? dialog box, review the Shift relationship and then click Yes, Delete.

The copy job is deleted.

Using the Qumulo CLI to Copy Files and Manage Relationships

This section describes how to use the Qumulo CLI to copy files from Amazon S3 to a Qumulo cluster, review Shift relationship details, stop a running copy job, repeat a completed copy job, and delete a relationship.

Copying Files from Amazon S3

To copy files, use the replication_create_object_relationship command and specify the following:

- · Local directory path on Qumulo cluster
- · Copy direction (copy-from)
- · S3 object folder
- · S3 bucket
- · AWS region
- · AWS access key ID
- AWS secret access key

The following example shows how to create a relationship between the directory /my-dir/ on a Qumulo cluster and the S3 bucket my-bucket and folder /my-folder/ in the us-west-2 AWS region. The secret access key is associated with the access key ID.

```
qq replication_create_object_relationship \
    --local-directory-path /my-dir/ \
    --direction COPY_FROM_OBJECT \
    --object-folder /my-folder/ \
    --bucket my-bucket \
    --region us-west-2 \
    --access-key-id AKIAIOSFODNN7EXAMPLE \
    --secret-access-key wJalrXUtnFEMI/K7MDENG/bPxRfiCYEXAMPLEKEY
```

The CLI returns the details of the relationship in JSON format, for example:

```
"access_key_id": "ABC",
  "bucket": "my-bucket",
  "object_store_address": "s3.us-west-2.amazonaws.com",
  "id": "1c23b4ed-5c67-8f90-1e23-a4f5f6ceff78",
  "object_folder": "my-folder/",
  "port": 443,
  "ca_certificate": null,
  "region": "us-west-2",
  "local_directory_id": "3",
  "direction": "COPY_FROM_OBJECT",
}
```

Viewing Configuration Details and Status of Shift Relationships

- To view configuration details for all Shift relationships, use the replication_list_object_relationships command.
- To view configuration details for a specific relationship, use the
 replication_get_object_relationship command followed by the --id and the Shift
 relationship ID (GUID), for example:

```
qq replication_get_object_relationship --id 1c23b4ed-5c67-8f90-1e23-a4f5f6cef
f78
```

- To view the status of a specific relationship, use the
 replication_get_object_relationship_status command followed by the --id and the
 Shift relationship ID.
- To view the status of all relationships, use the replication_list_object_relationship_statuses command.

The CLI returns the details of all relationships in JSON format, for example:

```
[
 {
    "direction": "COPY_FROM_OBJECT",
    "access key id": "AKIAIOSFODNN7EXAMPLE",
    "bucket": "my-bucket",
    "object store address": "s3.us-west-2.amazonaws.com",
    "id": "1c23b4ed-5c67-8f90-1e23-a4f5f6ceff78",
    "object_folder": "my-folder/",
    "port": 443,
    "ca certificate": null,
    "region": "us-west-2",
    "local directory id": "3",
    "local directory path": "/my-dir/",
    "state": "REPLICATION_RUNNING",
    "current job": {
      "start time": "2020-04-06T17:56:29.659309904Z",
      "estimated end time": "2020-04-06T21:54:33.244095593Z",
      "job progress": {
        "bytes transferred": "178388608",
        "bytes unchanged": "0",
        "bytes_remaining": "21660032",
        "bytes total": "200048640",
        "files transferred": "17",
        "files unchanged": "0",
        "files remaining": "4",
        "files total": "21",
        "percent_complete": 89.0368314738253,
        "throughput current": "12330689",
        "throughput_overall": "12330689"
      }
    },
    "last job": null
  }
]
```

The state field shows the REPLICATION_RUNNING status and the current_job field shows the job's progress. When Qumulo Core copies files from S3, details for the most recently completed job become available in the last_job field, the state field changes to REPLICATION_NOT_RUNNING, and the current_job field reverts to null.

O Note

If you already ran a job for a relationship, it is possible for both the current_job and last_job fields to be non-null while you run a new job.

The bytes_total and files_total fields represent the total amount of data and number of files to be transferred by a Shift job. The bytes_remaining and files_remaining fields show the amount of data and number of files not yet transferred. The values of these four fields don't stabilize until the work estimation for the job is complete.

The percent_complete field displays the overall job progress and the estimated_end_time field displays the time at which the job is estimated to be complete. The values of these two fields are populated when the work estimation for the job is complete.

Shift-From performs a single task that estimates the amount of content to copy by listing all files and summing up their contents. Until this task is complete, the percent_complete
field is set to "None" and the estimated_end_time
field is set to "". To list the bucket prefix content in sets of 5,000 objects, this task uses the ListObjectV2
S3 action.

Stopping a Copy Job in Progress

To stop a copy job already in progress, use the replication_abort_object_relationship command followed by the --id and the Shift relationship ID.

Repeating a Completed Copy Job

To repeat a completed copy job, use the replication_start_object_relationship command followed by the --id and the Shift relationship ID.

This command begins a new job for the existing relationship and downloads any content that changed in the S3 bucket or on the Qumulo cluster since the time the previous job ran.

Deleting a Shift Relationship

After your copy job is complete, you can delete your Shift relationship. To do this, run the replication_delete_object_relationship command followed by the --id and the Shift relationship ID.

O Note

You can run this command only against a relationship that doesn't have any active jobs running.

This command removes the copy job's record, leaving locally stored objects unchanged. Any storage that the relationship used to track downloaded objects becomes available when you delete the relationship.

Troubleshooting Copy Job Issues

Any fatal errors that occur during a copy job cause the job to fail, leaving a partially copied set of files in the directory on your Qumulo cluster. However, to let you review the Shift relationship status any failure messages, the Shift relationship continues to exist. You can start a new job to complete the copying of objects from the S3 bucket—any successfully transferred files from the previous job aren't retransferred to your Qumulo cluster.

Whenever Qumulo Core doesn't complete an operation successfully and returns an error from the API or CLI, the error field within the last_job field (that the replication_list_object_relationship_statuses command returns) contains a detailed failure message. For more troubleshooting details, see qumulo-replication.log on your Qumulo cluster.

Best Practices

We recommend the following best practices for working with Qumulo Shift-From for Amazon S3.

• Inheritable Permissions: Because the system user creates the files that Shift-From for S3 copies, the system owns these files. By default, everyone is granted read permissions and administrators always have full access to the files.

O Note

To ensure that the copied files and subdirectories have the correct permissions, you must assign the necessary inheritable permissions to the root directory of the relationship *before* you create a Shift-From S3 relationship. To edit directory permissions, you can use the Windows Security Dialog or the qq fs_modify_acl command. For more information, see Qumulo File Permissions Overview on Qumulo Care.

- VPC Endpoints: For best performance when using a Qumulo cluster in AWS, configure a VPC endpoint to S3. For on-premises Qumulo clusters, we recommend AWS Direct Connect or another high-bandwidth, low-latency connection to S3.
- Repeated Synchronization: If you need to repeatedly synchronize an S3 folder with a Qumulo directory, we recommend reusing the same relationship. This lets you avoid repeated downloading of unchanged objects that already exist locally.
- Completed Jobs: If you don't plan to use a Shift relationship to download updates from S3, delete the relationship to free up any storage associated with it.
- Concurrent Replication Relationships: To increase parallelism, especially across distinct
 datasets, use concurrent replication relationships from S3. To avoid having a large number
 of concurrent operations impact client I/O to the Qumulo cluster, limit the number of
 concurrent replication relationships. While there is no hard limit, we don't recommend
 creating more than 100 concurrent replication relationships on a cluster (including both
 Shift and Qumulo local replication relationships).

Restrictions

- S3-Compatible Object Stores: S3-compatible object stores aren't supported. Currently, Qumulo Shift-From supports replication only from Amazon S3.
- HTTP: HTTP isn't supported. All Qumulo connections are encrypted by using HTTPS and verify the S3 server's SSL certificate.

- Anonymous Access: Anonymous access isn't supported. You must use valid AWS credentials.
- Replication without Throttling: Replication provides no throttling and might use all available bandwidth. If necessary, use Quality of Service rules on your network.
- Amazon S3 Standard Storage Class: Qumulo Shift-From supports downloading only
 objects stored in the Amazon S3 Standard storage class. You can't download objects
 stored in the Amazon S3 Glacier or Deep Archive storage classes and any buckets that
 contain such objects cause a copy job to fail.
- Disallowed Amazon S3 Paths in Qumulo Clusters: Certain allowed Amazon S3 paths can't be copied to Qumulo clusters and cause a copy job to fail. Disallowed paths contain:
 - A trailing slash (/) character (with non-zero object content length)
 - Consecutive slash (/) characters
 - Single and double period (, , . .) characters
 - The path component .snapshot
- Disallowed Conflicting Types: When content in an S3 bucket or Qumulo directory changes over time, a conflict related to type mismatches might arise, the Shift-from job fails, and an error message gives details about the conflict. For example, a conflict might occur when a remote object maps to a local file system directory entry which:
 - Is a regular file with two or more links
 - Isn't a regular file (for example, a directory or a special file)
- Disallowed Amazon S3 Path Configurations: Because of conflicting type requirements, Qumulo Core can't recreate certain allowed Amazon S3 path configurations on Qumulo clusters. For example, if an S3 bucket contains objects a/b/c and a/b, then path a/b must be both a file and directory on a Qumulo cluster. Because this isn't possible, this configuration causes a copy job to fail.
- Directories in Multiple Relationships: A directory on a Qumulo cluster for one Shift relationship can't overlap with a directory used for another Shift relationship, or with a remote directory for a Qumulo-to-Qumulo replication relationship. This causes the relationship creation to fail.
- Changes to S3 Folder During Copy Job: Currently, Shift-From assumes that the S3 folder remains unchanged throughout the copy job. Any changes (deleting, archiving, or modifying an object) during the copy job might cause a copy job to fail.
- Read-Only Local Directory: When the Shift-From copy job begins, the local directory on the Qumulo cluster becomes read-only. While no external clients can modify anything in the directory or its subdirectories, all content remains readable. When the copy job is complete, the directory reverts to its previous permissions.

• Partially Downloaded Files: If a copy job is interrupted or encounters a fatal error (that can't be resolved by retrying the operation), Qumulo Core attempts to delete partially downloaded files. Because this is a best-effort process, certain interruptions can prevent the cleanup of partially downloaded files.

File System Changes

How File System Change Notifications Work in Qumulo Core

This section describes how file system change notifications work in Qumulo Core and explains request filtering, recursion, and the three configuration modes for notification requests.

Qumulo Core can stream file system change notifications to a client whenever someone modifies a file or directory. The client can specify for which directories in the file system to receive notifications and what notification types Qumulo Core sends for these directories.

Qumulo Core supports two protocols for streaming file system change notifications. Both protocols provide roughly the same functionality.

- SMB: For more information, see Watching for File Attribute and Directory Changes by Using SMB2 CHANGE_NOTIFY (page 147)
- REST: For more information, see Watching for File Attribute and Directory Changes by Using REST (page 151).

The Qumulo Core notification system guarantees that:

- · The system never misses a notification.
- · The system sends notifications in real time.
- A client connected to any node in a cluster receives a notification which might originate from any node in the cluster.
- The system sends notifications in accurate chronological order. (For example, the system doesn't send a child_file_added event after a child_file_removed event.)
- In case a client can't keep up with the amount of events that the system emits, the cluster stops queuing events and produces an error the next time the client attempts to contact the cluster.

How Request Filtering Works

A client can request Qumulo Core to filter notifications. Although available filters differ between protocols, they work in a similar way.

☑ Tip

We recommend using a filter whenever you have an idea of the kind of events for which you want to receive notifications. Filtering notifications reduces back-end and front-end load and helps your client keep up with the data the cluster streams to it.

For example, you want to configure a client to receive notifications only about files being created but not deleted. In this scenario, you can make a notification request with the child_file_added filter for your protocol.

How Recursion Works

A Important

- Because recursive notification lets you monitor large regions of the file system tree without having to "walk" through these regions to look for changes, recursive notification is a powerful feature. Use this feature carefully: Watching too large a file tree can lead to the system sending too many notifications.
- In case a client can't keep up with the amount of events that the system emits, the cluster stops queuing events and produces an error the next time the client attempts to contact the cluster.
- When Qumulo Core sends too many messages, there can also be a slight performance impact for your cluster. For example, thousands of recursive watches at the file system root can have a measurable performance impact on a write-heavy IOPS workload.
- In the Windows implementation, recursive notifications over SMB require permissions only for the directory that the system watches. The system doesn't check permissions anywhere below the directory. Before you enable recursive notification, consider whether this policy is appropriate for your organization.

When you don't use recursion, the system sends notifications for changes that occur immediately within a target directory (changes to files that are children of the watched directory).

Consider the following example with recursion disabled:

- 1. You watch the /a/b directory.
- 2. You create the /a/b/f file.
- 3. You receive a notification that a user created the /a/b/f file.
- 4. You create the /a/b/c/f file.

In this example, you receive no additional notifications.

When you use recursion, the system sends notifications for changes that occur in the sub-tree below the target directory. Both protocols let you use recursion.

If you repeat the previous example scenario with recursion enabled, you receive two notifications:

· You receive a notification that a user created the /a/b/f file.

· You receive a notification that a user created the /a/b/c/f file.

Configuration Modes for Notification Requests

Qumulo Core has three global configuration modes that affect all recursive notification requests for both protocols.

- DISABLED ERROR: Recursive change notification requests return errors immediately.
- DISABLED_IGNORE: The system accepts recursive change notification requests but sends
 notifications only for the top directory that it watches. (The system behaves as if the user
 doesn't specify the recursive flag.)

Use this mode to improve compatibility with applications that request recursive behavior but don't depend on it.

A Important

For scenarios that require recursive behavior, this mode can cause an application to become unresponsive or exhibit other unexpected behavior.

• ENABLED: This is the default mode. This mode provides full support for recursive change notification requests. The system pushes notifications for all descendants of the watched directory to the watcher.

A Important

This mode can affect system performance. For example, thousands of recursive watches at the file system root can have a measurable performance impact on a write-heavy IOPS workload.

To select the configuration mode, use the /v1/file-system/settings/notify REST API or the qq fs_set_notify_settings command.

Supported Functionality

Functionality	Supported from Qumulo Core Version
Fully featured REST API for streaming file system notifications	6.0.2

Functionality	Supported from Qumulo Core Version
 Full support for SMB CHANGE_NOTIFY Parity with Windows Server Full support for recursion No limit on maximum directory size Configuration options for controlling recursive behavior 	6.0.1
 Improved compatibility with Windows applications Configuration options for controlling behavior of unsupported features 	5.3.3
 Support for all possible SMB CHANGE_NOTIFY types except for Alternate Data Streams (ADS) 	5.3.1
SMB2 CHANGE_NOTIFY support for adding and removing files and directories	5.3.0

Watching for File Attribute and Directory Changes by Using SMB2 CHANGE_NOTIFY

This section lists the completion filters that an SMB client can request and the corresponding actions that Qumulo Core returns for a matched change.

O Note

Qumulo provides REST access to the same change notification system (page 143) that backs SMB2 CHANGE_NOTIFY. Because the notification interface and types are easier to work with compared to SMB2, we recommend programming by using REST rather than SMB CHANGE NOTIFY.

Qumulo Core can watch for changes in file attributes and directory entries with a combination of SMB2 CHANGE_NOTIFY filters. Depending on the requested filter—and activity in the file system—an SMB client or an application remains current by receiving a variety of notifications.

Commonly, these requests help limit the amount of traffic required to keep a current cache of entries for an open directory. The requests also help operating system applications such as Windows Explorer and macOS Finder update automatically when changes take place. It is also possible to make requests programmatically. For more information about language bindings, see the Windows Protocol documentation, such as ReadDirectoryChangesW function (winbase.h) for Win32 and FileSystemWatcher Class for .NET.

Completion Filter Types

Each request uses a *completion filter* to specify the events to watch for. When events occur, the system batches them into a NOTIFY response that contains a list of FILE_ACTION items, each tagged with the names of changed entries. As long as the handle for the watched directory remains open, events queue up on the server, so that no events are lost between NOTIFY requests.

- · Watching for Name Changes: A name change can include four event types.
 - Renaming
 - Deleting
 - Moving into watched directory
 - Moving out of watched directory

The returned action specifies to your application whether an entry has been added, renamed, or removed.

 Watching for Metadata Changes: A metadata change can include six supported attribute types.

- File attributes
- File size
- Last-write time
- Last-access time
- Creation time
- Security (the permissions or access control list for the file or directory)

O Note

Qumulo doesn't support mutating extended attributes (EA). If the system requests only the FILE_NOTIFY_CHANGE_EA filter, no events propagate.

Completion Filters and Corresponding Actions

The following table shows the requested completion filters (grouped by the number of inode reads required to support them), the changes they watch for, and the actions that correspond to them.

Completion Filters	Actions	Description
The following filters watch for name changes (readdir-without-attrs). • FILE_NOTIFY_CHANGE_DIR_NAME • FILE_NOTIFY_CHANGE_FILE_NAME	 FILE_ACTION_ADDED FILE_ACTION_MODIFIED FILE_ACTION_REMOVED FILE_ACTION_RENAMED_NEW_NAME FILE_ACTION_RENAMED_OLD_NAME 	When Qumulo Core watches names, it notifies the client when there is an added, removed, or renamed file or directory in the watched directory. A rename event sends separate, consecutive events for OLD and NEW names, for example: [REMOVED, file_old_name], [ADDED, file_new_name]

Completion Filters	Actions	Description
The following filters watch for metadata changes (readdir-with-attrs). • FILE_NOTIFY_CHANGE_ATTRIBUTES • FILE_NOTIFY_CHANGE_CREATION • FILE_NOTIFY_CHANGE_SECURITY • FILE_NOTIFY_CHANGE_SIZE • FILE_NOTIFY_CHANGE_LAST_ACCESS • FILE_NOTIFY_CHANGE_LAST_WRITE	FILE_ACTION_MODIFIED	When one of the watched attributes changes for an entry of the watched directory and the filter is requested, the client receives a MODIFIED event. 1 Note In Microsoft terminology, attrib-
• Note Qumulo doesn't support mutating extended attributes (EA). If the system requests only the FILE_NOTIFY_CHANGE_EA filter, no events propagate.		utes are flags. For more information, see File Attributes in the Open Specification documentation.
The following filters watch for alternative data stream (ADS) changes (readdir-attrs-and-stream-names). • FILE_NOTIFY_CHANGE_STREAM_NAME • FILE_NOTIFY_CHANGE_STREAM_SIZE • FILE_NOTIFY_CHANGE_STREAM_WRITE	 FILE_ACTION_ADDED_STREAM FILE_ACTION_REMOVED_STREAM FILE_ACTION_MODIFIED_STREAM 	Consider the following example command. echo "data" > wat ched_dir/file0:st ream This command generates the following event. [ADDED_STREAM, file0:stream]

O Note

If you don't supply a filter, the SMB server accepts the request but doesn't send any notifications.

Re-Enumerating Changes after the STATUS_ENUM_DIR Error

If the SMB client can't keep up with the notification stream from the server, the server returns the STATUS ENUM DIR error code to the client and stops sending notifications.

In this scenario, the client must re-enumerate any changes that concern it directly, by opening and inspecting files, rather than relying on notifications. This scenario can happen because the request is too broad, for example, a recursive watch on the filesystem root.

Configuring Full Recursion for the WATCH_TREE Flag

By default, when a client supplies the WATCH_TREE flag, the server sends an error to the client immediately. In this scenario, you can configure your Qumulo cluster to support full recursion (page 143).

Watching for File Attribute and Directory Changes by Using REST

This section describes how to configure Qumulo Core and watch for file attribute and directory changes by using REST.

O Note

Qumulo provides REST access to the same change notification system (page 143) that backs SMB2 CHANGE_NOTIFY. Because the notification interface and types are easier to work with compared to SMB2, we recommend programming by using REST rather than SMB CHANGE NOTIFY.

Qumulo Core streams notifications to the client by using HTML server-sent events (SSE). For more information about the SSE syntax, see Server-sent events in the HTML Living Standard documentation.

- The comment syntax—any line that begins with a colon (:)—shows that the call has registered successfully for notifications and periodic keep-alive connections.
- The data syntax (data: payload) shows the event content.

Qumulo Core continues to stream events until the client closes the connection.

▲ Important

- Standard file system permissions apply to API requests for non-recursive watching:
 The system compares the authenticated user that makes the API request with the
 defined access control list (ACL) permissions for the file and grants or denies access
 access. The authenticated user must have permission to read a directory in order to
 request notifications for its changes.
- Because of the complexity of representing and enforcing a permissions model for an arbitrary subtree of the file system, recursive notification requests require an authenticated API user to have the DATA_ADMINISTRATOR privilege. This requirement remains true even if you configure your Qumulo cluster to ignore the recursive notification mode.

How SSE Event Payloads are Structured for Recursive Notification Requests

An SSE event payload is a JSON-encoded list of notification objects. The following is a format example of the SSE event payload.

- type: One of the possible notification types (page 0).
- path: The path to the file for which the notification occurred.

This path is relative to the watched directory.

- spine: A representation of the file path that uses Qumulo file IDs (rather than path components).
 - The first file ID in the spine is the oldest ancestor in the path.
 - The last file ID in the spine is the file for which the system sends the notification.
- stream_name: The name of an alternate data stream (ADS) for the file.

When this value is **null**, the notification is for the file's default stream. Otherwise, the notification is for the listed stream.

SSE Payload Notification Types

The following is a list of available notification types in SSE event payloads (page 0).

- The type field shows a single notification type.
- The filter field shows multiple notification types in comma-separated format.

Notification Type	Description
child_acl_changed	The ACL for the listed or directory has been modified.

Notification Type	Description	
child_atime_changed	The atime (access time) of the listed file or directory has been modified.	
	 • When a client modifies the atime field for a file directly, Qumulo Cores sends atime notifications for the file. • If you have enabled atime monitoring on your Qumulo Cluster, Qumulo Core sends atime notifications automatically. To configure atime monitoring, use the /v1/file-system/settings/atime REST API or the qq fs_set_atime_settings command. 	
child_btime_changed	The btime (creation time) of the listed file or directory has been modified.	
child_mtime_changed	The mtime (modification time) of the listed file or directory has been modified.	
child_data_written	Data has been written to the listed file.	
child_dir_added	The listed directory has been created.	
child_dir_removed	The listed directory has been removed.	
child_dir_moved_from	A directory has been moved from the listed location. ① Note The combination of the *_moved_to and *_moved_from notification type constitutes the renaming of the listed directory.	

Notification Type	Description
child_dir_moved_to	A directory has been moved to the listed location.
	⑤ Note The combination of the *_moved_to and *_moved_from notification type constitutes the renaming of the listed directory.
child_extra_attrs_changed	The additional attributes for the listed file or directory have been modified.
	The additional attributes are Windows-specific extra file attributes which include HIDDEN, READ_ONLY, and so on. For more information, see File Attributes in the Microsoft Open Attributes documentation.
child_file_added	The listed file has been created.
child_file_removed	The listed file has been removed.
child_file_moved_from	A file has been moved from the listed location.
	♠ Note The combination of the *_moved_to and *_moved_from notification type constitutes the renaming of the listed file.
child_file_moved_to	A file has been moved from the listed location.
	① Note The combination of the *_moved_to and *_moved_from notification type constitutes the renaming of the listed file.
child_group_changed	The group for the listed file or directory has been changed.

Notification Type	Description
child_owner_changed	The owner for the listed file or directory has been changed.
child_size_changed	The size of the listed file has been changed.
child_stream_added	The listed alternate data stream (ADS) has been added to the listed file or directory.
child_stream_data_written	Data has been written to the listed ADS.
child_stream_moved_from	The listed ADS has been moved to the listed file or directory.
child_stream_moved_to	The listed ADS has been moved from the listed file or directory.
child_stream_removed	The listed ADS has been removed from the listed file or directory.
child_stream_size_changed	The size of the listed ADS for the listed file or directory has been changed.
self_removed	The directory from which then system streams notifications has been removed from the file system.
	① Note No notifications follow a self_removed notification.

Streaming Change Notifications by Using the qq CLI

Use the qq fs_notify command and specify the path to a directory. For example:

qq fs_notify --path /my/directory

In this example, Qumulo Core streams all notification types (page 0) for files immediately under the /my/directory directory.

To terminate the stream, send a **SIGQUIT** signal.

Streaming Change Notifications by Using the REST API

Make a **GET** request to the REST endpoint in the following format:

/v1/files/<ref>/notify&filter=<filter>&recursive>

In the following example:

- · ref : An absolute path or a numeric file ID for the directory to watch.
- filter: A comma-separated list of notification types (page 0).
- · recursive: When set to true, enables recursive change notifications.

/v1/files/my/directory/notify&filter=child_file_added,child_dir_removed&recursive=tr
ue

File System Protocols

Enabling and Using NFSv4.1 on a Qumulo Cluster

This section explains how to configure your cluster for a supported export configuration and enable or disable NFSv4.1 on your cluster.

For more information about NFSv4.1 and file access permissions, see Managing File Access Permissions by Using NFSv4.1 Access Control Lists (ACLs) (page 164).

▲ Important

- Currently, Qumulo Core 4.3.0 (and higher) supports only NFSv4.1. Mounting with version 4.0 or 4.2 isn't supported.
- The NFSv4.1 protocol requires clients to provide the server with globally unique identifiers. By default, the NFSv4.1 client for Linux uses the machine's hostname as co_ownerid. Because the NFSv4.1 protocol requires a unique identifier for every client, an unpredictable failure can occur if two clients have the same hostname. To configure unique identification for your NFS clients, set the nfs4_unique_id value for them. For more information, see The nfs4_unique_id parameter in the Linux kernel user's and administrator's guide.

Configuring and Using Exports for NFSv4.1

Qumulo's NFS exports can present a view of your cluster over NFS that might differ from the contents of the underlying file system. You can mark NFS exports as read-only, restricted (to allow access only from certain IP adresses), or configure specific user mappings. For more information, see Create an NFS Export on Qumulo Care.

While NFSv3 and NFSv4.1 share each cluster's NFS export configuration, exports behave differently when you access them by using NFSv4.1. This section explains these differences and the new requirements for export configurations with NFSv4.1.

Differences Between NFSv3 and NFSv4.1 Exports

In the following example, a Qumulo cluster has the following export configuration.

Export Name	File System Path	Read-Only
/home	/home	No
/files	/home/admin/files	No
/read_only/home	/home	Yes

Export Name	File System Path	Read-Only
/read_only/files	/home/admin/files	Yes

NFSv3 lets you mount one of these exports by specifying the full export name, for example:

```
mount -o nfsvers=3 cluster.qumulo.com:/read_only/home /mnt/cluster/home
```

This command gives read-only access to the home directory on the cluster by using the path /mnt/cluster/home. However, the following command fails with the <a href="https://no.ncm.ncm.no.ncm.ncm.no.ncm.

```
mount -o nfsvers=3 cluster.qumulo.com:/read_only /mnt/cluster/read_only
```

NFSv4.1 still lets you mount exports by specifying the full export name. However, NFSv4.1 also supports navigating *above* exports, as if they are part of the file system. The following command succeeds.

```
mount -o nfsvers=4.1 cluster.qumulo.com:/read_only /mnt/cluster/read_only
```

At the mount, the exports under <code>/read_only</code> are visible: <code>/mnt/cluster/read_only</code> displays virtual directories named <code>files/</code> and <code>home/</code> with the contents of the corresponding directories in the file system, for example:

```
/mnt/cluster/read_only/
|--- files/<file system contents>
|--- home/
|----- admin/files/<file system contents>
|----- <other file system contents>
```

This presentation of exports lets you view existing exports by using the file system's own interface. It also lets you view new exports as soon as someone creates or modifies them without remounting.

Preparing Export Configurations for NFSv4.1

Qumulo's implementation of NFSv4.1 distinguishes between navigating *above* exports and *inside* an export. To avoid confusion between paths that refer to a virtual directory above an export or a real file system directory inside an export, no export name can be a prefix of another export name when NFSv4.1 is enabled.

In the following example, a Qumulo cluster has the following export configuration.

Export Name	File System Path
/	/
/admin	/home/admin

Because / is a prefix of /admin, you can't enable NFSv4.1 with this export configuration. This restriction prevents the situation where the path /admin can refer to both the export of /home/admin or the actual file system path /admin.

To prepare this configuration for NFSv4.1, you can do one of the following:

- Delete the / export and use NFSv4.1 presentation of exports when mounting /.
- · Delete the /admin export.
- Give the / export a name that doesn't use other exports as a prefix, for example:

Export Name	File System Path
/root	/
/admin	/home/admin

Visibility of IP-Address-Restricted Exports

O Note

The names of exports are public to all NFSv4.1 clients, regardless of IP address restrictions. You can't disable this behavior.

NFSv4.1 respects IP address restrictions on exports: Only clients with allowed IP addresses can access the contents of an export. However, clients without access to an export can still view the export as a directory when they traverse *above* exports. The restrictions apply only when a client attempts to access the contents of the export.

32-Bit Sanitization

· In NFSv3, you can configure specific exports to return 32-bit sanitized data for individual

fields. NFSv3 converts any data larger than 32 bits in configured fields to 32-bit data and returns the data. For example, it can sanitize file size to 32-bit format. This truncates the field to max uint32 whenever the NFSv3 server returns the attribute.

 NFSv4.1 doesn't support 32-bit sanitization and ignores any sanitizations configured for an export.

Enabling NFSv4.1 on a Qumulo Cluster

O Note

Currently, you can enable NFSv4.1 only by using the qq CLI.

You can enable NFSv4.1 on your Qumulo cluster by using a single cluster-wide configuration command, for example:

qq nfs_modify_settings --enable-v4

When you enable NFSv4.1, all NFS exports are accessible through NFSv3 and NFSv4.1.

Specifying the NFS Mount Option

Typically, NFS clients find and use the highest version of the protocol that both the client and server support. For example, the following command mounts by using NFSv4.1 (if it is enabled) or by using NFSv3 otherwise.

mount -t nfs your.qumulo.cluster:/mount path/path/to/mountpoint

Because Qumulo's NFSv4.1 implementation currently doesn't have full feature parity with NFSv3, you must provide the nfsvers=3 option for any mounts that require features (such as snapshot access) that only NFSv3 supports, for example:

mount -t nfs -o nfsvers=3 your.qumulo.cluster:/mount_path /path/to/mountpoint

Note

We recommend specifying the nfsvers=4 or nfsvers=4.1 option for any mounts that use NFSv4.1.

Checking Whether NFSv4.1 is enabled

To check whether NFSv4.1 is enabled on your cluster, use the following qq CLI command:

qq nfs_get_settings

Disabling NFSv4.1 on a Qumulo Cluster

A Important

Disabling NFSv4.1 makes any NFSv4.1 mounts unusable immediately. We recommend switching any NFSv4.1 mounts to NFSv3 before disabling NFSv4.1.

To disable NFSv4.1 on an entire Qumulo cluster, use the following qq CLI command:

qq nfs modify settings --disable-v4

Configuring Floating IPs for Nodes

Currently, each Qumulo node is limited to 1,000 clients connected through NFSv4.1 simultaneously. To account for nodes going down, we recommend balancing the number of client connections across your nodes by configuring a sufficient number of floating IP addresses for each node. This prevents a node failover event from overloading the nodes to which the clients might fail over.

For example, if you configure only one IP address for each node, on a cluster with 600 clients for each node, a single node failure might overload one of the remaining nodes, preventing 200 clients from connecting. If you assign multiple floating IP addresses to each node, the clients' connections are distributed across multiple nodes.

Listing NFSv4.1 Byte-Range Locks

Rather than lock an entire file, byte-range locking lets you lock specific portions of a file or an entire file in use. This feature is available in Qumulo Core 5.1.3 (and higher). It doesn't require client mount configuration.

The NFSv4.1 implementation in Qumulo Core has a non-configurable lease of one minute. During each lease period, clients send a heartbeat to your Qumulo cluster. The cluster uses this heartbeat to detect lost client connections and to revoke the client leases. When the cluster revokes a lease, it releases any byte-range locks and makes them available to other clients.

A Important

- NFSv4.1 byte-range locks are interoperable with NLM (NFSv3) byte-range locks.
 NFSv4.1 clients view and respect locks that NFSv3 clients hold (the opposite is also true).
- NFSv4.1 and NLM locks aren't interoperable with SMB locks.

To list NFSv4.1 byte-range locks in your cluster, use the following qq CLI command:

qq fs_list_locks --protocol nfs4 --lock-type byte-range

A Note

- Currently, Qumulo Core doesn't support revoking NFSv4.1 byte-range locks by using the CLI.
- The time to acquire or release a lock scales linearly with the number of locks that the system already holds on a specific file. If a file has a very large number of locks, system performance can degrade.

Supported and Unsupported Features in Qumulo's Implementation of NFSv4.1

Qumulo's implementation of NFSv4.1 currently supports:

- Authentication with Kerberos (page 173)
- · General file system access (reading, writing, and navigating files)
- · Unstable writes
- Full use of the NFS exports configuration shared with NFSv3
- · Navigation in the pseudo-file system above your exports
- NFSv3-style AUTH SYS authentication (also known as AUTH UNIX)
- Fine-grained control over file permissions by using access control lists (ACLs)
- File locking (for example, by using the fcntl command)
- · Snapshots through NFSv4.1 (Qumulo Core 5.2.4 and higher)
- · Quotas through NFSv4.1 (Qumulo Core 5.2.5.1 and higher)

Qumulo Core doesn't currently support the following NFSv4.1 features:

· Delegations

Managing File Access Permissions by Using NFSv4.1 Access Control Lists (ACLs)

This section explains how to use Qumulo Core's implementation of NFSv4.1 with access control lists (ACLs) to manage access permissions for files.

The Qumulo Core implementation supports using AUTH_SYS credentials (also known as AUTH_UNIX), AUTH_NONE (which acts as AUTH_SYS but maps incoming UIDs and GIDs to nobody), and AUTH_KRB5, AUTH_KRB5P, or AUTH_KRB5I credentials. You can use the CLI tools in the nfs-acltools Linux package to allow or deny various operations.

For more information about NFSv4.1, see Enabling and Using NFSv4.1 on a Qumulo Cluster (page 157).

Using the NFSv4.1 CLI Commands to Manage ACLs

In most Linux distributions, the nfs-acl-tools package contains the NFSv4.1 commands that let you manage ACLs for files.

Showing the ACL of a File

To show the ACL of a file, use the nfs4_getfacl command. In the following example, we create
the file my-file and then show the ACL for it.

```
$ touch /mnt/qumulo/my-file
$ nfs4_getfacl /mnt/qumulo/my-file
A::userl@domain.example.com:rwatTnNcy
A:g:groupl@domain.example.com:rwatTnNcy
A::EVERYONE@:rtncy
```

The entries in the ACL have four parts separated by colons (:). For more information, see the nfs4 acl in the Linux documentation.

The ACL in this example corresponds to 664 mode: The owner (user1) and group (group1) of the file are allowed to read and write, while others (EVERYONE@) are allowed to only read. To check the current mode, use the stat command, for example:

```
$ stat -c %a /mnt/qumulo/my-file
664
```

Editing the ACL of a File

To edit the ACL of a file (by using the text editor specified in the \$EDITOR environment variable), use the nfs4_editfacl (or nfs4_setfacl -e) command. For more information, see the nfs4_editfacl and nfs4_setfacl in the Linux documentation.

Setting the ACL of a File

To set the ACL of a file, you can use one of the following commands:

- · Add a Single ACE: nfs4 setfacl -a <ace>
- Set an Entire ACL: nfs4 setfacl -s <acl>

Configuring Access Control Entries (ACEs) and Trustee Representation

O Note

The following guidance applies to all nfs4_acl scenarios, including getting, editing, and setting the ACL.

There are four fields in the nfs4_acl syntax, separated by colons (:):

- · The ACE type
- · Additional ACE flags
- · The trustee to which the ACE applies
- · The access types to which the ACE applies

ACE Type

In the example of the file ACL (page 164), all three ACEs are set to A (allow).

O Note

Qumulo Core supports only A and D ACEs.

- · A: Allow
- · D: Deny
- · U: Audit
- · L: Alarm

Additional ACE Flags

In the example of the file ACL (page 164), the second ACE has the flag g that shows that the ID in the following part represents a *group* (rather than a user).

Note

Qumulo Core doesn't support The S and F flags.

The Trustee to Which the ACE Applies

You can use the following trustee representation formats.

A Important

- Be careful when you copy *local users and groups* across different Qumulo clusters manually. Aside from UIDs and GIDs, local users and groups are the only identity types in this table that aren't globally unique (because a user or group name represents them). If the destination cluster interprets the named user or group differently, the permissions you set might be unexpected.
- · This consideration doesn't apply to replication copies of local user or group trustees.

Trustee Representation	Example	Description
<user>@<domain></domain></user>	userl@domain.example.com	A Kerberos principal that represents a user in the domain to which a Qumulo cluster is joined. You can use this format regardless of client mount security, but only when the cluster is joined to AD. For this trustee in the ACE, the system stores the corresponding AD SID for this user principal on disk. For more information about configuring your clients and Qumulo cluster for Kerberos, see the Using NFSv4.1 with Kerberos in Qumulo Core (page 173).

Trustee Representation	Example	Description
<group>@<domain></domain></group>	groupl@domain.example.com	A Kerberos principal that represents a group in the domain to which that a Qumulo cluster is joined. You can use this format regardless of client mount security, but only when the cluster is joined to AD. The group flag isn't necessary to show that this is a group. For this trustee in this ACE, the system stores the corresponding AD SID for this group principal on disk. For more information about configuring your clients and Qumulo cluster for Kerberos, see Using NFSv4.1 with Kerberos in Qumulo Core (page 173).
<s-r-x-y1-y2-yn-1-yn></s-r-x-y1-y2-yn-1-yn>	S-1-5-32-544	A raw SID. For more information, see Security Identifiers in the Microsoft documentation. To store a SID on disk for this trustee, you can use this format in place of a Kerberos principal. An AD SID must be a user or a group, but can't be both. However, the group flag isn't necessary for showing whether the SID represents a user or group. This can be useful if you have SIDs in a foreign domain (that is, a domain that the cluster isn't joined to). You can use this representation when the cluster isn't joined to a domain at all. When you retrieve an ACL by using nfs4_getfacl, the presentation for joined domain SIDs is <group>@<domain> and the presentation for foreign SIDs is <s-r-x-y1-y2-yn-1-yn>.</s-r-x-y1-y2-yn-1-yn></domain></group>
<numeric_uid></numeric_uid>	1234	A numerical UID for an AUTH_SYS user. For this trustee in the ACE, the system stores this UID on disk.

Trustee Representation	Example	Description
<numeric_gid></numeric_gid>	5678	A numerical GID for an AUTH_SYS user. To avoid having the group interpreted as a user, you must specify the group flag (page 165). For this trustee in the ACE, the system stores the GID on disk.
qumulo_local/ <username></username>	qumulo_local/localuser1	A user local to a Qumulo cluster (that is, a user that created by using Qumulo Core Web UI or the qq CLI. For the trustee in this ACE, the system stores this user as a local user.
qumulo_local/ <groupname></groupname>	qumulo_local/localgroup1	A group local to a Qumulo cluster (that is, a group created by using the Qumulo Core Web UI or the qq CLI. Because local Qumulo users and groups can't share a name, the group flag isn't necessary to show this is a group. For the trustee in this ACE, the system stores this group as a local group, on disk.
EVERYONE@	_	Any user of the file system.
GR0UP@	_	The group owner of a file.
OWNER@	_	The owner of a file.

You you can use all trusteee representations interchangeably, even within a single ACL. For example, the following ACL is possible for a file:

\$ nfs4_getfacl /mnt/qumulo/my-file
A::user1@domain.example.com:rwatTnNcy
A:g:group1@domain.example.com:rwatTnNcy

A::1234:rwatTnNcy
A:g:5678:rwatTnNcy
A::S-1-5-8-9:rwatTnNcy
A:g:S-1-5-32-544:rwatTnNcy

A::qumulo_local/localuser1:rwatTnNcy
A:g:qumulo local/localgroup1:rwatTnNcy

A::EVERYONE@:rtncy</code>

The Access Types to Which the ACE Applies

For example:

r: Read

• t : Read attributes

• w: Write

The nfs4_setfacl command also lets you use the following shorthand:

· R: Generic read

· W: Generic write

· X: Execute permissions

Managing NFSv4.1 Permissions with ACLs and POSIX-Style Modes

You can manage NFSv4.1 access permissions by using ACLs, POSIX-style modes, or a combination of both.

- If you set an ACL on a file and then also set a mode on it, the restrictions that the mode expresses also apply to the ACL. These restrictions change or remove ACEs that apply to the owner, group, or other users.
- If you use the OWNER@ or GROUP@ identifiers in an ACL that allows read, write, or execute permissions, the identifiers appear in the owner or group bits of the mode when you read the file's mode

Note

Because the EVERYONE@ identifier includes the owner and group of a file and the other bits of a mode don't apply to the owner or group, the permissions you grant to the EVERYONE@ identifier are more broad than a mode's other bits.

Using NFSv4.1 ACLs with SMB Access Control

NFSv4.1 ACLs are interoperable with SMB access controls. You can write and read by using both protocols. When you edit over NFS, the system represents SMB SIDs Kerberos principals.

Changing File Owners

When you change the owner of a file, the ACEs that refer to the owner change to the new owner, for example:

\$ nfs4_getfacl /mnt/qumulo/my-file
A::userl@domain.example.com:rwatTnNcy
A:g:groupl@domain.example.com:rwatTnNcy
A::EVERYONE@:rtncy

\$ sudo chown user2 /mnt/qumulo/my_file

\$ nfs4_getfacl /mnt/qumulo/my-file
A::user2@domain.example.com:rwatTnNcy
A:g:group1@domain.example.com:rwatTnNcy
A::EVERYONE@:rtncy

Using Equivalent NFSv4.1 and Qumulo ACL Commands

The syntax for the nfs4_setfacl command is <type>:<flags>:<principal>:<permissions> , for
example A:fd:GROUP@:rwaDdxtTnNcCoy . You can use equivalent NFS (nfs4_setfacl) and Qumulo
(qq fs modify acl) CLI commands to set ACL permissions.

The following tables compare elements of NFS and Qumulo ACL permissions.

NFSv4.1 ACL Type	Qumulo ACL Type
Α	Allowed
D	Denied

NFSv4.1 ACL Flag	Qumulo ACL Flag
d	Container inherit
f	Object inherit

NFSv4.1 Rights	Qumulo Rights
a	Extend file
С	Read ACL
С	Write ACL
d	Delete
n	Read EA
0	Take Ownership
r	Read contents
R	Read, Synchronize
t	Read attr
Т	Write attr
W	Write data
W	Read ACL, Read attr, Synchronize, Write ACL, Write file
Х	Execute/Traverse
X	Execute/Traverse, Read ACL, Read attr, Synchronize
у	Synchronize

The following table gives examples of permissions and equivalent NFS and Qumulo CLI commands.

Permissions	NFSv4.1 Command	Qumulo Command
Add Read Permission to File	<pre>nfs4_setfacl -a "A::OWN- ER@:R" myfile.ext</pre>	<pre>qq fs_modify_aclpath /myfile.ext add_entry -y Allowed -t "File Owner" -r Read</pre>
Add Read and Execute Permissions to File	<pre>nfs4_setfacl -a "A::EVERY- ONE@:rtRX" myfile.ext</pre>	<pre>qq fs_modify_aclpath /myfile.ext add_entry -y Allowed -t "EVERYONE" -r Execute/Traverse, Read</pre>

Permissions	NFSv4.1 Command	Qumulo Command
Add Read, Write, and Execute Per- missions to File	<pre>nfs4_setfacl -a "A::GROUP@:rtwRWX" my- file.ext</pre>	<pre>qq fs_modify_aclpath /myfile.ext add_entry -y Allowed -t "File Group Owner" -r Execute/Traverse, Read, Write ACL, Write file</pre>
Add Full Access to File	<pre>nfs4_setfacl -a "A::GROUP@:rtwRWX" my- file.ext</pre>	<pre>qq fs_modify_aclpath /myfile.ext add_entry -y Allowed -t "File Group Owner" -r Execute/Traverse, Read, Write ACL, Write file</pre>
Remove Write and Execute Permis- sion to File	nfs4_setfacl -a "D::OWN- ER@:wx" myfile.ext	<pre>qq fs_modify_aclpath /myfile.ext add_entry -y Denied -t "File Owner" -r Execute/Traverse, Write data</pre>
Add Full Access to Group File and Di- rectory Inheri- tances to Directory	<pre>nfs4_setfacl -a "A:fd:GROUP@:rwaDdxtTnNcCoy" mydirectory</pre>	<pre>qq fs_modify_aclpath /mydirecto- ry add_entry -y Allowed -t "File Group Owner" -r All -f 'Container in- herit' 'Object inherit'</pre>

NFSv4.1 with Kerberos

How NFSv4.1 Works with Kerberos in Qumulo Core

This section provides an overview of how NFSv4.1 works with Kerberos in Qumulo Core.

Kerberos is a network authentication protocol that works by using a three-way trust between a key distribution center (KDC), a service server (for example, NFSv4.1 on Qumulo Core), and a client system (for example, a Linux system). This section of the Qumulo Administrator Guide explains how to configure and use the three entities involved in the trust and provides troubleshooting directions. For more information, see Kerberos on Wikipedia and the MIT Kerberos documentation.

Active Directory (AD) simplifies Kerberos requirements by providing a globally unique security identifier for every user and group (SID) and a KDC implementation with a ticket-granting service (TGS) and an authentication service (AS).

Configuring Kerberos for Qumulo Core

Qumulo Core 5.1.5 (and higher) supports Kerberos for authenticating AD users over NFSv4.1. The following is an overview of the Kerberos configuration process following the configuration of your AD domain.

- 1. Join your Qumulo cluster to your AD domain.
- 2. Join Linux systems to your AD domain.
- 3. Log in to a Linux system and mount the Qumulo cluster by using the -o sec=krb5 mount option.

Known Kerberos Limitations for Qumulo Core

Qumulo Core supports only the following features:

- · NFSv4.1
- · Linux clients
- AES-128 and AES-256 encryption algorithms—for more information, see Network security:
 Configure encryption types allowed for Kerberos in the Microsoft documentation
- Microsoft Windows Active Directory (Windows Server 2008 and higher)

Prerequisites for Joining a Qumulo Cluster to Active Directory

This section describes the prerequisites for joining a Qumulo Cluster to Active Directory for using NFSv4.1 with Kerberos.

For more information, see Join Your Qumulo Cluster to Active Directory on Qumulo Care.

Using Active Directory (AD) for POSIX Attributes (RFC2307)

While using AD for POSIX attributes is optional, it helps avoid issues with Linux ID mapping. We recommend enabling RFC 2307 to match your client's functionality.

- Enabling RFC 2307 might simplify AUTH_SYS -based Linux clients that access the cluster by using known UIDs and GIDs. In this way, the cluster can map the UIDs and GIDs to the user or group objects on the AD server and enforce the appropriate permissions.
- If you configure sssd on Kerberos-mounted Linux clients for mapping by SID, disabling RFC 2307 can help avoid ascribing special meaning to randomly assigned Linux UIDs and GIDs.

Specifying the Base Distinguished Name (Base DN)

Qumulo uses LDAP to query the AD domain for users and groups. For this functionality, a Base DN must cover any identities intended for use with Kerberos. For example, if multiple organizational units (OUs) contain users, you must include them all in the Base DN (separated with semicolons).

Alternatively, a parent container can hold all nested containers of interest. It is possible to set a top-level domain (TLD) as the Base DN (however, this can cause queries to perform poorly in certain scenarios). We recommend using as specific a Base DN as possible. If you don't configure the Base DN correctly, Linux clients might present permissions such as nobody or 65534.

In the following example, there is an OU with the AD domain my.example.com. The TLD Base DN for this domain is as follows.

DC=my, DC=example, DC=com

If a Users container holds users and a Computers container holds machine accounts, you can set the Base DN as follows.

CN=Users, DC=my, DC=example, DC=com; CN=Computers, DC=stuff, DC=example, DC=com

O Note

This example is a very common configuration for user and computer objects in AD.

Using the Active Directory Domain Controller as the NTP Server

Kerberos is very sensitive to clock skew. It is important for all systems involved in a Kerberos relationship—the KDC, your Qumulo cluster, and any Linux clients—to have as little clock skew as possible. We recommend using the same NTP server for all three components.

- You can use your AD domain controller as an NTP server. In the Web UI, on the Active Directory page, for Use Active Directory as your primary time server, click Yes.
- To configure any other NTP server in the Web UI, click Cluster > Date & Time.

Configuring Active Directory for Use With Kerberos

This section describes the Active Directory Domain Controller (DC) configuration changes necessary for enabling NFSv4.1 with Kerberos.

Configuring DNS in Active Directory

Kerberos relies on DNS to identify machines involved in authentication. NFS clients and servers require DNS A records for forward-DNS lookups and PTR records for reverse-DNS lookups.

You can use a variety of DNS implementations with Kerberos. In some cases, for example, it might be convenient to use the DNS server that the AD DC provides. For this reason, this section discusses DNS configuration in general terms.

Modifying the Default DNS Configuration

By default, the Qumulo domain-join operation creates a machine account on the domain in the organizational unit (OU)—that you specify during the join process—automatically. This machine account represents all nodes in the cluster, not a single machine.

By default, this machine account has a single, automatically created DNS A record that refers to the node on which the system performs the domain-join operation. This DNS record exists on the AD DC used for the domain-join operation and the record refers to a single, public IP address for the node.

The default DNS configuration is generally not useful without additional modifications because:

- It applies to the DNS server for the DC: If the environment doesn't use this DNS server, you must create the entry on the DNS server manually.
- It creates only a DNS A (forward) record: You must create the PTR record (a reverse record that maps an IP address to a hostname) manually. This can require creating a reverse zone for the subnet and then adding the specific PTR record to the zone.
- We don't recommend assigning a single IP address to an entire cluster: In such a configuration, any client that mounts the cluster points at the same node.

Configuring DNS for Distributing Workflows Across Nodes

The Qumulo distributed file system works best when you spread the workload evenly across multiple nodes. We recommend configuring round-robin DNS in Active Directory.

This approach provides a list of IP addresses which refer to different nodes in the cluster. Successive DNS queries for the single cluster hostname return different IP addresses. From the perspective of Kerberos, all nodes that comprise a Qumulo cluster act as one host and have the same Kerberos key table. In this way, the Kerberos experience is the same regardless of the selected node.

Unless you need direct access to a specific node through a DNS fully qualified domain name (FQDN), it isn't necessary to use individual DNS A records for each node in the cluster (for example, qumulo1.example.com, qumulo2.example.com, qumulo3.example.com, and so on). Instead, we recommend creating a DNS A record for the cluster and then duplicating this A record for each IP address in the cluster (for example, qumulo.example.com \rightarrow 203.0.113.0, qumulo.example.com \rightarrow 203.0.113.1, and so on).

To Configure Round-Robin DNS

- 1. Join your Qumulo cluster to AD (page 174).
- 2. Find the DNS entry for the cluster on the DNS server.
 - Unless you renamed the cluster after joining it to AD, this entry is generally the cluster's name. To find the machine account name in the Web UI, click Cluster > Active Directory and write down the name under Machine Account.
- Update the list of IP addresses for this host record. Include the IP addresses for all nodes.
 To find the IP addresses in the Web UI, click Cluster > Network Configuration.
- 4. Configure the DNS resolver to point to the DNS server.

To find the IP addresses, look up the hostname for the DC. For example:

```
nslookup stuff.example.com
```

5. Confirm that successive ping <cluster_name> requests connect to a different IP address
every time.

Configuring the Service Principal Name (SPN) for NFS

The SPN is a string that identifies the Kerberos services that a particular host provides. We recommend configuring the Qumulo cluster to provide the NFS service. When you configure the SPN, clients can enumerate the cluster and the NFS service as part of a service-ticket-granting request.

To Configure the SPN for NFS by Using the Windows Server Attribute Editor



To maximize compatibility with Linux, we recommend formatting SPN entries in lowercase.

- 1. Use RDP to log in to the DC for your AD domain.
- 2. Open Active Directory Users and Computers.
- 3. Find the machine account for your Qumulo cluster.

To find the machine account name in the Web UI, click Cluster > Active Directory and write down the name under Machine Account.

- 4. Right-click the account and then click Properties > Attribute Editor.
- 5. On the Attribute Editor tab, find the servicePrincipalName attribute and edit its value to include a new SPN in the nfs/<machine_account>.<domain_fqdn> format, for example:

nfs/<qumulo-cluster>.ad.eng.example.com

☑ Tip

You can use the other, automatically generated entries as syntax examples.

To Configure the SPN for NFS by Using the Windows Server Command Prompt

O Note

- To maximize compatibility with Linux, we recommend formatting SPN entries in lowercase.
- The SPN formatting in the following example is generally sufficient for Linux service ticket requests. However, depending on your environment and client configuration, additional entries might be necessary.
- 1. Open a command prompt with administrative privileges.
- 2. Use RDP or SSH to connect to your AD domain.
- 3. Run the setspn command with the machine account (in this example, <qumulo-cluster>) followed by a period (.) and the FQDN (in this example, ad.eng.example.com). For example:

setspn -s nfs/<qumulo-cluster>.ad.eng.example.com

4. Confirm the configuration by using the **setspn** command with the machine account name. For example:

setspn <qumulo-cluster>

To Troubleshoot Your SPN Configuration

If your SPN is configured incorrectly, a client is likely to display the following error:

mount.nfs: access denied by server while mounting <qumulo-cluster>.ad.eng.qumulo.co
m:/

- 1. Take a client-side packet capture and find the logs for the client and AD Kerberos.
- 2. Search the logs for the S PRINCIPAL UNKNOWN error.
- 3. Add the required client parameters to the SPN configuration.

Configuring SPN with DNS

For Kerberos authentication to work correctly, SPN entries must correspond to DNS A records exactly. Although the machine account is sometimes the same as the DNS A record created during the domain-join process, depending on your the DNS environment, this might not always be true.

In the following example, a Qumulo cluster has a machine account with the SPN nfs/qumulo.example.com and two DNS A records that point to the same Qumulo cluster IP, 203.0.113.0:

- qumulo.example.com
- storage.example.com

Because the storage.example.com doesn't have a corresponding SPN, you can perform Kerberos authentication by using the qumulo.example.com record. However, if you add the second SPN (nfs/storage.example.com) to the machine account account SPN list, the account can authenticate by using either of the two hostnames.

CNAME (alias) records are an exception to this arrangement. CNAME records that point to a correctly-configured A record, and which have a corresponding SPN entry in the machine account, don't require the CNAME host to be added to the SPN. For example, the CNAME record storage-alias.example.com that points to storage.example.com requires the SPN list to contain only nfs/storage.example.com to authenticate against storage-alias.example.com.

Performing Additional Cluster Configuration after Joining Active Directory

This section describes additional Qumulo cluster configuration that can affect the behavior of NFSv4.1 with Kerberos.

When your Qumulo cluster is joined to AD (page 174), you must configure the NFSv4.1 server (page 157) and NFSv4.1 security settings.

To Configure Security Settings by Using the qq CLI

Qumulo provides configuration for the permitted NFSv4.1 authentication flavors in the qq CLI or directly through the REST API.

1. Use the qq CLI to get the current settings:

```
$ qq nfs_get_settings
{
    "auth_sys_enabled": true,
    "krb5_enabled": true,
    "krb5p_enabled": true,
    "krbi_enabled": true,
    "v4_enabled": false
}
```

This is the default configuration:

- · NFSv4.1 is disabled by default.
- AUTH_SYS, AUTH_KRB5, AUTH_KRB5P, and AUTH_KRB5I are enabled by default (however, Qumulo Core doesn't support Kerberos configuration on NFSv3).
- 2. To harden security, configure your cluster to use only Kerberos by disabling AUTH_SYS (without changing AUTH_KRB5). For example:

A Important

Because it uses authentication based on a simple UID and GID passed over the wire in plain text, RPC AUTH_SYS is inherently insecure. In a trusted environment, AUTH_SYS might be sufficient for enforcing basic permissions and preventing good-faith actors from making mistakes. In all other cases, you must treat AUTH_SYS as if it provides no security whatseover.

```
$ qq nfs_modify_settings --disable-auth-sys
{
    "v4_enabled": false,
    "auth_sys_enabled": false,
    "auth_krb5_enabled": true,
    "auth_krb5p_enabled": true,
    "auth_krb5i_enabled": true
}
```

3. (Optional) You can also use the following commands.

Command	Description
<pre>qq nfs_modify_settingsenable- auth-sys</pre>	Enables AUTH_SYS without changing AUTH_KRB5
qq nfs_modify_settingsenable-krb5	Enables AUTH_KRB5 without changing AUTH_SYS
<pre>qq nfs_modify_settingsenable- krb5p</pre>	Enables AUTH_KRB5P without changing AUTH_SYS
qq nfs_modify_settingsenable- krb5i	Enables AUTH_KRB5I without changing AUTH_SYS
qq nfs_modify_settingsenable-v4	Enables NFSv4.1
qq nfs_modify_settingsdisable-v4	Disables NFSv4.1
<pre>qq nfs_modify_settingsdisable- krb5</pre>	Disables AUTH_KRB5 without changing AUTH_SYS
qq nfs_modify_settingsdisable- krb5p	Disables AUTH_KRB5P without changing AUTH_SYS
qq nfs_modify_settingsdisable- krb5i	Disables AUTH_KRB5I without changing AUTH_SYS

O Note

- Security configuration options apply to all versions of NFS (NFSv3 and NFSv4.1). Thus, disabling AUTH_SYS also disables NFSv3, because AUTH_SYS is the only authentication flavor that NFSv3 supports by design.
- In a secure environment, where Kerberos is required, AUTH_SYS NFSv3 connections aren't allowed.
- · These configuration options apply cluster-wide to all NFS exports and files.

Configuring Export Configuration

You can use NFSv4.1 exports (page 157) to configure access to the Qumulo file system.

The user-mapping portion of the export configuration has no effect on Kerberos configuration. Specifying root or any user mapping for a particular export applies only to AUTH_SYS mounts that access this export.

Otherwise, exports and IP address restrictions (that you specify in exports) behave identically for all authentication flavors: AUTH_SYS, AUTH_KRB5, AUTH_KRB5P, and AUTH_KRB5I.

Using Kerberos Permissions in the Qumulo File System

This section describes how NFSv4.1 interacts with the secure file permissions that Kerberos enables for the Qumulo Core file system.

For more information, see Qumulo File Permissions Overview on Qumulo Care.

Listing Permissions for Files

Note

- This section uses the Kerberos term *trustee* and Qumulo term *identity* (or auth_id) interchangeably.
- The term file in the Qumulo file system can refer to:
 - A file
 - A directory
 - A symbolic link
 - A special block device

All files in the Qumulo file system have the following fields associated with them:

- Owner
- · Group owner
- · Access control list (ACL)—a list of access control entries (ACEs)

These fields, stored in the metadata for a file or directory, determine the access permissions that a trustee or identity has to files.

For any file operation, the system checks the authenticated user against file permissions to determine whether the operation should be allowed. When you create a new file, the authenticated user becomes the owner of the new file.

In the following example, we create a file in a mount over NFS.

O Note

- Because this example uses an AUTH_SYS mount, it has UID and GID identity values set to 1000.
- We recommend becoming familiar with the following commands to better understand the various elements for permissions types that the system stores on disk.

```
touch /mnt/mount_point/filename
```

To view the exact permissions metadata for this file, use the qq fs_file_get_attr command. For example:

```
$ qq fs_file_get_attr --path /filename
{
    "group_details": {
        "id_type": "NFS_GID",
        "id_value": "1000"
},
    "owner_details": {
        "id_type": "NFS_UID",
        "id_value": "1000"
},
    ...
}
```

To view the permissions configured in an ACL, use the qq fs get acl command. For example:

```
$ qq fs_get_acl --path /filename
Control: Present
Posix Special Permissions: None
Permissions:
Position Trustee
                 Type
                         Flags Rights
Delete child, Read, Write file
1
        uid:1000 Allowed
2
        gid:1000 Allowed
                               Delete child, Read, Write file
3
        Everyone Allowed
                               Read
```

Listing Security Identifiers (SIDs)

The SID is a globally unique identifier for a user or group object in a domain. For more information, see Security identifiers in the Microsoft documentation.

Because Qumulo's Kerberos implementation requires AD, every user is also an Active Directory user. The domain controller (DC) has an equivalent mapping for AD users and SIDs. Qumulo uses LDAP to determine the AD-user \leftrightarrow SID mapping. For this reason, it is important to configure the Base DN for your cluster correctly.

Qumulo's Kerberos implementation stores SIDs on disk for files that have Kerberos identities in the user, group, or ACL. When a user authenticates by using Kerberos and creates a file, Qumulo Core configures the user, group, and ACL automatically.

To set the identity for an AD user, you can modify the permissions for an existing file by using the chown or nfs4_setfacl command.

In the following example, the Kerberos-authenticated AD domain user AD\myusername creates a file over NFSv4.1 and the system gives an ACL response from the REST API. The response contains an ACE entry for the owner and group owner of the user AD\myusername, with corresponding SIDs for both.

```
$ qq fs_get_acl --path /filename --json
{
  "aces": [{
    "trustee": {
      "name": "AD\\myusername",
      "sid": "S-1-5-21-4202559609-EXAMPLE158-3224923410-13507",
    },
    . . .
  }, {
    "trustee": {
      "name": "AD\\Domain Users",
      "sid": "S-1-5-21-4202559609-EXAMPLE158-3224923410-513",
      . . .
    },
  }]
}
```

Using Kerberos Principals

Although Qumulo stores SIDs on disk, SIDs appear rarely when you use NFSv4.1 on Linux systems. Instead, the system represents Kerberos identities as Kerberos principals. A *Kerberos principal*, a string in the <user@domain> or <ure>qroup@domain> format, is easier to read.

O Note

There is an equivalent mapping between AD users, SIDs and Kerberos principals. Each of these representations is unique (a primary key to the AD identity database).

Qumulo's implementation of the SID ↔ Kerberos principal mapping uses the sAMAccountName field, which is always present and unique for all AD users and groups. The system forms the Kerberos principal by concatenating the name and domain in the sAMAccountName>@<domain> format.

AD has fields with similar content but without the guarantee of uniqueness (such as the name, distinguishedName, CN, and servicePrincipalName). However, AD permits setting these fields to unrelated values. For this reason, it is unlikely but possible that certain environments use special values in these fields. Qumulo's Kerberos implementation ignores these fields and uses only the value in the sAMAccountName field.

O Note

The fields can diverge significantly if an administrator edits them.

The following example shows how the system represents the SIDs from the previous example as Kerberos principals.

\$ nfs4_getfacl filename
A::test2@ad.eng.qumulo.com:rwatTnNcy
A:g:Domain Users@ad.eng.qumulo.com:rtncy

A::EVERYONE@:rtncy

Although the system stores raw SIDs on disk, the nfs_getfact command displays users and groups as Kerberos principals. This format is valid for setting identities on a file by using commands such as nfs4 setfact, chown, and so on.

Understanding Kerberos Principal Caveats

This section explains some of the caveats of working with Kerberos principals.

Machine Account Object Names

When you work with machine accounts, AD stores the sAMAccountName as the object name and appends \$ to it. If a client named myclient is joined to the domain stuff.example.com, the name of the machine account object in Active Directory Users or Computers appears as myclient while the Kerberos principal representation over NFS appears as myclient\$@stuff.example.com.

This functionality is different from other account types in AD, where the object name usually matches the samaccountName exactly.

ID Mapping on Linux systems

Linux systems perform their own ID mapping separately from the Qumulo cluster ID mapping.

Linux systems also use SAMAccountName as the AD user primary key when joined to an AD domain.

However, Linux systems use CN when looking up groups. Thus, in groups where the SAMAccountName and CN don't match (possibly due to edits by an administrator), a Linux system and Qumulo Core might understand differently the group that the Kerberos principal refers to.

Ensure the two fields are in sync to prevent the following possible scenarios:

- · An error appears when you configure the group.
- · Group configuration succeeds but the configured group is incorrect.

Unicode Characters in Kerberos Principals

For most standard Linux tools, Qumulo Core supports all arbitrary Unicode characters in Kerberos principals. However, we don't recommend using the period (.) character in principals, except in the domain name.

Using the chown Tool With Kerberos

chown is a Linux tool that changes the owner or group owner for a file. You can generally use chown with Kerberos principals. On most Linux systems, chown requires the root user (sudo chown).

The AUTH SYS Root User

AUTH_SYS has the concept of the root user. Using sudo on a Linux NFS client fills in 0 for the UID and GID. As long as the mounted export doesn't root squash—maps a client's UID 0 (root) to 65534 (nobody) or to another non-root user—the Linux client receives root permissions on the Qumulo file system, where the client can perform chown operations.

The Kerberos Root User

Kerberos doesn't have the concept of the root user. However, you can still use it to run chown operations under the following conditions.

- The ACL for the file must grant the CHANGE OWNER privilege to an authenticated user.
- The currently authenticated user must be a member of the destination group (if provided) or a member of the current group (if the group isn't being modified).

If both conditions are true, a **chown** operation on files performed as a Kerberos user over NFSv4.1 succeeds. For example:

\$ chown user3:group4 filename

O Note

Including @<domain> for the destination user and group is optional.

Viewing the Owner and Group

The following examples show how to display user and group membership by using the ls -l and stat -c commands.

```
$ ls -l filename
-rw-r--r-- 1 user3 group4 0 Jun 9 23:18 filename
```

```
$ stat -c '%U, %G' filename
user3, group4
```

O Note

The Kerberos restrictions for chown also apply to other Linux tools that use the chown system call, such as cp and rsync, when you run them in ownership-preserving modes.

Using the Linux ACL Editor

The Linux ACL Editor consists of the following tools:

- nfs4 editfacl
- nfs4 getfacl
- nfs4 setfacl

You can use the editor to read and write ACLs on a Qumulo cluster that uses NFSv4.1 with Kerberos. For more information, see Managing File Access Permissions by Using NFSv4.1 Access Control Lists (ACLs) (page 164).

Configuring a Linux Client for NFSv4.1 with Kerberos

This section describes how to configure a Linux client for using NFSv4.1 with Kerberos.

1 Note

Qumulo Core supports only Linux for using NFSv4.1 with Kerberos.

Linux systems implement Kerberos support as a series of loosely related packages and configuration files. For this reason, configuration depends on the Linux distribution and version. This section refers to tools, packages, dæmons, configuration files, and other elements in Ubuntu 18.04 LTS.

Joining a Linux Client to a Domain

There are two common ways of joining a Linux client to an Active Directory (AD) domain automatically, by using samba or realmd. Both methods require creating the /etc/krb5.conf configuration file and defining a default domain and the relationships between domains and realms.

Configuring the /etc/krb5.conf File

The following is an example configuration for joining a domain.

```
[libdefaults]
  default_realm = MY-DOMAIN.EXAMPLE.COM

[realms]
  MY-DOMAIN.EXAMPLE.COM = {
    kdc = my-domain.example.com:88
    admin_server = my-domain.example.com:749
  }

[domain_realm]
  my-domain.example.com = MY-DOMAIN.EXAMPLE.COM
  .my-domain.exmaple.com = MY-DOMAIN.EXAMPLE.COM
```

To Join a Linux Client to a Domain by using samba

samba is a suite of Linux tools that provides Windows-like functionality on Linux. The net-ads join command creates a machine account on the domain.

1. To specify how the domain-join process behaves, edit the /etc/samba/smb.conf file. For example:

```
workgroup = my-domain
server role = member server
realm = my-domain.example.com
kerberos method = system keytab
```

2. To join the domain, run the net ads join command. For example:

```
$ net ads join my-domain.example.com -U Administrator
```

3. samba doesn't create configuration files. Configure the sssd and idmapd tools manually. For more information, see Mapping External Identities to Linux Identities (page 191).

To Join a Linux Client to a Domain by using realmd

realmd is a tool that allows managing realm-based authentication. It can be somewhat more difficult to use than samba. However, it creates a more complete configuration. For example, it configures the sssd tool during the domain-join process.

1. To join a domain, use the realm join command. For example:

```
$ realm join my-domain.example.com -U Administrator
```

2. Configure the sssd and idmapd tools manually. For more information, see Mapping External Identities to Linux Identities (page 191).

To Configure DNS and Service Principal Name (SPN)

Kerberos relies on DNS to identify machines involved in authentication. NFS clients and servers require DNS A records for forward-DNS lookups and PTR records for reverse-DNS lookups.

1. After you configure DNS, check DNS resolution from your client. For example:

```
$ nslookup my-client-machine.my-domain.example.com
```

2. In addition to DNS configuration, Linux clients require a standard host SPN on the machine account created while joining the domain. We recommend configuring the SPN by using the setspn command on the domain controller after the join procedure. For example:

O Note

Running this command resets the SPN to the default value for your machine.

setspn -r my-client_machine

Mapping External Identities to Linux Identities

During the ID mapping process, a Linux system converts external identities to Linux identities.

- · For Qumulo Core, external identities are equivalent to Kerberos principals.
- · For Linux, identities are simple integers: UIDs and GIDs.

O Note

Because Linux can't use complex external identities in system calls, a Linux system must perform identity conversion before operating on files.

ID mapping is bidirectional. A system call, such as **chown**, that takes a UID or GID as input requires mapping the UID or GID be mapped to a domain user or group *before* passing it to your Qumulo cluster over NES.

A system call, such as stat, that returns a UID or GID, requires that the domain user or group that returned from your Qumulo cluster over NFS be converted to a UID or GID before the system can present it to the user.

Configuring Active Directory Authentication by using sssd

sssd (System Security Services Daemon) is a tool responsible for managing authentication with external providers in Linux. To use NFSv4.1 with Kerberos, you must configure sssd with AD as the identity provider.

- If you join domains by using samba, you must create the /etc/sssd.conf file.
- If you join domains by using realmd, you might already have a /etc/sssd.conf file. For detailed configuration information, see sssd-ldap in the Linux documentation.

In the following example, the sssd.conf file configures basic ID mapping for AD.

```
[sssd]
domains = my-domain.example.com
config_file_version = 2
services = nss, pam

[domain/my-domain.example.com]
ad_domain = my-domain.example.com
krb5_realm = MY_DOMAIN.EXAMPLE.COM
cache_credentials = True
id_provider = ad
krb5_store_password_if_offline = True
default_shell = /bin/bash
ldap_id_mapping = False
use_fully_qualified_names = False
fallback_homedir = /home/%u@%d
access_provider = ad
```

Configuring LDAP Queries against the Domain Controller (DC) by using sssd

Like Qumulo clusters, Linux systems can resolve details about user and group objects by querying the DC over LDAP. In particular, a Linux system looks for an object with a matching sAMAccountName (user) or CN (group)

- 1. To toggle RFC 2307 for mappings in the sssd.conf file, configure the ldap_id_mapping field.
 - When you set the field to False, the client checks whether the RFC 2307 uidNumber or gidNumber are set on an object.
 - If the number is set, it becomes the Linux UID or GID for the operation.

A Important

AD doesn't prevent duplicate UID or GID numbers from being added to RFC 2307 values. For this reason, incorrect configuration can lead or UID or GUID collisions. When a Linux system determines that a collision has occurred, it chooses the first UID or GID it finds.

Otherwise, the UID or GID becomes nobody or nogroup (65534).

O Note

In most cases, an owner or group becomes 65534 as a result of incorrect user mapping configuration in the client. To understand which LDAP queries run and why they have trouble finding the correct information, check your logs.

• When you set the field to True, the client assigns locally a new unique UID or GID to each objectSID that it finds on the DC.

O Note

This is a more flexible approach than requiring RFC 2307. However, this also means that UIDs and GIDs aren't the same across different Linux systems within the same domain.

In both cases, the client communicates with the DC by using its machine account.

2. To pick up changes to the /etc/sssd.conf file on a live system, restart the sssd service.

Configuring the Conversion of Local Identities to NFS Representations by Using idmapd

idmapd (or nfsidmap), is a tool that lets you convert local identities to their on-the-wire NFS representations. Although idmapd works with sssd, it has additional configuration options.

In the following example, the (etc/idmapd.conf file configures a Linux client joined to AD:

```
[General]
Domain = my-domain.example.com
Verbosity = 0
Pipefs-Directory = /run/rpc_pipefs

[Mapping]
Nobody-User = nobody
Nobody-Group = nogroup
```

O Note

Depending on your Linux distribution and configuration, you might have to add the **Domain** field to the default configuration file.

Authenticating as an AD User and Mounting Your Qumulo Cluster

Qumulo Core supports three methods of authenticating as an AD user and mounting your cluster over NFSv4.1 as the AD user. These methods, from least to most complex, and in an increasing order of utility, are:

- By using a machine account
- · By using manual authentication with the kinit tool
- · By using the autofs tool

To Authenticate as an AD User by Using a Machine Account and Mount Your Qumulo Cluster

Machine account authentication uses one AD user for each Linux system. This *machine account user* is the same as the *machine account* created on the domain during the domain-join operation. Any user on the Linux system who has access to the machine account mount point can operate as the machine account user on a Qumulo cluster.

Machine account authentication can be useful for simple scenarios in which trusted users on trusted Linux machines require a secure mechanism for communicating with a Qumulo cluster. Because this is also the easiest authentication method to configure, it can be a good starting point for administrators who configure NFSv4.1 with Kerberos for the first time.

Note

Both machine account authentication and kinit have limited usefulness because they limit the mount point to a single authenticated user. Between the two authentication options, kinit has an advantage because of the way it handles ID mapping.

1. Confirm that your /etc/nfs.conf file, contains the following flag.

```
[gssd]
use-machine-creds=true
```

The use-machine-creds flag specifies whether authentication uses machine credentials when sudo mount is invoked for NFSv4.1 with Kerberos. When you set the flag to true, gssd authenticates as the machine account for the system on behalf of the NFS client. (It performs a kinit operation as the machine account). The credential cache that results from the kinit is usually located in /tmp. To search for the cache, use the ls/tmp/*krb5* command.

Note

In versions of Ubuntu lower than 22.04 (and possibly on other Linux distributions), you can't use the /etc/nfs.conf file to configure gssd. If this is the case for your system, we recommend starting the rpc.gssd service by using the -n flag.

2. Mount your cluster by using the krb5 security mechanism. For example:

```
$ sudo mount -o vers=4.1,sec=krb5 my-cluster.my-domain.example.com:/ /mnt/poin
t
```

3. Use the Qumulo file system.

▲ Important

The machine account is the owner of any new files.

If the machine name isn't visible, make sure that the AD container holds this machine in the Qumulo cluster's Base DN configuration (typically, CN=Computers, DC=...). If the machine name is still not visible, configure the Linux client ID mapper to provide local mappings when no RFC 2307 mapping is available. It is uncommon for machine accounts to have RFC 2307 mappings.

To Authenticate as an AD User Manually by Using kinit and Mount Your Qumulo Cluster

kinit authentication is very similar to machine account authentication. The main difference is that you must create the credentials for the mount manually. You can use any user in the AD domain. However (this is also true for machine accounts), any local Linux user that can access the mount point can operate on the Qumulo cluster as this single user.

O Note

Both machine account authentication and kinit have limited usefulness because they limit the mount point to a single authenticated user. Between the two authentication options, kinit has an advantage because of the way it handles ID mapping.

In environments where Linux systems map exactly to end users that have kinit -based Kerberos mounts on their Qumulo clusters, kinit might be sufficient.

1. Authenticate by using kinit . For example:

```
$ sudo kinit my-user
```

- 2. When prompted for a password, use the AD domain password for the user.
- 3. To confirm the result of the authentication operation, use the sudo klist command.
- 4. Confirm that the /etc/nfs.conf file contains the following flag:

```
[gssd]
use-machine-creds=false
```

The use-machine-creds flag specifies whether authentication uses machine credentials when sudo mount is invoked for NFSv4.1 with Kerberos. When you set the flag to false, gssd searches for an existing credential cache (which you created by running kinit) in /tmp/krb5cc_0 for authenticating with the Qumulo cluster.

5. Mount your cluster by using the krb5 security mechanism. For example:

```
$ sudo mount -o vers=4.1,sec=krb5 my-cluster.my-domain.example.com:/ /mnt/poin
t
```

6. Use the Qumulo file system.

▲ Important

The kinit user is the owner of any new files.

To Authenticate as an AD User Manually by Using autofs and Mount Your Qumulo Cluster

autofs is a dæmon that manages mount points for individual Linux users. For this reason, Linux users have different views of a mount point. autofs can authenticate an AD user through ssh, the Linux file system, or a Qumulo cluster mounted on a Linux system.

▲ Important

When you use autofs, the Linux system maps the root user to the machine account user for the Linux system on the Qumulo cluster. However, the machine account user doesn't have all the privileges of the root user, such as special permissions for the Qumulo cluster. You must specify all permissions in ACLs.

1. Log in to an AD domain and configure sssd to authenticate with this domain. For example:

```
$ sudo login my-domain-user
```

Alternatively, you can use the following command.

```
$ ssh my-domain_user@my-linux-system
```

- 2. Configure the autofs mappings. For more information, see auto.master in the Linux documentation. The following is an example of a simple configuration that provides a single (direct) mount point which authenticates AD users automatically.
 - a. To define a mount point and the path to its map file, add the following line to the /etc/auto.master file.

```
/- /etc/auto.kerberos_nfs_mount_example --timeout 60
```

For more information, see Autofs in the Ubuntu documentation.

b. Add the following line to the /etc/auto.kerberos_nfs_mount_example map file.

```
/mnt/qumulo_mount_point -vers=4.1,sec=krb5 <qumulo-cluster>.my-domain.e
xample.com:/
```

3. Restart autofs.

```
$ sudo systemctl restart autofs
```

autofs creates the /mnt/qumulo_mount_point directory and mounts it as necessary for any user. For example:

```
$ ssh domain_user_1@my-linux-system touch /mnt/qumulo_mount_point/user1_file
$ ssh domain_user_2@my-linux-system touch /mnt/qumulo_mount_point/user2_file
$ ssh domain_user_3@my-linux-system ls -l /mnt/qumulo_mount_point
-rw-r--r-- 1 user1 domain users 0 Jun 9 23:18 user1_file
-rw-r--r-- 1 user2 domain users 0 Jun 9 23:18 user2_file
```

A Important

The user you logged in to the AD domain with is the owner of any new files.

Network Time Protocol (NTP) Server

Kerberos is very sensitive to clock skew. It is important for all systems involved in a Kerberos relationship—the KDC, your Qumulo cluster, and any Linux clients—to have as little clock skew as possible. We recommend using the same NTP server for all three components.

- You can use your AD domain controller as an NTP server. In the Web UI, on the Active Directory page, for Use Active Directory as your primary time server, click Yes.
- To configure any other NTP server in the Web UI, click Cluster > Date & Time.

There are many NTP dæmons for Linux. For example, Ubuntu uses the NTP functionality in systemd (timedatectl and timesyncd).

Configuring Cross-Domain Active Directory Trusts

This section describes how the configuration of cross-domain Active Directory (AD) trusts supports NFSv4.1 with Kerberos.

Trusts are relationships between different AD domains. For more information, see Trust Technologies in the Microsoft documentation.

NFSv4.1 with Kerberos and the general AD configuration in Qumulo Core support the same forms of trust relationships.

- · Child or parent trusts can:
 - Authenticate as a user from the child domain against the parent domain's AD domain controller (DC).
 - Authenticate as a user from the parent domain against the child domain's AD DC.
- Transitive trusts can authenticate as a user from any of the domains in the transitive trust, against any of the other trusted domains' AD DC.

Configuring the Base DN

For identity mapping to work, you must configure LDAP Base DNs correctly on your Qumulo cluster and on your client. This helps avoid nobody or 66534 identity responses that occur when you inspect files that contain trusted users (stored as identities) from other domains. For more information about configuring the Base DN, see Using Active Directory for POSIX Attributes on Qumulo Care.

The following example has trust between between parent.example.com and child.example.com. In order for both domains' identities to authenticate against a Qumulo cluster, you must configure the cluster and your client with the following Base DN.

CN=Users,DC=parent,DC=example,DC=com;CN=Users,DC=child,DC=parent,DC=example,DC=com

O Note

AD doesn't prevent duplicate UID or GID numbers from being added to RFC 2307 values. Such improper configuration can cause UID and GID collisions across trusted domains. On Linux, if any collisions occur, the system chooses the first UID or GID that it finds.

Enabling More Secure Trust Encryption Types

While Linux systems disallow deprecated encryption types for Kerberos, Windows prefers RC4 for cross-domain traffic (which Linux systems consider to be deprecated).

For certain trust configurations, you must enable a more secure encryption type for trusted traffic. To enable AES-128 (or SHA1) and AES-256 (or SHA1) for a particular trust, use the ksetup command in a Windows Administrator console. For example:

- \$ ksetup /getenctypeattr <domain>
- \$ ksetup /setenctypeattr <domain> RC4-HMAC-MD5 AES128-CTS-HMAC-SHA1-96 AES256-CTS-HM
 AC-SHA1-96

O Note

This example doesn't disable RC4. Instead, it enables new encryption types *in addition* to RC4. When working with Windows systems, we recommend making additive changes whenever possible. We also recommend staging changes in a safe environment before applying them to a production environment.

Troubleshooting NFSv4.1 with Kerberos

This section describes common troubleshooting procedures for configuring NFSv4.1 to work with Kerberos.

Following General Debugging Techniques

This section lists common debugging techniques.

To Turn Up Logging Levels for Client-Side Tools

- 1. In the /etc/sssd.conf file, set debug level = 9.
- 2. In the /etc/idmapd.conf file, set Verbosity = 9.
- 3. In the [gssd] section of the /etc/nfs.conf file, set verbosity=9 and rpc-verbosity=9.

1 Note

In versions of Ubuntu lower than 22.04 (and possibly on other Linux distributions), you can't use the /etc/nfs.conf file to configure gssd. If this is the case for your system, we recommend starting the rpc.gssd service by using the -n flag.

4. Turn on rpcdebug, for example:

```
rpcdebug -m nfs -s all && rpcdebug -m rpc -s all
```

Taking a Client-Side Packet Capture

Normally, there should be:

- · Kerberos and LDAP traffic between the client and the domain controller
- · DNS traffic between the client and DNS server
- · RPC or NFS traffic between the client and the Qumulo cluster

Because a Kerberos mount requires the client to perform a series of steps, in most cases, the last traffic that the client issues indicates the source of failure. To view encrypted Kerberos traffic, use Wireshark with a Kerberos keytab file. For more information, see Kerberos in the Wireshark documentation.

For help with interpreting logging and metrics from your Qumulo cluster and for insights from the telemetry of our Kerberos implementation, contact the Qumulo Care team.

Resolving Incorrect Display of Users or Groups

Under certain conditions, users or groups display as **nobody** when you run the **ls -l** or **stat** command.

Differentiating Client and Cluster Issues

To resolve this issue, determine whether it is with the client or with the cluster by running the nfs4_getfacl command on a file. If the presentation in the ACL editor appears correct, the issue is with the client. Otherwise, the issue is with the cluster.

O Note

The ACL editor doesn't perform any ID mapping. It only passes ACE trustees through, in plaintext.

Resolving Client-Side Issues

If the issue is with the client, it is most often an ID mapping issue. Confirm that your mappings are configured correctly. For more information, see User-Defined Identity Mappings on Qumulo Care.

If the issue persists, investigate logging and packet captures.

Resolving Cluster-Side Issues

If the issue is with the cluster, confirm that your cluster's Active Directory settings include the Base DNs that contain the expected users. For more information, see Prerequisites for Joining a Qumulo Cluster to Active Directory (page 174).

Diagnosing Mount-Failed Errors

Under certain conditions, you might receive mount-failed errors from mount.nfs. To diagnose this type of error, you can try the following procedures.

- 1. Confirm that the rpc.gssd service is running.
- 2. Confirm that the cluster and client both resolve from the client. It should be possible to reach the cluster and client through a fully qualified domain name (FQDN), such as my-machine.my-domain.example.com.
- 3. Confirm that reverse DNS works for the IP addresses on both the client and the cluster.
- 4. Confirm that the client has a **host** service principal name (SPN) and that the cluster has an **nfs** SPN that matches the DNS records.
- 5. Do one of the following:
 - If you use a machine account or kinit authentication, confirm that the
 credentials are correct. You can use the keytab ktutil command or the
 credential cache klist command to list the encryption methods.

- Confirm that Kerberos tickets use AES-128 or AES-256 for service encryption by examining a packet capture or your Active Directory Kerberos settings.
- 6. If you use domain trusts, confirm that trust has AES-128 or AES-256 enabled.
- 7. Confirm that the clocks on the client, cluster, and domain controller are synchronized to the same time.
- 8. Inspect logs and packet captures.

S3 API

Configuring and Using the S3 API in Qumulo Core

This section explains how to configure and get started working with the S3 API. This API lets clients and applications interact with the Qumulo file system natively, by using the Amazon S3 API.

Prerequisites

To use the S3 API, you must install the aws and qq CLI tools.

A Important

The following instructions are for Ubuntu 18.04 (and higher).

Step 1: Configure HTTPS

The Qumulo Core S3 API accepts only HTTPS requests by default. To enable HTTPS support for your Qumulo cluster, you must install a valid SSL certificate on it.

Every Qumulo cluster is preconfigured with a self-signed SSL certificate. However, because certain applications don't accept the default certificate, we recommend installing your own.

To configure HTTPS for your cluster, follow the steps in SSL: Install a Signed SSL Certificate on Qumulo Care.

Enabling and Disabling Plaintext HTTP Connections

A Important

If you configure the S3 API service to accept only plaintext HTTP connections, no requests made through the S3 API are encrypted.

- To enable HTTP connections, use the qq s3_modify_settings --insecure command.
- To revert to encrypted HTTPS requests, use the qq s3_modify_settings --secure command.

Step 2: Enable the S3 API for Your Qumulo Cluster

To let your Qumulo cluster accept S3 traffic, you must enable the S3 API by using the qq s3_modify_settings --enable command.

After you run the command, all nodes in your cluster begin to accept S3 API traffic on TCP port 9000.

Step 3: Create an Access Key Pair

To create and manage S3 buckets you must have a valid S3 access key pair associated with a specific user in your Qumulo cluster or in a connected external identity provider (such as Active Directory). For more information, see Creating and Managing S3 Access Keys (page 209).

Use the qq s3 create access key and specify the username. For example:

```
$ qq s3 create access key my-username
```

Note

After Qumulo Core initially creates your secret access keys, it never logs or displays them again. If you lose your secret access key, it isn't possible to recover it and you must create a new access key pair.

Step 4: Configure the AWS CLI for Use with Qumulo Core

To create and manage S3 buckets, you must configure AWS CLI to work with your Qumulo cluster.

Note

- We recommend configuring a dedicated profile for Qumulo in your AWS CLI S3 Configuration.
- Qumulo Core listens for S3 API traffic on TCP port 9000. It isn't possible to change this setting.
- Currently, Qumulo Core supports only path-style bucket addressing. For more information, see Bucket Addressing Style (page 240).
- 1. Configure the AWS CLI to use path-style bucket addressing by using the aws configure command and specify your profile.

```
$ aws configure \
   --profile my-qumulo-profile set s3.addressing_style path
```

- 2. Use the access key pair that you have created earlier (page 0) and the aws configure command to:
 - a. Specify your profile and access key ID (page 209). For example:

```
$ aws configure
--profile my-qumulo-profile set aws_access_key_id \
    000000000001fEXAMPLE
```

b. Specify your profile and secret access key (page 209). For example:

```
$ aws configure
--profile my-qumulo-profile set aws_secret_access_key \
    TEIT4liMZ8A32iI7JXmqIiLWp5co/jmkjEXAMPLE
```

3. Because it isn't possible to specify your cluster's URI persistently, create a shell alias to specify your cluster's URI, in the following format:

```
$ alias aws="aws --endpoint-url https://<qumulo-cluster>:9000 --profile my-qum
ulo-profile"
```

O Note

If you haven't installed an SSL certificate, append --no-verify-ssl to the end of the command.

4. (Optional) If you haven't configured your machine to trust the SSL certificate installed on your Qumulo cluster, to configure the path to the trusted SSL certificate bundle that you have created and installed earlier (page 0) manually, use the aws.configure command. For example:

```
$ aws configure \
   --profile my-qumulo-profile set ca_bundle MySpecialCert.crt
```

5. To test your configuration, send an S3 API request to your Qumulo cluster by using the aws s3api list-buckets command.

A successful response includes an empty JSON array named **Buckets**.

```
{
   "Buckets": []
}
```

Step 5: Create an S3 Bucket

O Note

Creating buckets requires the PRIVILEGE_S3_BUCKETS_WRITE role-based access control (RBAC) privilege and permission to create a directory under the cluster's root directory.

Use the aws s3api create-bucket command and specify the bucket name. For example:

```
$ aws s3api create-bucket \
--bucket my-bucket
```

The S3 API creates the new directory /my-bucket/. All of the bucket's objects are located under this directory. For more information, see Creating and Working with S3 Buckets in Qumulo Core (page 217).

Step 6: Test Writing and Reading S3 Objects

1. To test writing data to your Qumulo cluster, perform a Put0bject S3 API action by using the aws s3api put-object command. For example:

```
$ aws s3api put-object \
  --bucket my-bucket \
  --key archives/my-remote-file.zip \
  --body my-local-file.zip
```

The S3 API uploads the contents of my-local-file.zip into an object named my-remote-file.zip.

2. To test reading read data from and S3 bucket, perform a Get0bject S3 API action by using the aws s3api get-object command. For example:

```
$ aws s3api get-object \
  --bucket my-bucket \
  --key archives/my-remote-file.zip local-file.zip
```

The S3 API downloads the contents of the my-remote-file.zip object into local-file.zip and returns the object metadata. For example:

```
{
  "AcceptRanges": "bytes",
  "LastModified": "Wed, 14 Dec 2022 20:42:46 GMT",
  "ETag": "\"-gUAAAAAAAAAAAAAA\\"",
  "ContentType": "binary/octet-stream",
  "Metadata": {}
}
```

Creating and Managing S3 Access Keys in Qumulo Core

This section explains how to create and manage credentials that S3 API actions in Qumulo Core require to access file system resources, such as access key pairs that sign requests.

O Note

You can configure an S3 bucket to allow read-only, anonymous access (page 227). This approach requires no credentials but limits users to non-modifying operations. For more information, see To Enable Anonymous Access to S3 Buckets by Using the qq CLI (page 228).

Prerequisites

Managing S3 access keys requires the following role-based access control (RBAC) privileges:

- PRIVILEGE_S3_BUCKETS_WRITE: Create and delete S3 access keys
- PRIVILEGE_S3_BUCKETS_READ: List S3 access keys

How S3 Access Keys Work in Qumulo Core

An *identity* is a single principal from an identity provider (IdP). Examples of identities include SMB security identifiers (SIDs), Active Directory user principal names (UPNs), POSIX user identifiers (UIDs), and local users in a Qumulo cluster.

A Important

It isn't possible to create access keys for UIDs in an Active Directory environment that has POSIX extensions enabled. However, it is possible to use Active Directory identity identifiers (SIDs, UPNs, and so on).

An access key (or access key pair) is comprised of an S3 access key ID and an S3 secret access key.

- The access key ID is the public component of an S3 access key pair. It identifies the user that performs an S3 request.
- The secret access key (or secret key) is the private component of an S3 access key pair. The client uses the secret access key to sign requests and the server uses the secret access key to validate request signatures.

A Important

- Qumulo Core uses a cryptographically secure source, certified according to FIPS 140-2 requirements, to derive secret access keys.
- Because access keys are cluster-local, you can't use an access key for an identity in one Qumulo cluster on a different Qumulo cluster.

Qumulo Core creates an access key pair whenever an authorized user requests it. For more information, see Creating S3 Access Keys for a Qumulo Cluster (page 211).

The way in which Qumulo Core access keys let you access your Qumulo cluster makes the process similar to the way in which IAM Access Keys let you access Amazon S3 resources. For this reason, applications that access objects stored in a Qumulo cluster can use the Qumulo S3 API similarly to the native Amazon S3 API.

How S3 Access Keys work with Identities

An S3 access key doesn't grant any additional permissions. It associates an S3 API request with a specific identity (page 209) that the Qumulo cluster knows.

When Qumulo Core processes a request, it evaluates permissions by using the Qumulo ACL (QACL) mechanism that operates like the access control list (ACL) mechanism that all file system protocols use. When the QACL grants or denies permissions to an associated identity, it also grants or denies the same permissions to the request being processed.

For more information, see Managing Access to S3 Buckets in a Qumulo Cluster (page 226).

How Qumulo Core Stores S3 Access Keys

To authenticate S3 API requests, Qumulo Core retrieves existing access key pairs that it stores securely as configuration metadata in your Qumulo cluster. Qumulo Core encrypts secret access keys on disk and holds decrypted secret access keys in memory only while it processes a request.

A Important

Because (unlike secret access keys) your access key IDs aren't a cryptographic secret, Qumulo Core can log and display access key IDs. After Qumulo Core initially creates your secret access keys, it never logs or displays them again. If you lose your secret access key, it isn't possible to recover it and you must create a new access key pair.

S3 Access Key Lifecycle in Qumulo Core

Qumulo Core doesn't limit how long you can use an access key pair after you create it. Your system administrators must take responsibility for using the Qumulo REST API or qq CLI to view the creation dates for access keys and revoke any pair at their discretion.

For more information, see Listing S3 Access Keys for a Qumulo Cluster (page 0).

O Note

- To facilitate key rotation, each user identity (page 209) can have at most two S3 access key pairs associated with it. It is a good practice to delete a user's old access key after you create a new one and test that the new key works.
- If you revoke an access key pair, it isn't possible to restore it. Before you revoke an access key pair, ensure that no critical applications depend on it.

Creating S3 Access Keys for a Qumulo Cluster

To make S3 API requests to a Qumulo cluster as a specific user, you must create an S3 access key pair for that user identity (page 209) by using the Qumulo REST API or qq CLI.

To create S3 access keys, you must have an administrator account or have .

To Create an Access Key by Using the qq CLI

To create an S3 access key for a particular user identity (page 209), use the qqs3_create_access_key command and specify an identity. For example:

```
$ qq s3_create_access_key my_identity
```

You can specify an identity by using:

- · A name, optionally qualified with a domain prefix:
 - ∘ ad:MY NAME
 - AD\MY NAME
 - ∘ local:MY NAME
 - MY NAME
- An Active Directory Security Identifier. For example: SID:S-1-1-0
- A Qumulo auth ID, Qumulo Core's common representation for identities, in the form of a numeric identifier. For example: auth id:513

A Important

Currently, it isn't possible to associate an S3 access key with a POSIX group ID (GID).

The following is example output from the command.

```
{
   "access_key_id": "000000000001fEXAMPLE",
   "creation_time": "2022-12-12T21:37:53.553457928Z",
   "owner": {
        "auth_id": "501",
        "domain": "LOCAL",
        "gid": null,
        "name": "guest",
        "sid": "S-0-1-23-4567890123-456789012-345678901-234",
        "uid": null
},
   "secret_access_key": "TEIT4liMZ8A32iI7JXmqIiLWp5co/jmkjEXAMPLE"
}
```

In this example, the access key id is 00000000001fEXAMPLE and the secret access key is TEIT4liMZ8A32iI7JXmqIiLWp5co/jmkjEXAMPLE.

▲ Important

After Qumulo Core initially creates your secret access keys, it never logs or displays them again. If you lose your secret access key, it isn't possible to recover it and you must create a new access key pair.

To Create an S3 Access Key by Using the Qumulo REST API

Send a POST request to the /v1/s3/access-keys/ endpoint with the following body. You must include at least one of the following keys:

- auth id
- sid
- · uid

For example:

```
{
   "user": {
    "sid": "S-0-1-23-4567890123-456789012-345678901-234"
   }
}
```

The following is example output from the response.

```
{
   "access_key_id": "000000000001fEXAMPLE",
   "creation_time": "2022-12-12T21:37:53.553457928Z",
   "owner": {
        "auth_id": "501",
        "domain": "LOCAL",
        "gid": null,
        "name": "guest",
        "sid": "S-0-1-23-4567890123-456789012-345678901-234",
        "uid": null
},
   "secret_access_key": "TEIT4liMZ8A32iI7JXmqIiLWp5co/jmkjEXAMPLE"
}
```

In this example, the access key id is 00000000001fEXAMPLE and the secret access key is TEIT4liMZ8A32iI7JXmqIiLWp5co/jmkjEXAMPLE.

A Important

After Qumulo Core initially creates your secret access keys, it never logs or displays them again. If you lose your secret access key, it isn't possible to recover it and you must create a new access key pair.

Listing S3 Access Keys for a Qumulo Cluster

You can list every S3 access key that your Qumulo cluster knows, along with the identities associated with the key and the key creation times, by using the Qumulo REST API or qq CLI.

To list S3 access keys, you must have the PRIVILEGE S3 BUCKETS READ privilege.

O Note

Qumulo Core doesn't list access keys in any particular order. To sort keys according to fields such as creation time or owner you must process or filter the response.

To List S3 Access Keys by Using the qq CLI

To list the S3 access keys that your Qumulo cluster knows, use the qq
 s3 list access keys command:

The following is example output from the command. All times are in the UTC time zone.

• For JSON output, use the -- json flag.

The following is example output from the command. The command returns a single JSON object that contains the combined responses from calls to the $\frac{v1}{s3}/access-keys$ Qumulo REST API endpoint.

```
{
  "entries": [
      "access_key_id": "00000000001fEXAMPLE",
      "creation time": "2022-12-12T21:37:53.553457928Z",
      "owner": {
        "auth id": "501",
        "domain": null,
        "gid": null,
        "name": null,
        "sid": null,
        "uid": null
      }
    },
  ],
  "paging": {
    "next": null
  }
}
```

To List S3 Access Keys by Using the Qumulo REST API

To list the S3 access keys that your Qumulo cluster knows, send a GET request to the /v1/s3/access-keys/ endpoint.

O Note

To restrict the number of returned results, up to the maximum of 10,000 access keys (this is the default limit), include the optional limit query parameter in the request.

The following is example output from the response. The entries list contains the access keys, limited to the first 10,000. The paging.next field contains the URI to which you can send a GET request to retrieve the next page of access keys. By making GET requests with all returned paging.next values, you can iterate over all of the access keys in the cluster.

```
"entries": [
      "access_key_id": "00000000001fEXAMPLE",
      "creation time": "2022-12-12T21:37:53.553457928Z",
      "owner": {
        "auth_id": "501",
        "domain": null,
        "gid": null,
        "name": null,
        "sid": null,
        "uid": null
      }
    },
  ],
  "paging": {
    "next": null
  }
}
```

Revoking S3 Access Keys for a Qumulo Cluster

To revoke an S3 access key, you must delete the access key from your Qumulo cluster. You can delete an S3 access key by using the Qumulo REST API or qq CLI.

To revoke an access key, you must have the PRIVILEGE S3 BUCKETS WRITE privilege.

To Delete an S3 Access Key by Using the qq CLI

Use the qq s3_delete_access_key command and specify the access key ID. For example:

```
$ qq s3_delete_access_key \
--id 0000000001fEXAMPLE
```

To Delete an S3 Access Key by Using the Qumulo REST API

Send a DELETE request to the /v1/s3/access-keys/<access-key-id> Qumulo REST API endpoint and specify the access key ID.

Configuring Active Directory (AD) for S3

If you create S3 access keys for users that exist in an AD domain that has a trust relationship with the joined domain, you must append that domain's base DN to the base DN in your Qumulo cluster's AD configuration.

To append the trusted base DN to the base DN in use—with a semicolon (;) separating the two—use the Web UI or the qq ad_reconfigure command. For example:

```
$ qq ad_reconfigure \
   --base-dn 'CN=Users,DC=joined_domain,DC=example,DC=com;CN=Users,DC=trusted_domai
n,DC=example,DC=com'
```

For more information, see Configuring Cross-Domain Active Directory Trusts (page 199)

Creating and Managing S3 Buckets in Qumulo Core

This section explains how to create and manage S3 buckets for a Qumulo cluster. These buckets expose a part of your Qumulo file system to applications that use the Amazon S3 API.

You can create and work with S3 buckets by using the Qumulo REST API or qq CLI. You can also use the S3 API directly.

Prerequisites

To create and manage S3 buckets by using the Qumulo REST API or qq CLI, you need the following role-based access control (RBAC) privileges:

• PRIVILEGE S3 BUCKETS WRITE: Create and delete S3 buckets

O Note

If you perform create and delete operations on directories by using the qq CLI, you also need this privilege.

PRIVILEGE_S3_BUCKETS_READ: List S3 buckets

To create and manage S3 buckets by using the S3 API, you also need:

- A valid Qumulo S3 access key (page 209)
- A configured AWS CLI (page 0)

How S3 Buckets Map to the Qumulo File System

An S3 bucket exposes a portion of your Qumulo file system to applications that use the Amazon S3 API.

The bucket root directory (or bucket root) is the directory to which you attach an S3 bucket. All files under the bucket root directory (and all of its subdirectories) are objects in the bucket. The presence of the slash (/) in objects' keys determines the directory hierarchy.

A Important

Because S3 buckets can use any directory in the file system as a root directory, the same file can be an object in multiple buckets.

How the Qumulo File System Determines Object Keys

The *object key* in a Qumulo S3 bucket is its file system path, relative to the bucket's root directory. Only objects that are directories have a trailing slash (/) in their keys.

The following example shows the contents of a Qumulo file system.

In this example, if you have the S3 bucket bucket1 with its root directory at /application-data/deployment/">/application-data/deployment/, the bucket contains objects with the following keys:

- data1.dat
- data2.dat

However, if you have the S3 bucket bucket2 with its root directory at /application-data/, the bucket contains objects with the following keys:

- develop.dat
- processing/
- deployment/data1.dat
- deployment/data2.dat

O Note

In this example:

- Both buckets contain /application-data/deployment/data1.dat and /application-data/deployment/data2.dat as objects.
- The processing/object in bucket2 has a trailing slash because it is a directory.

How to Name an S3 Bucket

When you create an S3 bucket, you name it. A bucket's name doesn't need to be related to its root directory.

Except for names that contain the period (.), Qumulo Core accepts all names that conform to the following Amazon S3 bucket naming rules.

- Bucket names must be between 3 and 63 characters long.
- Bucket names can consist only of lowercase ASCII letters, numbers, and hyphens (-).

Bucket names must start with a letter or a number.

How to Choose a Bucket Root

You specify the bucket root directory (page 217) depending on how you create your S3 bucket.

- When you create an S3 bucket by using the Qumulo REST API or qq CLI, you can choose a directory to use as the bucket root.
- When you create an S3 bucket by using the CreateBucket S3 API action, the API creates a new directory with the same name as the bucket under the default bucket directory prefix. For more information, see Configuring the Default Bucket Directory Prefix for S3 Buckets (page 0).
- If you don't specify a directory, the Qumulo REST API and qq CLI use the default bucket directory prefix (page 0).

The user that creates a new directory for a new bucket owns the directory. For more information, see Managing Access to S3 Buckets in a Qumulo Cluster (page 226).

Creating S3 Buckets

You can create an S3 bucket by using the Qumulo REST API or qq CLI. You can also use the S3 API directly.

While the Qumulo REST API and qq CLI let you use an existing directory as the new bucket root, the S3 API always creates a new directory for the bucket root.

▲ Important

- All S3 buckets in a Qumulo cluster share the same namespace: It isn't possible to create two buckets with the same name, even if they use different directories as their bucket root.
- · All S3 buckets must follow the bucket naming rules (page 0).

Configuring the Default Bucket Directory Prefix for S3 Buckets

The default bucket directory prefix is the directory under which Qumulo Core creates new bucket root directories when it creates S3 buckets by using the CreateBucket S3 API action or when you create an S3 bucket without specifying a directory by using the Qumulo REST API or qq CLI.

By default, the default bucket directory prefix for newly created buckets is the cluster's root directory (/). Thus, if you create a bucket named my-bucket, its root directory is /my-bucket.

- To view the current default bucket directory prefix by using the Qumulo REST API or qq
 CLI, you need the PRIVILEGE_S3_BUCKETS_READ privilege.
- To change the default bucket directory prefix, you need the PRIVILEGE_S3_BUCKETS_WRITE privilege.

To Configure the Default Bucket Directory Prefix by Using the qq CLI

1. To view the current default bucket directory prefix, use the qq s3_get_settings command.

The following is example output from the command.

```
{"enabled": true, "base_path": "/buckets/", ...}
```

2. To change the setting, use the qq s3_modify_settings command and specify the new default bucket directory prefix. In the following example, we specify /buckets.

```
$ qq s3_modify_settings \
--base-path /buckets
```

Creating an S3 Bucket by Using the qq CLI

To create an S3 bucket by using the Qumulo REST API or qq CLI, you need the PRIVILEGE_S3_BUCKETS_WRITE privilege.

A Important

In Qumulo Core 6.0.1.1 (and higher), the qq CLI command changed from s3_create_bucket to s3_add_bucket and the flag for specifying the directory path has changed from --path to --fs-path.

When you use the qq CLI to create a bucket, you can use a new or existing directory as the bucket root.

O Note

If an entry with the specified name or directory already exists, or if you don't have permission to create a directory, the command returns an error. For more information, see Configuring the Default Bucket Directory Prefix for S3 Buckets (page 0).

To create a new, empty bucket from the default bucket directory prefix (page 0), use the
 qq s3_add_bucket
 command and specify the bucket name. For example:

```
$ qq s3_add_bucket \
--name my-bucket
```

Qumulo Core creates a new directory named my-bucket under the default bucket directory prefix.

• To create a bucket from an existing directory, use the qq s3_add_bucket command and specify the bucket name and the directory path. For example:

```
$ qq s3_add_bucket \
   --name my-bucket \
   --fs-path /products/web/appliances/
```

• To create a bucket for a path that doesn't exist yet, specify the name and path and add the --create-fs-path flag. For example:

```
$ qq s3_add_bucket \
   --name my-bucket \
   --fs-path /products/web/appliances/ \
   --create-fs-path
```

Creating an S3 Bucket by Using the S3 API

Use the aws s3api create-bucket command and specify the bucket name. This command uses the CreateBucket S3 API action. For example:

```
$ aws s3api create-bucket \
--bucket my-bucket
```

Qumulo Core creates the bucket root directory under the default bucket directory prefix (page 0) and names it the same as the bucket. In this example, if the default bucket directory prefix is /buckets/, the new bucket root directory is /buckets/my-bucket/.

O Note

When you use the CreateBucket S3 API action with the LocationConstraint parameter, the Qumulo S3 API supports only the local region.

Configuring S3 Buckets

You can view and modify the settings for individual buckets by using the Qumulo REST API or qq CLI.

Although you can configure global settings, such as the default bucket directory prefix (page 0) for S3 buckets, the only individual S3 bucket setting that you can configure in Qumulo Core is anonymous access. For more information, see Enabling Anonymous Access for an S3 Bucket (page 0).

- To view the current bucket configuration by using the Qumulo REST API or qq CLI, you need the PRIVILEGE_S3_BUCKETS_READ privilege.
- To change the bucket configuration, you need the PRIVILEGE_S3_BUCKETS_WRITE privilege.

Viewing the Current S3 Bucket Configuration by Using the qq CLI

Use the qq s3 get bucket command and specify the bucket name. For example:

```
$ qq s3_get_bucket \
--name my-bucket
```

The following is example output from the command. All times are in the UTC time zone..

```
{
  "anonymous_access_enabled": false,
  "creation_time": "2022-12-20T19:42:26.833076147Z",
  "name": "my-bucket",
  "path": "/buckets/my-bucket"
}
```

Listing S3 Buckets

You can list all S3 buckets in your Qumulo cluster by using the Qumulo REST API or qq CLI. You can also use the S3 API directly.

To List S3 Buckets by Using the qq CLI

To list your S3 buckets by using the Qumulo REST API or qq CLI, you need the PRIVILEGE S3 BUCKETS READ privilege.

Use the qq s3 list buckets command.

The following is example output from the command. All times are in the UTC time zone.

```
name creation_time path
========
my-bucket 2022-12-13T22:18:01.406433425Z /my-bucket
```

For JSON output, use the -- ison flag.

The following is example output from the command. All times are in the UTC time zone. The JSON output contains an array named **Buckets** that contains the individual buckets as objects.

```
{
    "buckets": [
        {
            "anonymous_access_enabled": false,
            "creation_time": "2022-12-13T22:18:01.406433425Z",
            "name": "my-bucket",
            "path": "/my-bucket"
        }
    ]
}
```

To List S3 Buckets by Using the S3 API

Use the aws s3api list-buckets command. This command uses the ListBuckets S3 API action.

The following is example output from the command. All times are in the UTC time zone. The JSON output contains an array named Buckets that contains the individual buckets as objects.

```
{
    "Buckets": [
        {
            "Name": "my-bucket",
            "CreationDate": "2022-12-13T22:18:01.406Z"
        }
        }
        ]
}
```

Deleting S3 Buckets

You can delete an S3 bucket by using the Qumulo REST API or qq CLI. You can also use the S3 API directly.

While the Qumulo REST API and qq CLI let you choose whether to also delete the bucket root directory, the S3 API always deletes the bucket root directory.

O Note

Before you delete your S3 bucket, you must either let all in-progress upload operations for the bucket (UploadPart, PutObject, or CopyObject) complete or you must abort the operations.

Deleting an S3 Bucket by Using the qq CLI

To delete an S3 bucket by using the Qumulo REST API or qq CLI, you need the PRIVILEGE_S3_BUCKETS_WRITE privilege.

When you use the qq CLI to delete a bucket, you can choose to also delete the bucket root directory.

• To delete an S3 bucket, but not its root directory, use the qq s3_delete_bucket command and specify the bucket name. For example:

```
$ qq s3_delete_bucket \
--name my-bucket
```

This command doesn't delete the bucket root directory. It deletes all metadata related to the bucket from your Qumulo cluster.

If any of the following conditions are true, the command returns an error:

- The specified bucket doesn't exist.
- You don't have the PRIVILEGE_S3_BUCKETS_WRITE privilege.
- The bucket has in-progress upload operations (UploadPart, PutObject, or CopyObject).
- To delete a bucket together with its root directory, use the qq s3_delete_bucket, specify the bucket name, and use the --delete-root-dir flag. For example:

```
$ qq s3_delete_bucket \
   --delete-root-dir \
   --name my-bucket
```

If any of the following conditions are true, the command returns an error:

- You don't have permission to delete the bucket root directory.
- The bucket root directory isn't empty.

Deleting an S3 Bucket by Using the S3 API

Use the aws s3api delete-bucket command and specify the bucket name. This command uses the DeleteBucket S3 API action. For example:

```
$ aws s3api delete-bucket \
  --bucket my-bucket
```

This command deletes the bucket root directory and all metadata related to the bucket from your Qumulo cluster.

If any of the following conditions are true, the command returns an error:

- · The specified bucket doesn't exist.
- · You don't have permission to delete the bucket root directory.
- · The bucket root directory isn't empty.
- The bucket has in-progress upload operations (UploadPart, Put0bject, or Copy0bject).

Managing Access to S3 Buckets in a Qumulo Cluster

This section explains how to manage access to S3 buckets in a Qumulo cluster.

Managing user access to S3 buckets in a Qumulo cluster is very similar to managing access to SMB shares and NFS exports, with the following exceptions:

- To let a user access S3 buckets in the cluster, you must assign an S3 access key (page 209)
 to the user. Alternatively, you can create presigned URLs (page 0) or enable read-only,
 anonymous access (page 227) for the entire S3 bucket.
- Because Amazon S3 provides access to an entire S3 bucket but Qumulo Core provides access to individual files and directories, a bucket might behave differently from user expectations.

O Note

There are no bucket-level permissions in Qumulo Core, only file-level permissions. To configure an S3 bucket in Qumulo Core to work more like an Amazon S3 bucket, use inheritable access control entries (ACEs) to imitate bucket-level permissions (page 0).

How S3 Bucket Permissions Work in Qumulo Core

To process an S3 API request, Qumulo Core performs one or more file system operations. Qumulo Core processes these operations by checking the user's access against the access control lists (ACLs) for each file that is part of the request.

For authenticated requests signed with Amazon Signature Version 4, Qumulo Core maps the access key ID (page 209) in the request to its corresponding auth ID (page 211), and then processes the request as that user. Qumulo Core processes unsigned, anonymous requests as the Guest user.

While Qumulo Core processes an S3 request, the ownership of any newly created files and directories belongs to the user that makes the request. These files and directories inherit access control entries (ACEs) from their parents (this process is the same for all protocols).

Granting Access to S3 Buckets by Using Presigned URLs

To let trusted users perform S3 API actions—such as **GetObject** or **UploadPart** —as if using your user account, you can generate a *presigned URL* (also known as *query parameter authentication*), associate the URL with specific API actions, and then share it with trusted users. Every presigned URLs has a configurable expiration time that ensures that the URL stops working at the configured time.

For more information, see Authenticating Requests: Using Query Parameters (AWS Signature Version 4) in the Amazon Simple Storage Service API Reference.

O Note

Qumulo Core accepts only presigned requests that use the PUT, GET, HEAD, and DELETE HTTP methods. Qumulo Core rejects presigned requests for POST requests, such as the following:

- AbortMultipartUpload
- CompleteMultipartUpload
- CreateMultipartUpload
- DeleteObjects

To create a presigned URL, use the AWS CLI **presign** command. In the following example, the presigned URL expires in 10 minutes (600 seconds).

```
$ aws2 s3 presign s3://my-bucket/my-file.txt \
--endpoint-url https://203.0.113.0:9000 \
--profile my-qumulo-profile \
--expires-in 600
```

The following is example output from the command, with line breaks inserted for readability. The X-Amz-Expires header is set to 10 minutes.

```
https://203.0.113.0:9000/my-bucket/my-file.txt?

X-Amz-Algorithm=AWS4-HMAC-SHA256

&X-Amz-Credential=0000000000000000888527%2F20230217%2Fus-east-1%2Fs3%2Faws4_request

&X-Amz-Date=20230217T205559Z

&X-Amz-Expires=600

&X-Amz-SignedHeaders=host

&X-Amz-Signature=141fa5b10caaa8575ba9c065d2270a24ce14b2ff58bb2c2e98382c76297b21ee
```

Enabling Anonymous Access for an S3 Bucket

In certain cases, it might be more practical to allow anonymous (unauthenticated) requests to access the contents of S3 buckets, for example, if you want to let users access objects from the S3 bucket by using a web browser or if the number of users who need read access is very large. When you enable anonymous access to an S3 bucket, your users can perform read-only S3 operations without authenticating their requests.

A Important

Anonymous requests can never perform modifying operations. Qumulo Core requires all modifying operations on an S3 bucket to be authenticated.

When you enable anonymous access for an S3 bucket, Qumulo Core performs all anonymous requests as the **Guest** user. The **Guest** user is a member of the **Everyone** group, but not of the **Users** group.

To ensure that anonymous requests have permission to read files in a bucket, grant read permission to the Everyone group or to the Guest user. For more information, see Imitating Bucket-Level Read-Only Access (page 0).

O Note

If a file's ACL doesn't allow reads for the Guest user, an anonymous request can't read the file.

- To view the current bucket configuration by using the Qumulo REST API or qq CLI, you need the PRIVILEGE_S3_BUCKETS_READ privilege.
- To change the bucket configuration, you need the PRIVILEGE_S3_BUCKETS_WRITE privilege.

To Enable Anonymous Access to an S3 Bucket by Using the qq CLI

1. To view the current bucket configuration, use the qq s3_get_bucket command and specify the bucket name. For example:

```
$ qq s3_get_bucket \
  --name my-bucket
```

The following is example output from the command. All times are in the UTC time zone.

```
{
  "anonymous_access_enabled": false,
  "creation_time": "2022-12-20T19:42:26.833076147Z",
  "name": "my-bucket",
  "path": "/buckets/my-bucket"
}
```

2. Anonymous access to S3 buckets is disabled by default. To enable anonymous access, use the qq s3_modify_bucket command, specify the bucket name, and use the --enable-anonymous-access flag.

3. To disable anonymous access, use the qq s3_modify_bucket command, specify the bucket name, and use the --disable-anonymous-access flag.

Using Inheritable ACEs to Imitate Bucket-Level Permissions

To grant multiple users access to all paths in a bucket and ensure that newly created directories inherit the correct permissions, use inheritable access control entries (ACEs).

In Amazon S3, permission to read objects from —and write objects to— an S3 bucket applies to the entire bucket. In Qumulo Core, each object key (page 0) corresponds to a file path relative to a bucket's root directory. Qumulo Core grants permissions for individual files and directories.

When users create objects in an S3 bucket in a Qumulo cluster, they might also create new directories. The user that creates these directories owns them. However, without the correct access control entries (ACEs) in your bucket, these directories might have restrictive permissions that prevent other users from creating objects with the same prefix.

How Permissions with Inheritable ACEs Work

Access control entries (ACEs) control the permissions that users have for files and directories in a Qumulo cluster. When you add ACEs to a directory and mark them as *inheritable*, all new files and directories created in that directory inherit those ACEs and pass them on.

You can use inheritable ACEs to:

- Imitate bucket-level permissions by ensuring that any files and directories that your users create in an S3 bucket receive the same permissions.
 - To make all paths in an S3 bucket inherit the same set of ACEs, add the ACEs to the bucket's root directory and mark them as inheritable.
- · Configure default permissions for newly created buckets.

To make a set of ACEs the default for buckets that your users create by using the S3 API, add the ACEs to the default bucket directory prefix.

To add ACEs to a directory, use the qq CLI or use the File Explorer on a Windows client with a mapped SMB share that contains the directory.

O Note

Adding inheritable ACEs to a directory doesn't affect any files that already exist in that directory. For more information, see To Recursively Add a New ACL (with Multithreading) on Oumulo Core.

Imitating Bucket-Level Permissions by Using the qq CLI

The following sections show how to use the qq CLI to imitate bucket-level permissions by adding inheritable ACEs.

Imitating Bucket-Level Read-Write Access

Use the qq fs_modify_acl command. In the following example, we add the access control entry
(ACE) to the bucket whose root directory is /buckets/my-bucket for the user group MyWriters.

```
$ qq fs_modify_acl \
    --path /buckets/my-bucket add_entry \
    --trustee MyWriters \
    --type Allowed \
    --flags 'Container inherit' 'Object inherit' \
    --rights 'Delete child' 'Execute/Traverse' 'Read' 'Write file'
```

The ACE imitates bucket-level read-write access for a user or group of users.

Imitating Bucket-Level Read-Only Access

Use the qq fs_modify_acl command. In the following example, we add the access control entry (ACE) to the bucket whose root directory is /buckets/my-bucket for the user group MyReaders:

```
$ qq fs_modify_acl
--path /buckets/my-bucket add_entry \
--trustee MyReaders \
--type Allowed \
--flags 'Container inherit' 'Object inherit' \
--rights 'Execute/Traverse' 'Read'
```

The ACE imitates bucket-level read-only access for a user or group of users.

```
Type Flags Rights
====== Allowed Object inherit, Container inherit Execute/Traverse, Read
```

Imitating Bucket-Level List-Only Access

Use the qq fs_modify_acl command. In the following example, we add two access control entries (ACEs) to the bucket whose root directory is /buckets/my-bucket for the user group MyListers.

```
$ qq fs_modify_acl
--path /buckets/my-bucket add_entry \
--trustee MyListers \
--type Allowed \
--flags 'Container inherit' \
--rights 'Execute/Traverse' 'Read'
```

```
$ qq fs_modify_acl
   --path /buckets/my-bucket add_entry \
   --trustee MyListers \
   --type Allowed \
   --flags 'Object inherit' \
   --rights 'Read attr'
```

The two ACEs imitate bucket-level list-only access for a user or group of users:

Managing Multipart S3 Uploads in Qumulo Core

This section explains how multipart S3 uploads affect usable capacity on a Qumulo cluster and how to abort and clean up multipart uploads manually or automatically.

Qumulo Core supports the multipart upload functionality of the S3 API, which lets you upload objects to a bucket in parts and then, at a later time, combine these parts into a single object.

1 Note

For objects above a certain size (typically, larger than 100 MiB), applications often use the multipart S3 uploads, rather than the Put0bject S3 API action. The limitation for the Put0bject action is 5 GiB. For more information about how Qumulo handles this type of operation, see System-Initiated Multipart S3 Uploads (page 0).

Prerequisites

To manage multipart S3 uploads by using the qq CLI, you need the following role-based access control (RBAC) privileges:

- PRIVILEGE S3 SETTINGS WRITE: Configure frequency of multipart upload cleanup
- PRIVILEGE_S3_UPLOADS_READ: List multipart uploads
- PRIVILEGE_S3_UPLOADS_WRITE : Abort multipart uploads

How Multipart S3 Uploads Affect Usable Capacity on a Qumulo Cluster

The following conditions are true for multipart S3 uploads in Qumulo Core.

- To let you resume large uploads in the event of an outage, Qumulo Core stores data on the cluster durably.
- Multipart upload data isn't visible in the Qumulo file system, and isn't included in file system snapshots, until you complete the upload successfully by making a call to the CompleteMultipartUpload S3 API.

Note

When you view the breakdown of a Qumulo cluster's capacity by using the Qumulo Core Web UI, REST API, or qq CLI, Qumulo Core doesn't distinguish between capacity that the file system and incomplete multipart uploads use.

• Qumulo Core doesn't delete multipart data unless it aborts and cleans up the multipart upload automatically (page 0) or you abort and clean up the multipart upload manually (page 0).

To check how much space incomplete multipart uploads use on your cluster, you can list the uploads by using the Qumulo REST API or qq CLI. For more information, see Listing Multipart Uploads (page 0).

How System-Initiated Multipart S3 Uploads Work

Occasionally, when you list your multipart uploads (page 0), you might see uploads that you didn't initiate. These are *system-initiated uploads* which Qumulo Core uses for PutObject and CopyObject S3 API actions for objects that exceed a certain size.

If Qumulo Core encounters an error while performing a system-initiated upload, it attempts to abort the upload and clean up the partial upload data immediately.

However, if Qumulo Core is unable to clean up the incomplete upload data immediately, it cleans up the incomplete upload data in the background, according to the expiry interval (page 0).

1 Note

The process for background clean-up after incomplete and user-initiated uploads is the same. For more information, see Aborting and Cleaning Up Multipart S3 Uploads Automatically (page 0).

Listing Incomplete Multipart S3 Uploads

You can list the incomplete multipart uploads for a single S3 bucket by using the Qumulo REST API or qq CLI.

Note

- If you use the ListMultipartUploads S3 API action, the system doesn't show system-initiated uploads (page 0) or how much space the uploads use on your cluster.
- If you use the Qumulo REST API or qq CLI, Qumulo Core shows system-initiated uploads (page 0) and how much space each upload uses on your cluster.
- To list incomplete uploads by using the qq CLI, use the s3_list_uploads command and specify the bucket name. For example:

```
$ qq s3_list_uploads \
  --bucket my-bucket
```

To list incomplete uploads by using the REST API, send a GET request to the /v1/s3/buckets/<bucket-name>/uploads/ endpoint and specify the bucket name.

The output from the qq CLI and REST API is the same. The following example output is a single JSON object that contains the list of objects for the specified bucket. The list shows information for each multipart S3 upload, including:

- · When each upload was initiated
- · Which identity initiated the upload
- · When the upload received data last
- How much space the upload uses on the cluster—by data, by metadata, and in total—in units of blocks (4,096 bytes per block)

```
"uploads": [
      "bucket": "my-bucket",
      "completing": false,
      "datablocks": "16384",
      "id": "00000000example1",
      "initiated": "2023-03-02T19:01:00.446468848Z",
      "initiator": {
        "auth id": "500",
        "domain": null,
        "gid": null,
        "name": null,
        "sid": null,
        "uid": null
      "key": "deployment/data1.dat",
      "last_modified": "2023-03-02T19:03:37.209271702Z",
      "metablocks": "3",
      "system initiated": false,
      "total_blocks": "16387"
    },
    {
      "bucket": "my-bucket",
      "completing": false,
      "datablocks": "24576",
      "id": "00000000example2",
      "initiated": "2023-03-02T19:09:04.530619255Z",
      "initiator": {
        "auth id": "500",
        "domain": null,
        "gid": null,
        "name": null,
        "sid": null,
        "uid": null
      },
      "key": "release.dat",
      "last_modified": "2023-03-02T19:09:06.436699236Z",
      "metablocks": "4",
      "system_initiated": true,
      "total blocks": "24580"
    }
 ]
}
```

Aborting and Cleaning Up Multipart S3 Uploads Automatically

Qumulo Core automatically aborts and cleans up an incomplete multipart S3 if the upload doesn't receive any data after the configured *expiry interval* (1 day by default).

When Qumulo Core removes a multipart upload, it frees up the space that the upload uses on the cluster. You can configure the expiry interval by using the Qumulo REST API or qq CLI.

To configure the expiry interval for all current and future multipart uploads by using the qq CLI, use the s3_modify_settings command and the --multipart-upload-expiry-interval flag and specify one of the following:

- The string never.
- A string in the format <quantity><units> (without a space), where <quantity> is a positive integer less than 100 and <units> is one of the following strings:
 - days
 - hours
 - minutes
 - months
 - weeks

In the following example, we instruct Qumulo Core to abort and clean up uploads that haven't received data in more than 30 days.

```
$ qq s3_modify_settings \
--multipart-upload-expiry-interval 30days
```

In the following example, we disable automatic cleanup.

```
$ qq s3_modify_settings \
  --multipart-upload-expiry-interval never
```

Aborting or Cleaning Up Multipart S3 Uploads Manually

Use the Qumulo REST API or qq CLI to abort and clean up the upload. You need the bucket name and upload ID. For more information about looking up this information, see Listing Incomplete Multipart S3 Uploads (page 0).

O Note

If you are an administrative user or the user who initiated the upload, you can use the AbortMultipartUpload S3 API action. In addition to the bucket name and upload ID, you also need the object key for the upload.

• To abort an upload by using the qq CLI, use the s3_abort_upload command and specify the upload ID. For example:

```
$ qq s3_abort_upload \
   --bucket my-bucket \
   --upload-id 00000000example
```

To abort an upload by using the Qumulo REST API, send a DELETE request to the /v1/s3/buckets/<bucket-name>/uploads/<upload-ID> endpoint and specify the upload ID. For example:

DELETE /v1/s3/buckets/my-bucket/uploads/00000000example

There is no response body for both the qq CLI and REST API. Qumulo Core returns a 204 No Content status code when the upload is aborted or the cleanup is complete.

Supported Functionality and Known Limits for S3 in Qumulo Core

This section documents Qumulo Core support for S3 API functionality and S3 API limits.

Supported S3 API Actions

The following table lists the S3 API actions that Qumulo Core supports and the version from which support begins. For the full list of S3 API actions, see Actions in the Amazon Simple Storage Service API Reference.

O Note

The S3 API became generally available in Qumulo Core 5.3.3. This guide doesn't document enabling or using API actions that became available with preview functionality in versions of Qumulo Core lower than 5.3.3.

API Action	Supported from Qumulo Core Version
AbortMultipartUpload	5.3.3
CompleteMultipartUpload	5.3.3
CopyObject	5.3.3
CreateBucket	5.2.3
CreateMultipartUpload	5.3.3
DeleteBucket	5.2.4
DeleteObject	5.2.1
DeleteObjects	5.2.2
GetBucketLocation	5.1.2
Get0bject	5.0.4
HeadBucket	5.1.2
HeadObject	5.0.4
ListBuckets	5.0.4

API Action	Supported from Qumulo Core Version
ListMultipartUploads	5.3.3
ListObjects	5.0.5
ListObjectsV2	5.0.4
ListParts	5.3.3
Put0bject	5.2.1
UploadPart	5.3.3
UploadPartCopy	6.1.0.1

Unsupported S3 Functionality

The following table lists some of the S3 API functionality that Qumulo Core doesn't support.

Unsupported Feature	Description
BitTorrent	
Bucket ACLs	For comparable functionality, use inheritable access control entries (ACEs) (page 0).
Bucket lifecycle config- urations	
Bucket notifications	
Control of server-side encryption	All Qumulo Core data is encrypted at rest. You can't control this functionality by using the S3 API.
Logging controls	
Multi-chunk payload signing	Qumulo Core doesn't support the streaming version of Amazon Signature Version 4 (SigV4), only the single-chunk version.
Object locks	
Object tagging and custom object metadata	

Unsupported Feature	Description
Object versioning	Qumulo objects have one object version. To preserve previous object contents in Qumulo Core, use snapshots.
Policies	
Signature Version 2	Qumulo Core supports only SigV4 signatures.
Storage classes	Qumulo Core doesn't use the storage class concept. All objects have the same storage class status.
Retention policies	
Temporary access credentials	
Virtual-hosted bucket addressing	Qumulo Core supports only path-style bucket addressing.
Web hosting configura- tion	

S3 API Limitations

This section describes the most important S3 API limitations in Qumulo Core.

Bucket Addressing Style

Because Qumulo Core supports only path-style bucket addressing, you must configure your client applications to use path-style addressing to send S3 API requests to a Qumulo cluster. For more information, see Configuring the AWS CLI for Use with Qumulo Core (page 0).

ETags

RESTful APIs, such as the S3 API, use HTTP ETags to identify different resource versions.

- · Qumulo Core uses a proprietary mechanism to generate an object's ETag.
- · Amazon S3 uses the MD5 checksum of an object's contents as its ETag.

A Important

Well-behaved applications shouldn't attempt to interpret the contents of an ETag. However, certain applications do assume that S3 object ETags contain the MD5 checksum of the object's contents. Such applications might not function properly with the Qumulo S3 API.

Listing Objects

The S3 API supports listing objects in a bucket by using the List0bjects and List0bjectsV2 API actions.

Function	Qumulo Core	Amazon S3
Returning results	Consistent but non-alphabet- ical order	Alphabetical order, by object key
Arbitrary prefix	Partial support for Prefix, only if Prefix is a path to a file or directory under the bucket root directory (page 217)	Prefix limits results to object keys that begin with the prefix
Arbitrary delimiter	Only the slash (/) character can act as Delimiter	Delimiter groups results into common prefixes

Note

Although Qumulo Core supports Prefix and Delimiter partially, it supports the most common use case—listing the contents of S3 buckets as a hierarchical file tree—fully.)))

Request Authentication

Qumulo Core supports authenticating requests by using only Amazon Signature Version 4. Most S3 client applications support this authentication type.

If your application attempts to use a previous Amazon signature version, you receive a 400 Bad Request response with the error code AuthorizationHeaderMalformed.

Comparison of Known Limits between S3 in Qumulo and Amazon

This section compares the Qumulo Core S3 API limits with native Amazon S3 limits.

Limits for S3 Buckets

Limit	Qumulo Core	Amazon S3
Maximum number of buckets	16,000	1,000
Maximum number of objects in one bucket	Nominally unlimited	Unlimited
Minimum bucket name length	3 characters	

Limit	Qumulo Core	Amazon S3
Maximum bucket name length	63 characters	

O Note

If all objects in a bucket are under the same directory—none of the object keys have the slash (/) character in them—the maximum number of objects in the bucket is limited to the maximum number of files in a directory. For more information, see Supported Configurations and Known Limits for Qumulo Core (page 0).

Limits for S3 Objects

Limit	Qumulo Core	Amazon S3
Minimum object size	0 bytes	
Maximum object size (by using PutObject)	5 GiB	
Maximum object size (by using MultipartUpload)	48.8 TiB (10,000 * 5 GiB)	5 TiB
Minimum object key length	1 character	
Maximum object key length	1,530 characters, if there are no slash (/) characters in the key	1,024 characters

Limits for S3 Multipart Uploads

Limit	Qumulo Core	Amazon S3
Minimum part ID	1	
Maximum part ID	10,000	
Minimum number of parts for each upload	1	
Maximum number of parts for each upload	10,000	

Limit	Qumulo Core	Amazon S3
Minimum part size	5 MiB (except for the last part o	f an upload)
Maximum part size	5 GiB	
Additional part size requirements	Must be a multiple of 4 KiB (4,096 bytes), except for the last part of an upload	_

Limits for S3 API Requests

Maximum Limit	Qumulo Core	Amazon S3
Object keys that DeleteObjects specifies	Nominally unlimited	1,000
Buckets that ListBuckets returns	16,000	1,000
Objects that ListObjects and ListObjectsV2 return	1,000	
Parts that ListParts returns	Unlimited	1,000
Uploads that ListMultipar- tUploads returns	1,000	

1 Note

DeleteObjects is subject to a 10 MiB request payload limit in Qumulo Core. This provides a practical upper limit on the number of object keys that the API action can specify.

In addition, the following API actions have the Qumulo-specific maximum payload size limit of 10 MiB.

- CompleteMultipartUpload
- CreateBucket
- DeleteObjects

Monitoring and Metrics

Enabling Cloud-Based Monitoring and Remote Support

This section explains how to enable Cloud-Based Monitoring and Remote Support for your Oumulo cluster.

A Important

To let the Qumulo Care team provide fast support when you need it most, we strongly recommend enabling both Cloud-Based Monitoring and Remote Support.

How Cloud-Based Monitoring Works

Enabling Cloud-Based Monitoring (page 248) lets the Qumulo Care team monitor your Qumulo cluster proactively.

A Important

Cloud-Based Monitoring *doesn't* collect file names, path names, client IP addresses, or account credentials.

Qumulo Care Response Times

We use a proprietary application that aggregates diagnostic cluster data and alerts the Qumulo Care team if an issue arises. Depending on the issue severity and cluster state, a member of the Qumulo Care team reaches out. The following table outlines Qumulo Care response times.

Severity Level	Service Availability	Response Time	Description	Example
0	24×7 2 hours	24×7	Outage, data un- availability, or data loss	The Qumulo cluster is down or there aren't enough online nodes to form a quorum.
			▲ Important For Severity 0 cases, call one of our toll-free numbers (page 246) or select Severity 0 when you open a case (page 246).	
1	24×7	2 hours	High business impact while the Qumulo cluster is still available and functional	A node is down but the cluster's quorum is still intact.
2	24×5	2 hours	A serious issue with an available workaround	When you use ls and dd commands from the same client, the performance is poor. The workaround is to mount to two different nodes and run the ls command for the first node and the dd command for the second.
3	24×5	6 hours	Poor or frustrating user experience	In the Web UI, a dialog box lingers for about five seconds after a change.
4	24×5	6 hours	Cosmetic issues or other concerns	The background color of a dialog box makes text difficult to read.

Ways to Get Help

The Qumulo Care team is always here to help you. You can contact us by using any of the following ways.

- See Browse the Qumulo Knowledge Base
- Open a case by filing a request or emailing us.
- Message us on Slack (if you are a current customer). For more information, see our Slack Tips and Tricks.
- Call one of our toll-free numbers:
 - North America: +1 855-577-7544 (page 0)
 - United Kingdom: +44 808-164-6656 (page 0)
 - Germany: +49 800-000-7047 (page 0)
 - Australia: +61 1800-954-952 (page 0)
 - U.S. Government Customers: +1 844-962-3777 (page 0)

How Remote Support Works

Enabling Remote Support (page 248) lets the Qumulo Care team access your Qumulo cluster solely to assist you with a software update or perform diagnostics or troubleshooting on your cluster from the command line.

When you install VPN keys in the /etc/openvpn directory, an authorized member of the Qumulo Care team uses SSH to connect to the epl.qumulo.com server and then uses SSH through a secure VPN connection to connect to your cluster (normally, this VPN connection is closed).

By default, the VPN tunnel remains open for four hours to allow members of the Qumulo Care team to perform operations such as uploading logs to monitor.qumulo.com or to a secured Amazon S3 bucket and sending diagnostic data to a private Amazon EC2 instance for analysis.

Note

Currently, Qumulo Core doesn't support VPN connections with IPv6.

You can configure the connection period and enable or disable Remote Support at any time.

What Data Gets Sent to Qumulo

Cloud-Based Monitoring and Remote Support let your cluster send the following detailed diagnostic data to Qumulo through an encrypted connection.

- Cluster name
- · Number of nodes in cluster
- · Hardware and software incidents
 - Drives

- CRC errors
- S.M.A.R.T. status alerts
- Capacity triggers
- Nodes
 - PSU failure
 - Fan failure
 - Recused node
 - Offline node
 - Unreachable cluster
- Qumulo Core
 - New process core dump
- · Configuration data (such as users, groups, SMB shares, and NFS exports)
- · Logs, stack traces, and code dumps

Prerequisites

Before you can use Cloud-Based monitoring and Remote Support, you must:

- Install VPN keys on your Qumulo cluster (page 0)
- Enable the following destination hostnames for TCP on port 443.

Hostname	Description
api.nexus.qumulo.com	Nexus monitoring
ep1.qumulo.com	Remote Support A Important If your organization has an intrusion detection device or a firewall that performs SSL or HTTPS deep-packet inspection, you must add an exception to the IP address that resolves to
	epl.qumulo.com. To identify this IP address, log in to your Qumulo cluster and run the nslookup epl.qumulo.com command.

Hostname	Description
missionq.qumulo.com	Cloud-Based Monitoring connectivity for Qumulo Core versions lower than 6.1.1
api.missionq.qumulo.com	Cloud-Based Monitoring connectivity for Qumulo Core 6.1.1 and higher
missionq-dumps.s3.amazon- aws.com	Proxy forwarding
monitor.qumulo.com	Cloud-Based Monitoring log uploads

Enabling Cloud-Based Monitoring

You can enable Cloud-Based Monitoring by using the Web UI or qq CLI.

To Enable Cloud-Based Monitoring by Using the Web UI

- 1. Log in to the Web UI.
- 2. Click Support > Qumulo Care.
- 3. On the Qumulo Care page, do the following:
 - a. In the Cloud-Based Monitoring section, click Edit.
 - b. Click Yes, I want Qumulo Cloud-Based Monitoring and then click Save.

If your configuration is valid, the Web UI shows the status Enabled | Connected.

To Enable Cloud-Based Monitoring by Using the gg CLI

- To enable Cloud-Based Monitoring, run the qq set_monitoring_conf --enabled command.
- To disable Cloud-Based Monitoring, run the qq set_monitoring_conf --disabled command.
- To check the status of Cloud-Based Monitoring, run the qq monitoring conf command.

Enabling Remote Support

You can enable Remote Support by using the Web UI or qq CLI.

To Enable Remote Support by Using the Web UI

- 1. Log in to he Web UI.
- 2. Click Support > Qumulo Care.
- 3. On the Qumulo Care page, do the following:

- a. In the Remote Support section, click Edit.
- b. Under Do you want to enable Qumulo Remote Support?, click Yes and then click Save.

If your configuration is valid, the Web UI shows the status **Enabled | Connected**.

To Enable Remote Support by Using the qq CLI

- To enable Remote Support, run the qq set_monitoring_conf --vpn-enabled command.
- To disable Remote Support, run the qq set_monitoring_conf --vpn-disabled command.
- To check the status of Remote Support, run the qq set_monitoring_conf command.

Connecting to Cloud-Based Monitoring and S3 by Using a Custom Proxy

This section explains how to connect to Cloud-Based Monitoring and S3 by using a custom proxy.

In Qumulo Core 2.6.4 (and higher), you can specify a custom proxy for both Cloud-Based Monitoring and S3 in different ways.

Configuring a Custom Proxy for Cloud-Based Monitoring and S3

The following examples show some common configurations for custom proxies by using the qq set_monitoring_conf command.

A Important

Remote VPN support isn't available when you connect to Cloud-Based Monitoring by using a custom proxy.

You can specify a custom proxy hostname and port *only* for Cloud-Based Monitoring. For example:

```
qq set_monitoring_conf
  --enabled
  --mq-proxy-host mq-proxy.example.com
  --mq-proxy-port 123
```

You can also specify a custom proxy hostname and port *only* for the S3 proxy (for the endpoint that Qumulo Core uses to store core dumps and logs). For example:

```
qq set_monitoring_conf
  --enabled
  --s3-proxy-host s3-proxy.example.com
  --s3-proxy-port 456
```

Finally, you can specify a custom proxy hostname and port for *both* Cloud-Based Monitoring and your S3 proxy. For example:

```
qq set_monitoring_conf
  --enabled
  --all-proxy-host mq-s3-proxy.example.com
  --all-proxy-port 789
```

Restoring the Default Values for Cloud-Based and Nexus Monitoring

This section explains how to set the default values for Cloud-Based Monitoring and Nexus Monitoring.

When you no longer connect to Cloud-Based Monitoring by using a custom proxy (page 250), you can use the qq CLI to restore the default values for Cloud-Based and Nexus Monitoring.

To Restore the Default Values for Cloud-Based and Nexus Monitoring by Using the qq CLI

1. Connect to your cluster by using SSH. For example:

```
ssh admin@203.0.113.0
```

2. Log in to Qumulo Core by using the administrative account. For example:

```
qq login -u admin -p mypassword
```

3. To restore the default values for Cloud-Based and Nexus Monitoring, use the qq set_monitoring_conf command.

In the following example, we:

- · Enable Cloud-Based Monitoring (MQ)
- · Enable Nexus monitoring
- · Disable HTTPS for the S3 proxy
- · Set the hostname and port for Cloud-Based Monitoring
- Disable the proxy for Cloud-Based Monitoring (by setting it to 0)
- · Set the monitoring polling interval to 60 seconds
- · Set the hostname and port for the S3 proxy
- · Specify the custom VPN for your organization, example.qumulo.com

```
qq set_monitoring_conf
   --enabled
   --nexus-enabled
   --s3-proxy-disable-https
   --mq-host missionq.qumulo.com
   --mq-port 443
   --mq-proxy-port 0
   --period 60
   --s3-proxy-host monitor.qumulo.com
   --s3-proxy-port 443
   --vpn-host example.qumulo.com
```

4. To confirm that Cloud-Based Monitoring is working correctly, log in to the Qumulo Core Web UI and then click **Support**.

If your configuration is valid, the Web UI shows the status **Enabled | Connected**.

Qumulo OpenMetrics API Specification

This section lists the names, types, labels, and descriptions for the metrics that Qumulo Core 5.3.0 (and higher) emits in OpenMetrics API format.

The Qumulo OpenMetrics API has a single endpoint that provides a complete view of point-in-time telemetry from Qumulo Core to monitoring systems. These systems, such as Prometheus, can consume the OpenMetrics data format that the Qumulo REST API emits without custom code or a monitoring agent. For more information about data formats, see your monitoring system's documentation.

Accessing Qumulo Metrics

Qumulo metrics are available at the following endpoint.

https://<my-cluster-hostname>:8000/v2/metrics/endpoints/default/data

You can configure a monitoring system that supports the OpenMetrics Specification to use bearer token authentication (page 41) to access this endpoint.

Metric Types

All Qumulo metrics belong to one of the following OpenMetrics types.

Metric Type	Description
counter	An integer that increases monotonically from zero, stored in <met-ric_name>_count .</met-ric_name>
	① Note During normal operation, the value of counter never decreases.
gauge	A value that represents a single integer (similar to counter), stored in <met-ric_name>.</met-ric_name>
	① Note During normal operation, the value of a gauge metric might increase or decrease.

Metric Type	Description
histogram	A representation of a series of <i>buckets</i> , where each bucket tracks values within a specific range. A histogram has a count field and a sum field, stored in metric_name>_count (the total number of samples) and metric_name>_sum (the sum of all samples). Qumulo Core emits a single bucket that contains all samples.
	☑ Tip You can use histogram metrics to keep track of averages by dividing the sum field by the count field.
info	Informational text about the system, stored in <metric_name>_info . An info metric always has a value of 1 and labels that contain detailed information.</metric_name>

For more information, see Metric Types in the OpenMetrics Specification.

Metric Labels

The OpenMetrics format allows for metric labeling for communicating additional information. To provide context for metrics, Qumulo Core emits metric-specific labels. For example, the name of a protocol operation or the url of a remote server. For more information, see Available Labels (page 262).

Available Metrics

The following table lists metric names, types, labels, and descriptions.

O Note

For Qumulo as a Service, all metrics with a node_id label are unavailable because they refer to specific hardware.

Metric Name	Metric Type	Labels	Suppor- ted from Qumulo Core Version	Description
qumulo	info (page 255)	max_drive_failuresmax_node_failuresnameplatformservice_modeluuidversion	5.3.0	Qumulo Core information, including the cluster name, cluster UUID, and the current Qumulo Core version
qumulo_node	info (page 255)	node_id (page 264)node_model	6.0.2	Information about the nodes in the cluster, including the node ID and the node model
qumulo_ad_netlogon_request _errors	counter (page 254)	domain_url (page 263)server_url (page 265)	5.3.0	The total number of Active Directory (AD) NETLOGON requests that resulted in an error
qumulo_ad_netlogon_request _latency_seconds	histogram (page 255)	domain_url (page 263)server_url (page 265)	5.3.0	The total latency for AD NETLO- GON requests
qumulo_ad_netlogon_requests	counter (page 254)	domain_url (page 263)server_url (page 265)	5.3.0	The total number of completed AD NETLOGON operations

Metric Name	Metric Type	Labels	Suppor- ted from Qumulo Core Version	Description
qumulo_cpu_max_temperature _celsius	gauge (page 254)	cpu (page 263)node_id (page 264)	5.3.1	The maximum temperature threshold for each physical CPU
qumulo_disk_endurance _percent	gauge (page 254)	disk_type (page 263)drive_bay (page 263)node_id (page 264)	5.3.1	The remaining disk endurance value for each disk in the cluster, ranging 100 (no disk wear) to 0 (disk is worn fully)
qumulo_disk_transport _errors	counter (page 254)	disk_type (page 263)drive_bay (page 263)node_id (page 264)	5.3.2	The total number of communication errors between the specified drive and its host.
qumulo_disk_uncorrectable _media_errors	counter (page 254)	disk_type (page 263)drive_bay (page 263)node_id (page 264)	5.3.2	The total number of uncorrectable errors on the specified drive's physical media.
qumulo_cpu_temperature _celsius	gauge (page 254)	cpu (page 263)node_id (page 264)	5.3.0	The temperature for each physical CPU, in degrees Celsius
qumulo_disk_is_unhealthy	gauge (page 254)	disk_type (page 263)drive_bay (page 263)node_id (page 264)	5.3.0	The health of each disk in the cluster, ranging from 0 (the disk is healthy) to 1 (the disk is unhealthy)

Metric Name	Metric Type	Labels	Supported from Qumulo Core Version	Description
<pre>qumulo_disk_operation _latency_seconds</pre>	histogram (page 255)	disk_type (page 263)drive_bay (page 263)io_type (page 264)node_id (page 264)	5.3.0	The total latency for disk I/O op- erations
qumulo_fan_speed_rpm	gauge (page 254)	fan (page 263)node_id (page 264)	5.3.0	The fan speed, in RPM
qumulo_fs_capacity_bytes	gauge (page 254)	_	5.3.0	The total cluster space, in bytes
<pre>qumulo_fs_directory _tree_entries</pre>	gauge (page 254)	entry_type (page 263)path (page 264)	5.3.0	The number of file system objects on the cluster, sorted by object type
<pre>qumulo_fs_directory _used_bytes</pre>	gauge (page 254)	path (page 264)usage_type (page 265)	5.3.0	The amount of space that object types use, in bytes
qumulo_fs_free_bytes	gauge (page 254)	_	5.3.0	The free space on the cluster, in bytes
qumulo_fs_snapshots	gauge (page 254)	_	5.3.0	The number of snapshots on the cluster
qumulo_ldap_lookup _request_errors	counter (page 254)	domain_url (page 263)server_url (page 265)	5.3.0	The total number of LDAP requests that resulted in an error

Metric Name	Metric Type	Labels	Supported from Qumulo Core Version	Description
qumulo_ldap_lookup _request_latency_seconds	histogram (page 255)	domain_url (page 263)server_url (page 265)	5.3.0	The total latency of LDAP re- quests
qumulo_ldap_lookup _requests	counter (page 254)	domain_url (page 263)server_url (page 265)	5.3.0	The total number of completed LDAP requests
qumulo_ldap_operation _errors	counter (page 254)	domain_url (page 263)	5.3.0	The total num- ber of LDAP op- erations that re- sulted in an er- ror
<pre>qumulo_ldap_operation _latency_seconds</pre>	histogram (page 255)	domain_url (page 263)	5.3.0	The total latency for LDAP opera- tions
qumulo_ldap_operations	counter (page 254)	domain_url (page 263)	5.3.0	The total num- ber of complet- ed LDAP opera- tions
qumulo_memory_correctable _ecc_errors	counter (page 254)	node_id (page 264)	5.3.0	The total number of memory errors that Qumulo Core corrected automatically
qumulo_network_interface _is_down	gauge (page 254)	bond (page 262)interface (page 264)role (page 265)node_id (page 264)	5.3.0	The interface status, 0 (interface is up) or 1 (interface is down)

Metric Name	Metric Type	Labels	Supported from Qumulo Core Version	Description
<pre>qumulo_network_interface _link_speed_bits_per_second</pre>	gauge (page 254)	bond (page 262)interface (page 264)role (page 265)node_id (page 264)	5.3.0	The negotiated link speed for the specified interface
qumulo_network_interface _receive_errors	counter (page 254)	bond (page 262)interface (page 264)role (page 265)node_id (page 264)	5.3.0	The total number of receive errors on the specified interface
<pre>qumulo_network_interface _received_bytes</pre>	counter (page 254)	bond (page 262)interface (page 264)role (page 265)node_id (page 264)	5.3.0	The total bytes received on the specified interface
<pre>qumulo_network_interface _received_packets</pre>	counter (page 254)	bond (page 262)interface (page 264)role (page 265)node_id (page 264)	5.3.0	The total number of packets received on the specified interface
<pre>qumulo_network_interface _transmit_errors</pre>	counter (page 254)	bond (page 262)interface (page 264)role (page 265)node_id (page 264)	5.3.0	The total number of transmission errors on the specified interface

Metric Name	Metric Type	Labels	Suppor- ted from Qumulo Core Version	Description
<pre>qumulo_network_interface _transmitted_bytes</pre>	counter (page 254)	bond (page 262)interface (page 264)role (page 265)node_id (page 264)	5.3.0	The total number of bytes transmitted on the specified interface
qumulo_network_interface _transmitted_packets	counter (page 254)	bond (page 262)interface (page 264)role (page 265)node_id (page 264)	5.3.0	The total number of packets transmitted on the specified interface
qumulo_power_supply _is_unhealthy	gauge (page 254)	location (page 264)node_id (page 264)	5.3.0	PSU health, 0 (healthy) or 1 (unplugged, removed, or unresponsive)
<pre>qumulo_protocol_client _connections</pre>	counter (page 254)	protocol (page 264)	5.3.0	The total number of clients that have connected to the specified protocol
<pre>qumulo_protocol_client _disconnections</pre>	counter (page 254)	protocol (page 264)	5.3.0	The total number of clients that have disconnected from the specified protocol
<pre>qumulo_protocol_operation _bytes</pre>	counter (page 254)	data_type (page 263)io_type (page 264)op_name (page 264)protocol (page 264)	5.3.0	The total bytes that protocol operations have transferred

Metric Name	Metric Type	Labels	Suppor- ted from Qumulo Core Version	Description
<pre>qumulo_protocol_operation _latency_seconds</pre>	histogram (page 255)	data_type (page 263)io_type (page 264)op_name (page 264)protocol (page 264)	5.3.0	The total latency for protocol op- erations
qumulo_protocol_operations	counter (page 254)	data_type (page 263)io_type (page 264)op_name (page 264)protocol (page 264)	5.3.0	The total number of completed protocol operations
<pre>qumulo_quorum_node_is _offline</pre>	gauge (page 254)	node_id (page 264)	5.3.0	The online status for each node in the cluster, 0 (node online) or 1 (node offline)
qumulo_time_is_not_synchronizing	gauge (page 254)	node_id (page 264)	5.3.0	The time synchronization status for each node in the cluster, 0 (time is synchronized) or 1 (time isn't synchronized)

Available Labels

The following table lists metric label names, possible values, and descriptions.

Label Name	Possible Values	Description
bond	bond0bond1	The bond to which a network interface belongs

Label Name	Possible Values	Description
cpu	A non-negative integer	The CPU index in the node
data_type	 data: Read or write operations on the data of a file. metadata: Operations (such as look up, stat, or getattr) unrelated to a file's data none: Operations that operate on neither the file data nor the metadata. Note 	The data type that an operation transfers
	The protocol often requires the se operations for session negot iation and authentication.	
disk_type	hdd: Hard Disk Drivessd: Solid-State Drive	The underlying storage type
domain_url	An Active Directory domain (for example, my-domain.com) or an LDAP bind URI (for example, ldap://my-server.my-domain.com)	The URL of the domain
drive_bay	A drive bay name. For example: b3, 1.1	The physical drive bay in the chassis.
entry_type	alternate_data_streamdirectoryfileothersymlink	The file system object type
fan	A fan name, for example system fan 1	The fan name

Label Name	Possible Values	Description
interface	An interface name, for example eth0	The interface name
io_type	 none read wait: A blocking operation that tak es an indeterminate amount of tim e write 	The I/O that an operation performs
location	A location on the chassis, for example left or right	The location on the chassis. ① Note For PSU, this location is relative to the back of the node.
node_id	A positive integer that represents a node ID in the cluster.	A value that differentiates between the different nodes in a cluster
op_name	Any operation name, including NFSv3, NFSv4.1, SMBv2, SMBv3 or FTP	The recorded operation
path	Slash (/)	The path to a directory in the file system
protocol	nfs: NFSv3 or NFSv4.1smb2: SMBv2 or SMBv3ftp	The protocol of the recorded operation

Label Name	Possible Values	Description
role	frontendbackend	The role of the interface
		frontend includes protocol, management, and replication traffic. backend includes all intra-node communications.
server_url	A hostname (for example, ad.my-do-main.com) or an IP address	The URL of a remote server
usage_type	datametadatasnapshot	The data type that uses space